COMP201 Final 2021 Final Exam Duration: 120 minutes	Student ID: Lab Section:		
First Name(s):	_ Last Name: _		
Do not turn this page un In the meantime, please re Good	•		
This exam contains 7 multi-part questions and you have earn 100 marks.	e 120 minutes to	HC: # 1:	/ 1/ 16
 When the exam begins, please write your stude lecture section on top of this page, and sign the a code given below. Check that the exam booklet contains a total of 11 this one. 	cademic integrity	# 2: # 3: # 4:	/ 15 // 12 // 12 // 12
 This exam is an open book and notes exam. Show all work, as partial credit will be given segraded not only on the correctness and efficiency but also on your clarity that you express it. Be near 	of your answers,		/ 18 / 14 / 100
I hereby declare that I have completed this exam individual lateral forms are applied that only the below listed sources are applied (i) Coursebook, (ii) All material that is made available to students via Bin (iii) Notes taken by me during lectures. I have not used, accessed or taken any unpermitted informations to me	oproved to be used o	during this open-s	source quiz:
belongs to me.	G	'anatuvo	

Question 1A. Calling Functions in Assembly [16 Points]

In the following, you are provided some assembly code generated by compiling myprog.c using gcc compiler.

```
#include <stdio.h>
#include <stdib.h>
#include <string.h>

void mystery(char *str, int i, int n) {
   if (i >= n)
        return;
   str[i] += 'a'-'A';
   str++;
   mystery(str, i+1, n-1);
}

int main(int argc, char *argv[]) {
   char *str = argv[1];
   mystery(str, 1, 10);
   printf("%s", str);
   return 0;
}
```

```
mystery:
 cmpl %edx, %esi
  jge .L5
  subq $8, %rsp
 movslq %esi, %rax
 addb $32, (%rdi,%rax)
  addq $1, %rdi
 subl $1, %edx
 addl $1, %esi
 callq mystery
 addq $8, %rsp
.L5:
 rep reta
.LC0:
  .string "%s"
main:
 pushq %rbx
 movq 8(%rsi), %rbx
 movl $10, %edx
 movl $1, %esi
 movq %rbx, %rdi
 callq mystery
  movq %rbx, %rsi
 movl $.LC0, %edi
 movl $0, %eax
  callq printf
  movl $0, %eax
  popq %rbx
  reta
```

(a) [1 POINTS] What will the program output when it is called with ./myprog HAPPYNEWYEAR from the console?

HaPpYnEwYeAR

(b) [12 Points] Assume that the initial values of the stack pointer %rsp and %rsi have the values 0xFFF380 and 0x400540 just before the first instruction of main is executed. Fill in this table to give the contents of registers immediately before the first instruction of mystery is executed.

	%rdi	%esi	%edx	%rsp
First call to mystery	0x400548	0x1 / \$1 / 1	0xA / \$10 / 10	0xFFF370
Second call to mystery	0x400549	0x2 / \$2 / 2	0x9 / \$9 / 9	0xFFF360
Third call to mystery	0x40054a	0x3 / \$3 / 3	0x8 / \$8 / 8	0xFFF350

(c) [1 POINTS] Does main perform something to save and restore any caller saved registers? Yes/No

No

(d) [1 Points] Does main perform something to save and restore any callee saved registers? Yes/No

Yes

(e) [1 Points] Does main obey the x86-64/Linux calling conventions? Yes/No

Yes

Question 1B. Calling Functions in Assembly [16 Points]

In the following, you are provided some assembly code generated by compiling myprog.c using gcc compiler.

```
// myprog.c
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
void mystery(char *str, int i, int n) {
  if (i >= n)
    return;
  str[i] += 'a'-'A';
  str++;
  mystery(str, i+1, n-1);
                                                        .L5:
int main(int argc, char *argv[]) {
                                                        .LC0:
  char *str = argv[1];
  mystery(str, 2, 11);
                                                        main:
  printf("%s", str);
  return 0;
```

```
mystery:
  cmpl %edx, %esi
  jge .L5
  subq $8, %rsp
  movslq %esi, %rax
  addb $32, (%rdi,%rax)
  addq $1, %rdi
  subl $1, %edx
  addl $1, %esi
  callq mystery
  addq $8, %rsp
  rep reta
  .string "%s"
  pushq %rbx
  movq 8(%rsi), %rbx
  movl $11, %edx
  movl $2, %esi
  movq %rbx, %rdi
  callq mystery
  movq %rbx, %rsi
  movl $.LC0, %edi
  movl $0, %eax
  callq printf
  movl $0, %eax
  popq %rbx
  reta
```

(a) [1 POINTS] What will the program output when it is called with ./myprog HAPPYNEWYEAR from the console?

HApPyNeWyEaR

(b) [12 Points] Assume that the initial values of the stack pointer %rsp and %rsi have the values 0xFFF388 and 0x400548 just before the first instruction of main is executed. Fill in this table to give the contents of registers immediately before the first instruction of mystery is executed.

	%rdi	%esi	%edx	%rsp
First call to mystery	0x400550	0x2 / \$2 /2	0xB / \$11 / 11	0xFFF378
Second call to mystery	0x400551	0x3 / \$3 /3	0xA / \$10 / 10	0xFFF368
Third call to mystery	0x400552	0x4 / \$4 / 4	0x9 / \$9 /9	0xFFF358

(c) [1 Points] Does mystery perform something to save and restore any caller saved registers? Yes/No

No

(d) [1 Points] Does mystery perform something to save and restore any callee saved registers? Yes/No

No

(e) [1 Points] Does main obey the x86-64/Linux calling conventions? Yes/No

Yes

Question 2A. Data and Stack Frames [15 Points]

(a) [6 Points] Consider the following C program, in which A, B and C are constants expressed with #define directives, and the corresponding assembly code is generated by the gcc compiler. By inspecting the assembly code, try to find the values of B and C.

```
short arr1[A][B];
int arr2[B][C];

void update_values(int x, int y, int z) {
    arr2[x][y] = 3*arr1[y][z];
}
```

```
update_values:
  movslq %edx, %rdx
  movslq %esi, %rsi
  leaq (%rsi, %rsi, 2), %rax
  addq %rax, %rdx
  movswl arr1(%rdx, %rdx), %eax
  leal (%rax, %rax, 2), %edx
  movslq %edi, %rdi
  leaq (%rsi, %rdi, 4), %rax
  movl %edx, arr2(, %rax, 4)
  ret
```

```
B = 3
C = 4
```

(a) [9 Points] Determine the offset of each field and the total size (in bytes) of the structure given below, considering the alignment requirements of a 64-bit machine.

Structure	a	b	С	d	е	f	g	Total
struct mystruct {	0	4	6	24	32	40	48	56
int a;								
short b;								
char c[14];								
double d;								
short e;								
struct mystruct *f								
float g;								
}								

What is the size (in bytes) of the structure if the fields in part (b) are rearranged to have minimum wasted space? 48

Question 2B. Data and Stack Frames [15 POINTS]

(b) [6 Points] Consider the following C program, in which A, B and C are constants expressed with #define directives, and the corresponding assembly code is generated by the gcc compiler. By inspecting the assembly code, try to find the values of B and C.

```
short arr1[A][B];
int arr2[B][C];

void update_values(int x, int y, int z) {
    arr2[x][y] = 3*arr1[y][z];
}
```

```
update_values:
   movslq %edx, %rdx
movslq %esi, %rsi
leaq (%rdx,%rsi,4), %rax
movswl arr1(%rax,%rax), %eax
leal (%rax,%rax,2), %ecx
movslq %edi, %rdi
leaq (%rdi,%rdi,2), %rdx
leaq (%rdx,%rdx), %rax
addq %rax, %rsi
movl %ecx, arr2(,%rsi,4)
ret
```

```
B = 4
C = 6
```

(c) [9 Points] Determine the offset of each field and the total size (in bytes) of the structure given below, considering the alignment requirements of a 64-bit machine.

Structure	а	b	С	d	е	f	g	Total
struct mystruct {	0	4	8	16	32	40	44	48
short a;								
float b;								
double c;								
char d[9];								
struct mystruct *e;								
short f;								
int g;								
}								

What is the size (in bytes) of the structure if the fields in part (b) are rearranged to have minimum wasted space? 40

Question 3. Buffer Overflow [12 Points]

Thomas A. Anderson, also known as with the alias "Neo", is a highly skilled programmer who works at MetaCortex software company in his daily job. But at night he is obsessed with hacking the Matrix. Finally, he located the source code of a test program (test.c) which he thinks it is the key to unlock accessing the Matrix code. This function first retrieves the stored passcode and then compares the input string to check whether they match with each other or not.

```
#include <stdio.h>
...

void enter_matrix() {
    char passcode[8];
    char input[8];
    retrieve_passcode(passcode);
    gets(input);
    if (strcmp(input, passcode) == 0)
        found_the_rabbit();
}

int main(void) {
    printf("Password: ");
    enter_matrix();
    return 0;
}
```

Below, you are provided some part of the assembly code generated by compiling test.c using gcc compiler:

```
00000000004006a1 <found_the_rabbit>:
  4006a1:
                55
                                         push
                                                %rbp
0000000004006b1 <enter_matrix>:
  4006b1:
            55
                                                %rbp
                                         push
  4006b2:
               48 89 e5
                                         mov
                                                %rsp,%rbp
              48 83 ec 20
  4006b5:
                                                $0x20,%rsp
                                         sub
              48 8d 45 f0
  4006b9:
                                         lea
                                                -0x10(%rbp),%rax
  4006bd:
               48 89 c7
                                         mov
                                                %rax,%rdi
               e8 b8 ff ff ff
                                         callq 40067d <retrieve_passcode>
  4006c0:
  4006c5:
               48 8d 45 e0
                                         lea
                                                -0x20(%rbp),%rax
  4006c9:
               48 89 c7
                                                %rax,%rdi
                                         mov
                e8 9f fe ff ff
                                         callq 400570 <gets@plt>
  4006cc:
  4006d1:
                48 8d 55 f0
                                         lea
                                                -0x10(%rbp),%rdx
  4006d5:
                48 8d 45 e0
                                         lea
                                                -0x20(%rbp),%rax
  4006d9:
                48 89 d6
                                                %rdx,%rsi
                                         mov
                48 89 c7
  4006dc:
                                         mov
                                                %rax,%rdi
                                         callq 400550 <strcmp@plt>
  4006df:
                e8 6c fe ff ff
  4006e4:
                85 c0
                                         test
                                                %eax,%eax
                75 0a
                                                4006f2 <enter matrix+0x41>
  4006e6:
                                         jne
                b8 00 00 00 00
  4006e8:
                                         mov
                                                $0x0,%eax
                e8 af ff ff ff
                                         callq 4006a1 <found_the_rabbit>
  4006ed:
  4006f2:
                c9
                                         leaveq
  4006f3:
                                         retq
                c3
00000000004006f4 <main>:
  4006f4:
                55
                                         push
                                                %rbp
                48 89 e5
  4006f5:
                                         mov
                                                %rsp,%rbp
                bf d9 07 40 00
  4006f8:
                                         mov
                                                $0x4007d9,%edi
                b8 00 00 00 00
  4006fd:
                                         mov
                                                $0x0,%eax
                e8 29 fe ff ff
  400702:
                                         callq 400530 <printf@plt>
                b8 00 00 00 00
                                                $0x0,%eax
  400707:
                                         mov
  40070c:
                e8 a0 ff ff ff
                                         callq
                                                4006b1 <enter_matrix>
  400711:
                b8 00 00 00 00
                                                $0x0,%eax
                                         mov
                5d
  400716:
                                         pop
                                                %rbp
  400717:
                                         retq
```

(a) [4 Points] Which line is the root cause of buffer overflow attack here.

```
gets(input);
```

- (b) [4 Points] When the buffer overflow occurs, which one of the following addresses gets overwritten on the stack?
 - 1. 0x00000000004006a1
 - 2. 0x00000000004006b1
 - 3. 0x00000000004006d1
 - 4. None of the above

The gets function can write to the stack locations where both the passcode and input are stored, or you can directly try to change the next instruction address at main (0x0000000000011)

(c) [4 Points] Assume that you are working at the directory in which the test program exists and that you execute the following commands from the terminal:

```
./hex2raw < exploit | ./test</pre>
```

Write down the content of input (in hexadecimal digits) so that it performs the attack that will perform a function call to found_the_rabbit function.

Note: In your answer, if you want to repeat a specific random byte sequence, you can use the expression (n) to denote consecutive bytes of size n bytes to refer to that byte sequence.

(15) 00 (15) 00 (Two identical 16 byte strings – including the NULL terminating symbol)

Also true: Directly write onto the next instruction address in main pushed to the stack when call enter_matrix

00(40) a1 06 40 00 00 00 00 00

Question 4. Locality [12 POINTS]

Consider the following C functions which calculate the pairwise squared Euclidean distance matrix consisting of n rows where each row corresponds to the coordinates of an n-dimensional point.

For example, for the matrix $A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 10 \end{bmatrix}$, the distance matrix is calculated as $D = \begin{bmatrix} 0 & 27 & 121 \\ 27 & 0 & 34 \\ 121 & 34 & 0 \end{bmatrix}$.

```
void pdist1(int N, int A[N][N], int D[N][N]) {
  int i, j, k;
  for (i = 0; i < N; i++) {
      for (j = 0; j < N; j++) {
          int temp = 0;
          for (k = 0; k < N; k++) {
              temp += (int)pow(A[i][k] - A[j][k], 2);
          }
          D[i][j] = temp;
      }
  }
}
void pdist2(int N, int A[N][N], int D[N][N]) {
  int i, j, k;
  for (j = 0; j < N; j++) {
      for (i = 0; i < N; i++) {
          int temp = 0;
          for (k = 0; k < N; k++) {
              temp += (int)pow(A[i][k] - A[j][k], 2);
          D[i][j] = temp;
      }
 }
 }
```

- (a) [3 POINTS] The functions pdist1 and pdist2 both accomplish the same task. Which one is more cache friendly?
- 1. pdist1 is more cache friendly than pdist2
- 2. pdist2 is more cache friendly than pdist1
- 3. They perform similar in term of cache utilization.

Here one needs to think about the worst case where the size of the matrix is very huge, hence as we discussed in the class, you only need to focus on the inner most loop.

- (b) [3 POINTS] The runtime performances of pdist1 and pdist2 can be improved with loop unrolling [T/F]
- (c) [3 Points] The memory writes to local variable temp has bad temporal locality [T/F]
- (d) [3 POINTS] The memory accesses to A[i][k] has stride-N reference pattern and thus have good spatial locality [T/F] (it has stride-1 pattern)

Question 5A. Cache Memories [12 POINTS]

Suppose that you are working on a system with the following specifications:

- The device is a 24-bit machine, i.e., the physical addresses are 24 bits wide.
- The memory is byte addressable, and memory accesses are to 1-byte words.
- The device contains a 4-way set associative cache of size 4 KB with a block size of 32 bytes.
- (a) [2 Points] How many bits are reserved for the cache block offset? 5 bits
- (b) [2 POINTS] How many bits are used for the index? 5 bits
- (c) [2 POINTS] How many bits are used for the tag?

Assume that the cache is initially empty when executing the following code:

```
int arr[1024][8];
for (int i=0; i<1024; i+=2) {
    arr[i][0] += arr[i][4];
}</pre>
```

(d) [2 POINTS] What is the hit rate?

- 2/3 (There is 2 read and 1 write operations)
- (e) [2 POINTS] What is the number of hits?
- 1024
- (f) [2 Points] The hit rate improves if one increases the associativity of the cache to 8 while keeping the other parameters the same. [T/F]

Question 5B. Cache Memories [12 POINTS]

Suppose that you are working on a system with the following specifications:

- The device is a 16-bit machine, i.e., the physical addresses are 16 bits wide.
- The memory is byte addressable, and memory accesses are to 1-byte words.
- The device contains a 2-way set associative cache of size 8 KB with a block size of 32 bytes.
- (a) [2 POINTS] How many bits are reserved for the cache block offset? 5 bits
- (b) [2 Points] How many bits are used for the index? 7 bits
- (c) [2 POINTS] How many bits are used for the tag?

 4 bits

Assume that the cache is initially empty when executing the following code:

```
int arr[512][8];
for (int i=0; i<512; i+=4) {
    arr[i][0] += arr[i][4];
}</pre>
```

(d) [2 Points] What is the hit rate?

- 2/3 (There are 2 reads and 1 write, only the first is miss)
- (e) [2 Points] What is the number of hits?
- 256
- (f) [2 Points] The hit rate improves if one increases the associativity of the cache to 8 while keeping the other parameters the same. [T/F]

Question 6. Assembly and Linking [18 POINTS]

(a) [12 POINTS] Fill in the missing parts of the C functions based on the corresponding Assembly code generated by gcc. Hint: The ASCII codes of the characters 'A', 'C' and 'Z' are 65, 67 and 90, respectively.

myfile1.c

```
#include <string.h>
extern char c;

int strfun(char *str) {
   const int x = 32;
   char * ret;
   ret = strchr(str, c);
   if (ret[0]>='A' && ret[0]<='Z') {
      // *ret is also true
      ret[0] = ret[0] & ~x;
      return 1;
   }

   return 0;
}</pre>
```

```
strfun:
  subq $8, %rsp
 movsbl c(%rip), %esi
  call strchr
  movzbl (%rax), %edx
  leal -65(%rdx), %ecx
  cmpb $25, %cl
  ja .L3
  and1 $-33, %edx
  movb %dl, (%rax)
 movl $1, %eax
 jmp .L2
.L3:
 movl $0, %eax
.L2:
  addq $8, %rsp
  ret
```

myfile1.s

```
myfile2.c
char c = 'C';
int strfun(char *);
typedef struct mystr {
  struct mystr *a;
  char b[10];
  char c;
} mystr;
void strfun2(mystr *s, int b) {
  s \rightarrow c = 'C';
  while (s) {
    if (strfun(s->b)) {
       s->c--; // or any equivalent expr.
    s = s \rightarrow a;
  }
  return;
}
```

```
myfile2.s
strfun2:
  pushq %rbx
  movq %rdi, %rbx
  movb $67, 18(%rdi)
  imp .L2
.L4:
  leaq 8(%rbx), %rdi
  call strfun
  testl %eax, %eax
  je .L3
  movzbl 18(%rbx), %eax
 subl $1, %eax
 movb %al, 18(%rbx)
.L3:
 movq (%rbx), %rbx
 testq %rbx, %rbx
  jne .L4
  popq %rbx
  ret
```

(b) [3 Points] Fill in the following table to name the strong, weak and external symbols defined in the C files. If there is no such definition just write 'None'. If there are more than one, separate the names with commas.

The strong	The weak	The external	e external The strong		The external
symbols defined	symbols defined	symbols defined	symbols defined	symbols defined	symbols defined
in myfile1.c	in myfile1.c	in myfile1.c	in myfile2.c	in myfile2.c	in myfile2.c
strfun	None	С	strfun2, c	None	strfun

(c) [3 Points] Suppose that you have another file (main.c) which includes the main function, and two static libraries libmyfile1.a and libmyfile2.a created from myfile1.c and myfile2.c, respectively. The correct way to compile the program is gcc -o a.out main.c -L. -lmyfile2 -lmyfile1 [T/F]

myfile2 should come before myfile1 to correctly resolve the symbol strfun.

Question 7. Heap Allocators [14 Points]

Consider the following state of the heap with <u>an explicit free-list allocator with coalescing</u> after performing the following memory allocation operations. The header is of 8 bytes and the allocated chunks of memory should be multiples of 16 bytes, and <u>a best-fit policy is employed</u>. For the sake of simplicity, the memory addresses are given in decimal numbers.

```
void *a = malloc(20);
void *b = malloc(24);
void *c = malloc(28);
void *d = malloc(16);
          Address:
                   72
                        80
                             88
                                  96
                                       104
                                            112
                                                120
                                                     128
                                                           136
                                                                144
                                                                     152
                                                                          160
                                                                              168
                                                                                   176
         Memory:
                                                                                       b
                                                                                            b
                                                                                                 b
                                          а
                                     24
                             216
                                  224
                                                                272
                   200
                        208
                                       232
                                            240
                                                248
                                                      256 264
                                                                     280
                                                                          288
                                                                              296
                                                                                    304
          Address:
                                                                                        312
                      Used
                                                                  Used
                            С
                                 С
                                     С
                                          С
                                                                        d
                                                                             d
         Memory:
```

- (a) [2 Points] What is the amount of internal fragmentation (in bytes) for the current state of the heap for the allocated parts of the heap memory? In your answer, please do count the headers.
 - 8 // I made a typo here, I wanted to say (do not count the headers). So I also consider 40 or 72 as true here!
- (b) [2 POINTS] Suppose that the next call to the heap allocator is malloc(24). Specify the address that the malloc function returns?

248 or 304

(c) [2 Points] What if the first-fit approach were used instead for the request in part (b)?

136

(d) [2 Points]) Suppose that the next call to the heap allocator is free(a). After this operation, can malloc (60) request be satisfied? Yes/No

Yes. Because of coalescing, there is enough free space.

(e) [2 Points] Does the heap currently have external fragmentation?

Yes. There are non-consecutive free blocks.

- (f) [2 Points] Considering the layout of the heap shown at the top, what is the maximum size one can ask for the heap allocator for realloc(c, size)?
 - 64. The explicit free-list allocator does not perform in-place realloc. Hence, the maximum such chunk is of 64 bytes.
- (g) [2 Points] Suppose that you have switched to using implicit free-list allocator. Does this change increase the throughput for allocation requests or not?

No. Because it is linear in the number of free+allocated chunks, which is worse than explicit free list allocator.