

COMP201 Fall 2020 Final Exam
Duration: **120 minutes**

Student ID: _____

Lab Section: _____

First Name(s): _____ Last Name: _____

*Do **not** turn this page until you are told to do so.
In the meantime, please read the instructions below.*

Good Luck!

This exam contains **8 multi-part questions** and you have **120 minutes** to earn 100 marks.

- When the exam begins, please write your student ID, name and lecture section on top of this page, and sign the academic integrity code given below.
- Check that the exam booklet contains a total of **19 pages**, including this one.
- This exam is an **open book and notes exam**.
- Show all work, as partial credit will be given since you will be graded not only on the correctness and efficiency of your answers, but also on your clarity that you express it. Be neat.

1: _____ / 12
2: _____ / 14
3: _____ / 12
4: _____ / 14
5: _____ / 15
6: _____ / 13
7: _____ / 10
8: _____ / 10
TOTAL: _____ / 100

I hereby declare that I have completed this exam individually, without support from anyone else.

I hereby accept that only the below listed sources are approved to be used during this open-source exam:

- (i) Coursebook,
- (ii) All material that is made available to students via Blackboard for this course, and
- (iii) Notes taken by me during lectures.

I have not used, accessed or taken any unpermitted information from any other source. Hence, all effort belongs to me.

Signature: _____

Question 1A. Calling Functions in Assembly [12 POINTS]

In the following, you are provided some assembly generated by compiling a C program using gcc compiler.

```

000000000400540 <mystery>:
400540:    41 54                push    %r12
400542:    31 c0                xor     %eax,%eax
400544:    49 89 fc             mov     %rdi,%r12
400547:    55                  push    %rbp
400548:    89 d5                mov     %edx,%ebp
40054a:    53                  push    %rbx
40054b:    89 f3                mov     %esi,%ebx
40054d:    03 5f 04             add     0x4(%rdi),%ebx
400550:    bf 10 06 40 00       mov     $0x400610,%edi
400555:    89 de                mov     %ebx,%esi
400557:    e8 b4 fe ff ff       callq   400410 <printf@plt>
40055c:    41 03 6c 24 08       add     0x8(%r12),%ebp
400561:    bf 10 06 40 00       mov     $0x400610,%edi
400566:    31 c0                xor     %eax,%eax
400568:    89 ee                mov     %ebp,%esi
40056a:    e8 a1 fe ff ff       callq   400410 <printf@plt>
40056f:    8d 04 2b             lea     (%rbx,%rbp,1),%eax
400572:    5b                  pop     %rbx
400573:    5d                  pop     %rbp
400574:    41 5c                pop     %r12
400576:    c3                  retq

```

(a) [3 POINTS] How many parameters does this function take?

3. It uses %rdi, %esi, %edx

(b) [3 POINTS] Does this function return any value?

Yes. It uses %eax.

(c) [3 POINTS] What is the purpose of the push instruction at the first line?

- A: allocate space for local variables
- B: save parameters
- C: save caller-saved registers**
- D: save callee-saved registers

(d) [3 POINTS] What is the type of the variable the register %rdi stores?

Array or Pointer

Question 1B. Calling Functions in Assembly [12 POINTS]

In the following, you are provided some assembly generated by compiling a C program using gcc compiler.

```
000000000400540 <mystery>:
400540: 41 54                push    %r12
400542: 31 c0                xor     %eax,%eax
400544: 49 89 fc             mov     %rdi,%r12
400547: 55                  push    %rbp
400548: 89 f5                mov     %esi,%ebp
40054a: 53                  push    %rbx
40054b: 8b 5f 04             mov     0x4(%rdi),%ebx
40054e: bf 10 06 40 00       mov     $0x400610,%edi
400553: 01 f3                add     %esi,%ebx
400555: 89 de                mov     %ebx,%esi
400557: e8 b4 fe ff ff       callq   400410 <printf@plt>
40055c: 41 03 6c 24 08       add     0x8(%r12),%ebp
400561: bf 10 06 40 00       mov     $0x400610,%edi
400566: 31 c0                xor     %eax,%eax
400568: 0f af eb             imul    %ebx,%ebp
40056b: 89 ee                mov     %ebp,%esi
40056d: e8 9e fe ff ff       callq   400410 <printf@plt>
400572: 8d 04 2b             lea     (%rbx,%rbp,1),%eax
400575: 5b                  pop     %rbx
400576: 5d                  pop     %rbp
400577: 41 5c                pop     %r12
400579: c3                  retq
```

(a) [3 POINTS] How many parameters does this function take?

2. It uses %rdi, %esi

(b) [3 POINTS] Does this function return any value?

Yes. It uses %eax.

(c) [3 POINTS] What is the purpose of the push instruction at the first line?

- A: allocate space for local variables
- B: save parameters
- C: save caller-saved registers**
- D: save callee-saved registers

(d) [3 POINTS] What is the type of the variable the register %rdi stores?

Array or Pointer

Question 2. x86-64 Procedures and Code Optimization [14 POINTS]

In the following, you are given some code fragments written in C and the corresponding assembly codes generated by the GCC compiler (some parts of the assembly code may be intentionally omitted). In these questions, each code fragment contains a variable, constant, or a function named `mando` (short for Mandalorian). As an officer of the New Republic, your task is to locate `mando` at a certain point marked with `***`. In particular, you are asked to provide either an assembly expression or constant (like a memory address) that holds the value of `mando`. Below is an example question and the corresponding answer to make what is expected more clear.

Where is Mando?

```
int identity(int mando) {
    return mando;
}
```

```
0000000004004ed <identity>:
 4004ed:      55                push    %rbp
 4004ee:      48 89 e5          mov     %rsp,%rbp
 4004f1:      89 7d fc          mov     %edi,-0x4(%rbp)
 4004f4:      8b 45 fc          mov     -0x4(%rbp),%eax
                ***
 4004f7:      5d                pop     %rbp
 4004f8:      c3              retq
```

Answer: Any of the expressions `%edi`, `-0x4(%rbp)`, `%eax`, and `%rax` are valid answers. They all hold the value of `mando` at the marked point. Note that, however, if the asterisks came before the *first* instruction (before `4004ed`), `%edi` would be the only correct answer.

(a) [2 POINTS] Where is Mando?

```
int myfunc1(int a, int b, int mando) {
    if (a > b)
        return mando;
    else
        return 4 * mando;
}
```

```
000000000400505 <myfunc1>:
 400505:      55                push    %rbp
 400506:      48 89 e5          mov     %rsp,%rbp
 400509:      89 7d fc          mov     %edi,-0x4(%rbp)
 40050c:      89 75 f8          mov     %esi,-0x8(%rbp)
 40050f:      89 55 f4          mov     %edx,-0xc(%rbp)
                ***
 400512:      8b 45 fc          mov     -0x4(%rbp),%eax
 400515:      3b 45 f8          cmp     -0x8(%rbp),%eax
 400518:      7e 05             jle     40051f <myfunc1+0x1a>
 40051a:      8b 45 f4          mov     -0xc(%rbp),%eax
 40051d:      eb 06             jmp     400525 <myfunc1+0x20>
 40051f:      8b 45 f4          mov     -0xc(%rbp),%eax
 400522:      c1 e0 02          shl     $0x2,%eax
 400525:      5d                pop     %rbp
 400526:      c3              retq
```

Answer: `%edx` or `-0xc(%rbp)`

(b) [2 POINTS] Where is Mando?

```
int int_array_sum(int* a, int len) {
    int mando = 0;
    for (int i = 0; i < len; i++)
        mando += a[i];
    return mando;
}
```

```
000000000400527 <int_array_sum>:
400527:    55                push    %rbp
400528:    48 89 e5          mov     %rsp,%rbp
40052b:    48 89 7d e8        mov     %rdi,-0x18(%rbp)
40052f:    89 75 e4          mov     %esi,-0x1c(%rbp)
400532:    c7 45 fc 00 00 00 movl    $0x0,-0x4(%rbp)
400539:    c7 45 f8 00 00 00 movl    $0x0,-0x8(%rbp)
400540:    eb 1d            jmp     40055f <int_array_sum+0x38>
400542:    8b 45 f8          mov     -0x8(%rbp),%eax
400545:    48 98            cltq
400547:    48 8d 14 85 00 00 lea     0x0(,%rax,4),%rdx
40054e:    00
40054f:    48 8b 45 e8        mov     -0x18(%rbp),%rax
***
400553:    48 01 d0          add     %rdx,%rax
400556:    8b 00            mov     (%rax),%eax
400558:    01 45 fc          add     %eax,-0x4(%rbp)
40055b:    83 45 f8 01        addl    $0x1,-0x8(%rbp)
40055f:    8b 45 f8          mov     -0x8(%rbp),%eax
400562:    3b 45 e4          cmp     -0x1c(%rbp),%eax
400565:    7c db            jl      400542 <int_array_sum+0x1b>
400567:    8b 45 fc          mov     -0x4(%rbp),%eax
40056a:    5d              pop     %rbp
40056b:    c3              retq
```

Answer: -0x4(%rbp)

(c) [2 POINTS] Where is Mando?

```
int int_array_shift(int* a, int b, int mando, int len) {
    for (int i = 0; i < len; i++)
        a[i] += b * (mando + 2);
    return mando;
}
```

```
00000000040056c <int_array_shift>:
INSTRUCTIONS OMITTED
4005b4:    83 c2 02          add     $0x2,%edx
4005b7:    0f af 55 e4        imul    -0x1c(%rbp),%edx
4005bb:    01 ca            add     %ecx,%edx
4005bd:    89 10            mov     %edx,(%rax)
***
4005bf:    83 45 fc 01        addl    $0x1,-0x4(%rbp)
4005c3:    8b 45 fc          mov     -0x4(%rbp),%eax
4005c6:    3b 45 dc          cmp     -0x24(%rbp),%eax
4005c9:    7c bb            jl      400586 <int_array_shift+0x1a>
```

4005cb:	8b 45 e0	mov	-0x20(%rbp),%eax
4005ce:	5d	pop	%rbp
4005cf:	c3	retq	

Answer: **-0x20(%rbp)**

(d) [2 POINTS] Where is Mando?

```
int array2d_get(int** arr2d, int row, int col) {
    int* mando = arr2d[row];
    return mando[col];
}
```

```
00000000004005d0 <array2d_get>:
 4005d0: 55                push    %rbp
 4005d1: 48 89 e5          mov     %rsp,%rbp
 4005d4: 48 89 7d e8       mov     %rdi,-0x18(%rbp)
 4005d8: 89 75 e4          mov     %esi,-0x1c(%rbp)
 4005db: 89 55 e0          mov     %edx,-0x20(%rbp)
 4005de: 8b 45 e4          mov     -0x1c(%rbp),%eax
 4005e1: 48 98            cltq
 4005e3: 48 8d 14 c5 00 00 00 lea     0x0(,%rax,8),%rdx
 4005ea: 00
 4005eb: 48 8b 45 e8       mov     -0x18(%rbp),%rax
 4005ef: 48 01 d0          add     %rdx,%rax
 4005f2: 48 8b 00          mov     (%rax),%rax
 4005f5: 48 89 45 f8       mov     %rax,-0x8(%rbp)
 4005f9: 8b 45 e0          mov     -0x20(%rbp),%eax
      ***
 4005fc: 48 98            cltq
 4005fe: 48 8d 14 85 00 00 00 lea     0x0(,%rax,4),%rdx
 400605: 00
 400606: 48 8b 45 f8       mov     -0x8(%rbp),%rax
 40060a: 48 01 d0          add     %rdx,%rax
 40060d: 8b 00          mov     (%rax),%eax
 40060f: 5d                pop     %rbp
 400610: c3                retq
```

Answer: **%rax, %eax or -0x8(%rbp)**

(e) [2 POINTS] Where is Mando?

```
int myfunc2(int x) {
    if (x > 4) {
        return mando(63 * x);
    } else {
        return mando(124 * x);
    }
}
```

```

000000000400611 <myfunc2>:
 400611:      55                push    %rbp
 400612:     48 89 e5          mov     %rsp,%rbp
 400615:     48 83 ec 08       sub     $0x8,%rsp
 400619:     89 7d fc          mov     %edi,-0x4(%rbp)
 40061c:     83 7d fc 04       cmpl    $0x4,-0x4(%rbp)
 400620:     7e 13             jle     400635 <myfunc2+0x24>
 400622:     8b 55 fc          mov     -0x4(%rbp),%edx
 400625:     89 d0             mov     %edx,%eax
                ***
 400627:     c1 e0 06          shl     $0x6,%eax
 40062a:     29 d0             sub     %edx,%eax
 40062c:     89 c7             mov     %eax,%edi
 40062e:     e8 c6 fe ff ff    callq   4004f9
 400633:     eb 16             jmp     40064b <myfunc2+0x3a>
 400635:     8b 45 fc          mov     -0x4(%rbp),%eax
 400638:     c1 e0 02          shl     $0x2,%eax
 40063b:     89 c2             mov     %eax,%edx
 40063d:     c1 e2 05          shl     $0x5,%edx
 400640:     29 c2             sub     %eax,%edx
 400642:     89 d0             mov     %edx,%eax
 400644:     89 c7             mov     %eax,%edi
 400646:     e8 ae fe ff ff    callq   4004f9
 40064b:     c9               leaveq  %eax
 40064c:     c3               retq

```

Answer: 4004f9

In the remaining, you are given an assembly code which is compiled from one of the above functions using a different compiler or a different optimization level. Your task is to determine what the given assembly code corresponds to.

(f) [2 POINTS] What's mando? Circle one.

```

000000000400540 <mando>:
 400540:     85 c9             test    %ecx,%ecx
 400542:     89 d0             mov     %edx,%eax
 400544:     7e 1d             jle     400563 <mando+0x23>
 400546:     8d 52 02          lea     0x2(%rdx),%edx
 400549:     0f af f2          imul    %edx,%esi
 40054c:     8d 51 ff          lea     -0x1(%rcx),%edx
 40054f:     48 8d 4c 97 04     lea     0x4(%rdi,%rdx,4),%rcx
 400554:     0f 1f 40 00        nopl    0x0(%rax)
 400558:     01 37             add     %esi,(%rdi)
 40055a:     48 83 c7 04       add     $0x4,%rdi
 40055e:     48 39 cf          cmp     %rcx,%rdi
 400561:     75 f5             jne     400558 <mando+0x18>
 400563:     f3 c3             repz    retq

```

myfunc1

int_array_sum

int_array_shift

array2d_get

myfunc2

(g) [2 POINTS] What's mando? Circle one.

```
0000000000400580 <mando>:
400580: 83 ff 04          cmp     $0x4,%edi
400583: 7f 0b            jg      400590 <mando+0x10>
400585: 89 f8            mov     %edi,%eax
400587: ba 7c 00 00 00    mov     $0x7c,%edx
40058c: 0f af c2         imul    %edx,%eax
40058f: c3              retq
400590: 89 f8            mov     %edi,%eax
400592: c1 e0 06         shl     $0x6,%eax
400595: 29 f8            sub     %edi,%eax
400597: c3              retq
```

myfunc1

int_array_sum

int_array_shift

array2d_get

myfunc2

Question 3A. Data and Stack Frames [12 POINTS]

- (a) [4 POINTS] Consider the following C program, in which H and J are constants expressed with #define directives, and the corresponding assembly code generated by the gcc compiler. By inspecting the assembly code, try to find the values of H and J.

```
int arr1[H][J];
int arr2[J][H];

void update_array(int x, int y) {
    arr2[y][x] = 3*arr1[x][y];
}
```

```
0000000004004f0 <update_array>:
 4004f0: 48 63 f6          movslq %esi,%rsi
 4004f3: 48 63 ff          movslq %edi,%rdi
 4004f6: 48 8d 04 76       lea     (%rsi,%rsi,2),%rax
 4004fa: 48 8d 14 87       lea     (%rdi,%rax,4),%rdx
 4004fe: 48 8d 04 7f       lea     (%rdi,%rdi,2),%rax
 400502: 48 8d 04 46       lea     (%rsi,%rax,2),%rax
 400506: 8b 04 85 60 10 60 00 mov    arr1(,%rax,4),%eax
 40050d: 8d 04 40          lea     (%rax,%rax,2),%eax
 400510: 89 04 95 80 11 60 00 mov    %eax,arr2(,%rdx,4)
 400517: c3               retq
```

H = 12

J = 6

- (b) [4 POINTS] Determine the offset of each field and the total size (in bytes) of the structure given below, considering the alignment requirements of a 64-bit machine.

Structure	d1	c1	f1	c2	s1	p1	s2	Total
<pre>struct mystruct { int d1; char c1; float f1; char c2; short s1; int *p1; short s2; }</pre>	0	4	8	12	14	16	24	32

- (c) [4 POINTS] What is the size (in bytes) of the structure if the fields in part (b) are rearranged to have minimum wasted space?

24

Question 3B. Data and Stack Frames [12 POINTS]

- (a) [4 POINTS] Consider the following C program, in which H and J are constants expressed with #define directives, and the corresponding assembly code generated by the gcc compiler. By inspecting the assembly code, try to find the values of H and J.

```
int arr1[H][J];
int arr2[J][H];

void modify_array(int x, int y) {
    arr2[y][x] = 2*arr1[x][y]-1;
}
```

```
0000000004004f0 <modify_array>:
4004f0:    48 63 ff          movslq %edi,%rdi
4004f3:    48 63 f6          movslq %esi,%rsi
4004f6:    48 8d 14 3f       lea     (%rdi,%rdi,1),%rdx
4004fa:    48 8d 04 f7       lea     (%rdi,%rsi,8),%rax
4004fe:    48 c1 e7 04       shl     $0x4,%rdi
400502:    48 29 d7          sub     %rdx,%rdi
400505:    48 01 f7          add     %rsi,%rdi
400508:    8b 14 bd 60 10 60 00 mov     arr1(,%rdi,4),%edx
40050f:    8d 54 12 ff       lea     -0x1(%rdx,%rdx,1),%edx
400513:    89 14 85 20 12 60 00 mov     %edx,arr2(,%rax,4)
40051a:    c3              retq
```

H = 8

J = 14

- (b) [4 POINTS] Determine the offset of each field and the total size (in bytes) of the structure given below, considering the alignment requirements of a 64-bit machine.

Structure	c1	s1	d1	c2	f1	s2	p1	Total
<pre>struct mystruct { char c1; short s1; int d1; char c2; float f1; short s2; int *p1; }</pre>	0	2	4	8	12	16	24	32

- (c) [4 POINTS] What is the size (in bytes) of the structure if the fields in part (b) are rearranged to have minimum wasted space.

24

Question 4. Buffer Overflow [14 POINTS]

Elliott Alderson decided to perform attacks against FSociety machines to check the integrity of their security measures. As one of the new members, Tyrell Wellick has written a password checker that he thinks is unbreakable! Below is the front-end to his program:

```
int main(int argc, char *argv[]) {
    char passwd[20];
    printf("Password:");
    scanf("%s",passwd);
    if (check(passwd)) {
        enterFSociety();
        exit(0);
    }
    printf("Sorry, you are not awake!");
    return 0;
}
```

- (a) [3 POINTS] Elliott does not believe his eyes when he sees this. He thinks that it is so brittle. Briefly explain how Elliott can hack this program using a buffer overflow attack. Please explain with 2-3 sentences.

Elliot can enter a string long enough to overflow the buffer to overwrite the main function's return address so that it points to the address of the exploit.

Tyrell then objects that it is not that easy because he runs his program on a special system where the stack is not executable. In other words, it is impossible to execute any code on the stack, making a typical buffer overflow attack is not possible.

After a minute of silence, Elliott decides he can indeed perform a *return-to-libc* attack – utilizing an already present code in the program which can enable you execute arbitrary instructions. Here is some background on this attack type:

The `system(char *command)` function is a part of C standard library and it can be used to execute the string `command` as if it is typed in the shell of the operation system.

Using `gdb`, Elliott discovers the following:

```
(gdb) print system
$1 = {<text variable, no debug info>} 0xfc3406ac <system>
```

He knows that in every executable program, the environment variables are loaded at runtime, and guess what it involves the current SHELL:

```
(gdb) print (char *) 0xfc2356e2
$2 = 0xfc2356e2 "SHELL=/bin/bash"
```

Using this information, Elliot now knows that he can launch a shell from Tyrell's program, proving that he can in fact execute any arbitrary code that he wants!

- (b) [3 POINTS] What are the addresses of the `system()` function and the string `"/bin/bash"`?

Address of the `system()` function: **0xfc3406ac**

Address of the string `"/bin/bash"`: **0xfc2356e8**

- (c) [5 POINTS] Specify what Elliott can pass as input to Tyrell's program in order to cause it to launch a shell screen. In your answer, you don't need to be too specific about the sizes or lengths, but keep in mind where the arguments go in 32-bit mode.

some garbage bytes | old %ebp | addr. of system() | 4 byte garbage | addr. of "/bin/bash"

- (d) [3 POINTS] Explain how Elliot's exploit string will enable to execute a shell on Tyrell's program.

The main function will rather return to the system() function with the argument "/bin/bash", which will open a new shell to input any arbitrary command(s).

-1: If no mention of system() or "/bin/bash".

Question 5A. Locality [15 POINTS]

Consider the following C function that calculates the sum of the rows of a 2D matrix and assume that the elements of the array `sum` have been initialized to zero.

```
void sumRows(int nCols, int nRows, int sum[nRows], int A[nRows][nCols]) {
    for (int j = 0; j < nCols; j++) {
        for (int i = 0; i < nRows; i++) {
            sum[i] = sum[i] + A[i][j];
        }
    }
}
```

(a) [3 POINTS] Does the function `sumRows` exhibit *temporal locality*?

Temporal locality with respect to code: Yes. `sum[j]` computed and used many times.
 Temporal locality with respect to array A: No (not a single element `A[i][j]` accessed more than once.)
 Temporal locality with respect to array `sum`: Yes (`sum[i]` accessed frequently inside the inner loop).

(b) [3 POINTS] Does the function `sumRows` exhibit *spatial locality*?

Spatial locality with respect to code: Yes (the codes are executed in a contiguous section of the memory)
 Spatial locality with respect to array A: No (The inner loop traverses the column elements).
 Spatial locality with respect to array `sum`: Yes (Inner loop traverse `sum` with stride 1).

Now consider the following alternative implementation `sumRows2` where the loops are permuted.

```
void sumRows2(int nCols, int nRows, int sum[nRows], int A[nRows][nCols]) {
    for (int i = 0; i < nRows; i++) {
        for (int j = 0; j < nCols; j++){
            sum[i] = sum[i] + A[i][j];
        }
    }
}
```

(c) [3 POINTS] Does the function `sumRows2` exhibit *temporal locality*?

Temporal locality with respect to code: Yes. `sum[j]` computed and used many times.
 Temporal locality with respect to array A: No (not a single element `A[i][j]` accessed more than once.)
 Temporal locality with respect to array `sum`: Yes (`sum[i]` accessed frequently inside the inner loop).

(d) [3 POINTS] Does the function `sumRows2` exhibit *spatial locality*?

Spatial locality with respect to code: Yes (the codes are executed in a contiguous section of the memory)
 Spatial locality with respect to array A: Yes (the inner loop traverses A via stride 1).
 Spatial locality with respect to array `sum`: Yes (the inner loop traverse `sum` via stride 1).

(e) [3 POINTS] The functions `sumRows` and `sumRows2` accomplish the same task. Which one is more cache friendly?

`sumRows2` has more locality, hence more cache friendly.

Question 5B. Locality [15 POINTS]

Consider the following C function that calculates the sum of the columns of a 2D matrix and assume that the elements of the array `sum` have been initialized to zero.

```
void sumCols(int nCols, int nRows, int sum[nCols], int A[nRows][nCols]) {
    for (int j = 0; j < nCols; j++) {
        for (int i = 0; i < nRows; i++) {
            sum[j] = sum[j] + A[i][j];
        }
    }
}
```

(a) [3 POINTS] Does the function `sumCols` exhibit *temporal locality*?

Temporal locality with respect to code: **Yes. `sum[j]` computed and used many times.**
 Temporal locality with respect to array A: **No (not a single element `A[i][j]` accessed more than once.)**
 Temporal locality with respect to array `sum`: **Yes (`sum[j]` accessed frequently inside the inner loop).**

(b) [3 POINTS] Does the function `sumCols` exhibit *spatial locality*?

Spatial locality with respect to code: **Yes (the codes are executed in a contiguous section of the memory)**
 Spatial locality with respect to array A: **No (The inner loop traverses the column elements).**
 Spatial locality with respect to array `sum`: **Yes (Inner loop traverse `sum` with stride 1).**

Now consider the following alternative implementation `sumCols2` where the loops are permuted.

```
void sumCols2(int nCols, int nRows, int sum[nCols], int A[nRows][nCols]) {
    for (int i = 0; i < nRows; i++) {
        for (int j = 0; j < nCols; j++){
            sum[j] = sum[j] + A[i][j];
        }
    }
}
```

(c) [3 POINTS] Does the function `sumCols2` exhibit *temporal locality*?

Temporal locality with respect to code: **Yes. `sum[j]` computed and used many times.**
 Temporal locality with respect to array A: **No (not a single element `A[i][j]` accessed more than once.)**
 Temporal locality with respect to array `sum`: **Yes (`sum[j]` accessed frequently inside the inner loop).**

(d) [3 POINTS] Does the function `sumCols2` exhibit *spatial locality*?

Spatial locality with respect to code: **Yes (the codes are executed in a contiguous section of the memory)**
 Spatial locality with respect to array A: **Yes (the inner loop traverses A via stride 1).**
 Spatial locality with respect to array `sum`: **Yes (the inner loop traverse `sum` via stride 1).**

(e) [3 POINTS] The functions `sumCols` and `sumCols2` accomplish the same task. Which one is more cache friendly?

`sumCols2` has more locality, hence more cache friendly.

Question 6. Cache Memories [13 POINTS]

(a) [5 POINTS] Suppose that you are working on a system with the following specifications:

- The device is a 12-bit machine, i.e., the physical addresses are 12 bits wide.
- The memory is byte addressable, and memory accesses are to 1-byte words.
- The device contains a 4-way set associative cache (E = 4: Four lines per set)

Below table shows the current contents of the cache:

4-way Set Associate Cache (E = 4: Four lines per set)

Index	Valid	Tag	Byte 0	Byte 1	Valid	Tag	Byte 0	Byte 1	Valid	Tag	Byte 0	Byte 1	Valid	Tag	Byte 0	Byte 1
0	1	39	29	42	1	78	19	1A	0	7F	80	81	0	BC	36	34
1	0	FD	01	5F	1	D3	0F	10	1	3D	5E	5F	1	4F	A3	A6
2	1	A3	C4	54	1	BC	3A	3B	1	EC	4E	4C	1	D4	3E	4A
3	1	3A	CD	3E	0	0E	54	55	0	C3	33	12	1	57	6A	E2
4	0	22	69	65	1	D2	37	38	0	23	4C	45	1	22	7E	4C
5	0	CA	E4	39	0	CA	7E	7F	0	7B	C2	A2	1	CA	23	25
6	1	2C	1B	E4	0	3E	A9	AA	1	F9	D3	45	0	FD	FF	00
7	0	0F	40	DD	0	CD	E1	E2	0	D9	D4	F6	0	BB	63	64

For an access request to memory address $0xCAB$, fill in the following table with the corresponding values for the offset (O), index (I), tag (T), and whether a cache hit or miss occurred, and if it is a hit, what value will be returned?

	Value
Cache Offset (O)	$0x1$
Cache Index (I)	$0x5$
Cache Tag (T)	$0xCA$
Cache Hit? (Y/N)	Y
Cache Byte value	$0x25$

As a reporter working at MacRumors, you are given access to one of Apple's new iPhone prototypes. However, there is no information about the cache memories integrated to the device, and your goal is to perform some experiments to figure out some properties of its cache memory. Considering that the cache is empty (cold), here is the report of some address traces and whether it is a MISS or HIT.

Address Trace	Hit/Miss
$0x000$	MISS
$0x001$	HIT
$0x002$	HIT
$0x003$	HIT
$0x004$	HIT
$0x005$	HIT
$0x006$	HIT
$0x007$	HIT
$0x008$	MISS

(b) [2 POINTS] What is the block size of the cache (in bytes)?

8 bytes. The first one is a miss, but the request loads the next 7 sequential bytes as well.

- (c) [2 POINTS] The following table shows your observations from your second set of experiments, in which you incremented the addresses of the memory accesses with a certain pattern:

Address Trace	Hit/Miss
0x000	MISS
0x008	MISS
0x000	HIT
0x010	MISS
0x000	HIT
0x018	MISS
0x000	HIT
0x020	MISS
0x000	HIT
<i>(some more trials...)</i>	
0x100	MISS
0x000	MISS

What is the total size of the cache (in bytes)?

32*8 = 256 bytes. Only after 32 trials, 0x000 becomes a MISS.

- (d) [2 POINTS] The following table shows your observations from your third set of experiments:

Address Trace	Hit/Miss
0x000	MISS
0x100	MISS
0x000	HIT
0x200	MISS
0x000	HIT
0x300	MISS
0x000	HIT
0x400	MISS
0x000	MISS

What is the associativity of the cache, i.e., the number of lines per set?

4. If we increment the accesses by the size of the cache, i.e., 0x100, the access to 0x000 becomes a miss after the 4th increment.

- (e) [2 POINTS] What is the number of bits required to encode the cache index?

3 bits. There are $256 / (4 * 8) = 8$ sets. Hence 3 bits is enough.

Question 7A. Linking

- (a) [7 POINTS] The four stages of compiling a C program involves *preprocessing*, *compiling*, *assembly*, and *linking*. Select all the steps involved during the given events.

Compilation				Event
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Generation of assembly code
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Generation of relocation entries
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Generation of object code
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Insertion of code(s) from the contents of other files
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Resolution of relocation entries
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Figuring out which variables will go into the .rodata, .data, and .bss sections
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Merging separate code and data sections into single sections.

- (b) [3 POINTS] Consider the following executable main, generated with compiling and linking two source files, `main.c` and `foo.c`, using the command `gcc -o main main.c foo.c`

Considering the contents of `main.c` and `foo.c` given below,

```
/* main.c */
#include <stdio.h>

static int x = 0;
int y = 1;
int z = 2;

void foo();

int main()
{
    int y = 4;
    foo();
    printf("x=%d, y=%d, z=%d\n", x, y, z);
    return 0;
}
```

```
/* foo.c */
int x, y, z;

void foo()
{
    x = 9;
    y = 8;
    z = 7;
}
```

what will be the output of `main` when executed?

x=0, y=4, z=7

Question 7B. Linking [10 POINTS]

- (a) [7 POINTS] The four stages of compiling a C program involves *preprocessing*, *compiling*, *assembly*, and *linking*. Select all the steps involved during the given events.

Compilation				Event
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Generation of assembly code
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Generation of relocation entries
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Generation of object code
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Insertion of code(s) from the contents of other files
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Resolution of relocation entries
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Figuring out which variables will go into the <i>.rodata</i> , <i>.data</i> , and <i>.bss</i> sections
<i>preprocessing</i>	<i>compiling</i>	<i>assembly</i>	<i>linking</i>	Merging separate code and data sections into single sections.

- (b) [3 POINTS] Consider the following executable main, generated with compiling and linking two source files, *main.c* and *foo.c*, using the command `gcc -o main main.c foo.c`

Considering the contents of *main.c* and *foo.c* given below,

```
/* main.c */
#include <stdio.h>

int x = 0;
int y = 1;
static int z = 2;

void foo();

int main()
{
    int y = 6;
    foo();
    printf("x=%d, y=%d, z=%d\n", x, y, z);
    return 0;
}
```

```
/* foo.c */
int x, y, z;

void foo()
{
    x = 9;
    y = 8;
    z = 7;
}
```

what will be the output of main when executed?

x=9, y=6, z=2

Question 8. Heap Allocators [10 POINTS]

Consider the following memory allocation operations, which are performed on the following heap with an implicit free-list allocator. Also note that the header is of 8 bytes and the allocated chunks of memory should be multiples of 8 bytes. For the sake of simplicity, the memory addresses are given in decimal numbers.

```
void *a = malloc(16);
void *b = malloc(20);
void *c = malloc(12);
void *d = malloc(12);
void *e = malloc(20);
free(b);
free(d);
```

16 Used	a	a	24 Free				16 Used	c	c +pad	16 Free			24 Used	e	e	e +pad
64	72	80	88	96	104	112	120	128	136	144	152	160	168	172	180	188

(a) [2 POINTS] Does the heap currently have *internal* fragmentation?

Yes. There is padding added to the memory chunk allocated for c and e.

(b) [2 POINTS] Does the heap currently have *external* fragmentation?

Yes. There are non-consecutive free blocks.

(c) [2 POINTS] Suppose that the next call to the heap allocator is `malloc(4)`. Specify the address that the `malloc` function returns if a first-fit approach is employed? What if the best-fit approach is used?

Using first-fit: 96

Using best-fit: 152

(d) [2 POINTS] Considering the layout of the heap shown at the top, what is the maximum size one can ask for the heap allocator for `realloc(a, size)`?

24. The implicit free-list allocator does not perform in-place `realloc` and always move the data to a different location. The maximum such chunk is of 24 bytes.

(e) [2 POINTS] Suppose that you are considering an alternative heap allocator implementation, which uses an explicit free-list allocator. Does this change increase the throughput for allocation requests or not?

Yes. Because it is linear in the number of only free chunks.