



**KOÇ  
UNIVERSITY**

# **Database Management Systems**

## **Concurrency Control**

**M. Emre Gürsoy**

Assistant Professor  
Department of Computer Engineering

[www.memregursoy.com](http://www.memregursoy.com)



# Introduction

- We have been talking about:
  - "Good" schedules versus "bad" schedules (e.g., conflict serializability)
  - How to check if a schedule is "good" or "bad"
- But our previous methods have two shortcomings:
  - They require us to know all transactions and all actions ahead of time (all **R**s and **W**s)
  - They don't tell us how to **create** a "good" schedule
- Now we'll learn about **locking** for **concurrency control**

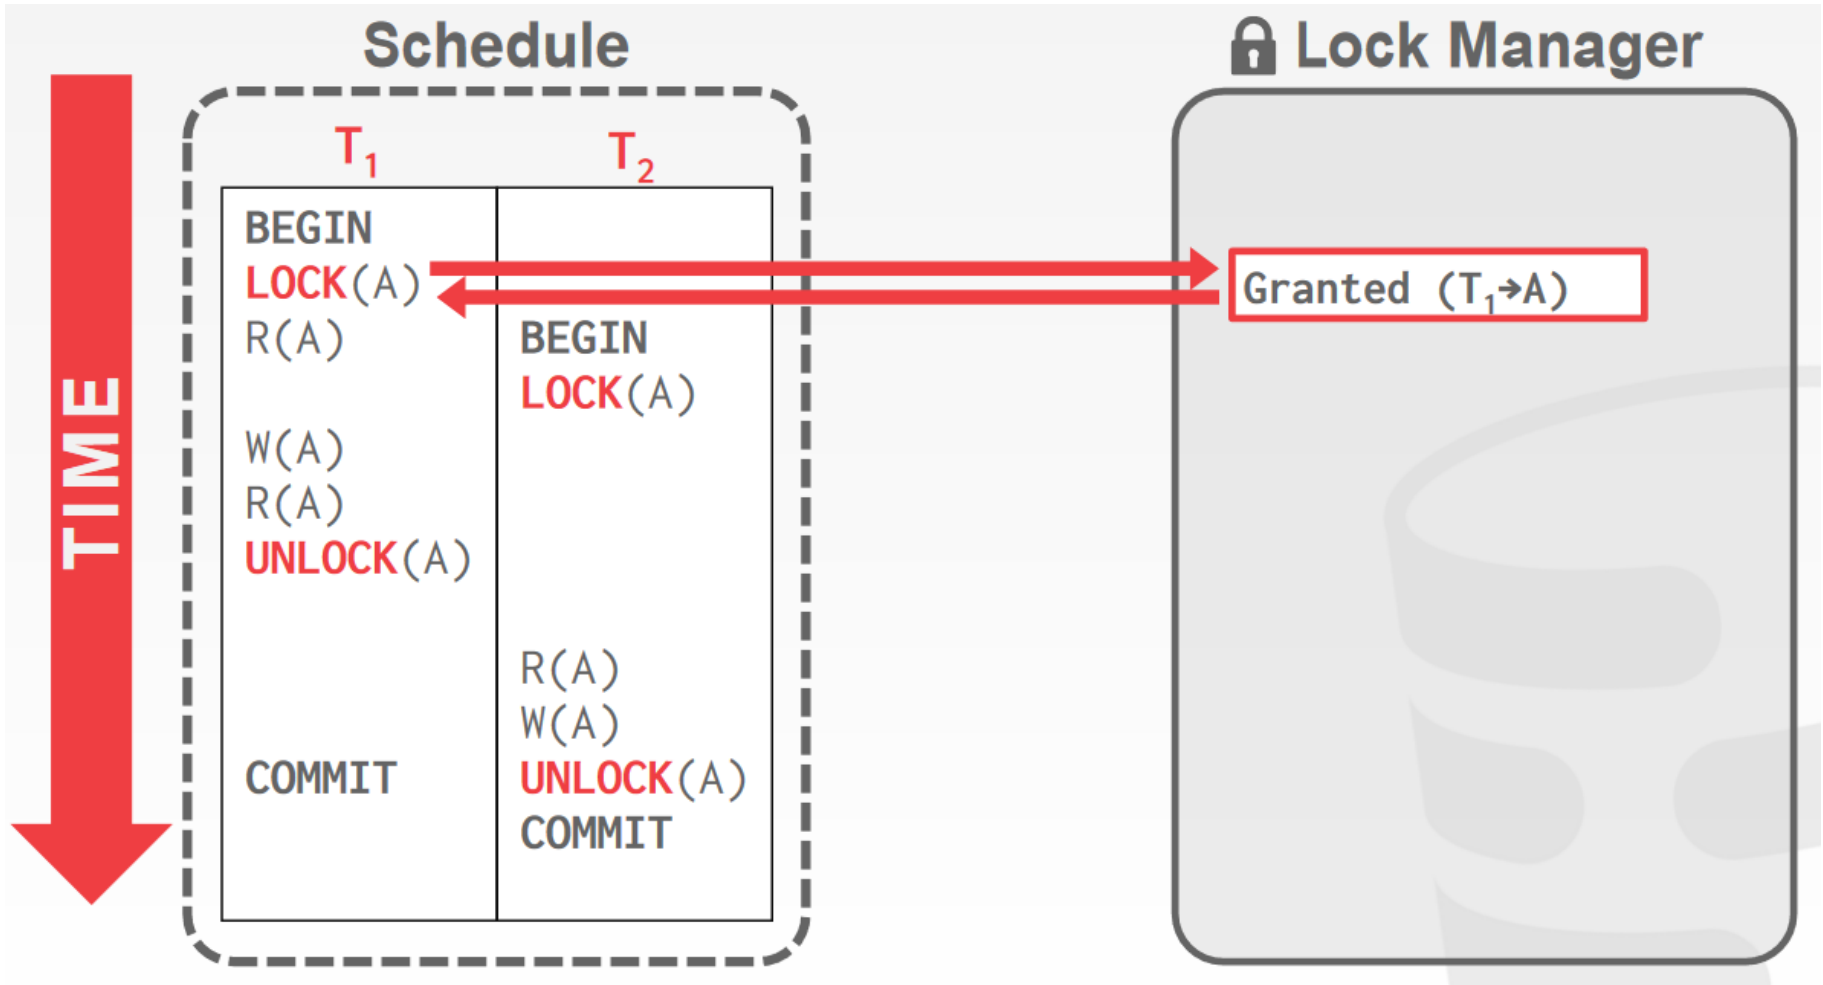


# Intuition

- Execution using locks:
  - Transactions **request** locks.
  - Lock manager **grants** or **denies** requests.
  - Once transactions are done with the DB object, they **release** their locks (**unlock**).
  - **Lock manager** keeps track of:
    - which transactions hold which locks on which objects
    - which transactions are waiting to acquire locks

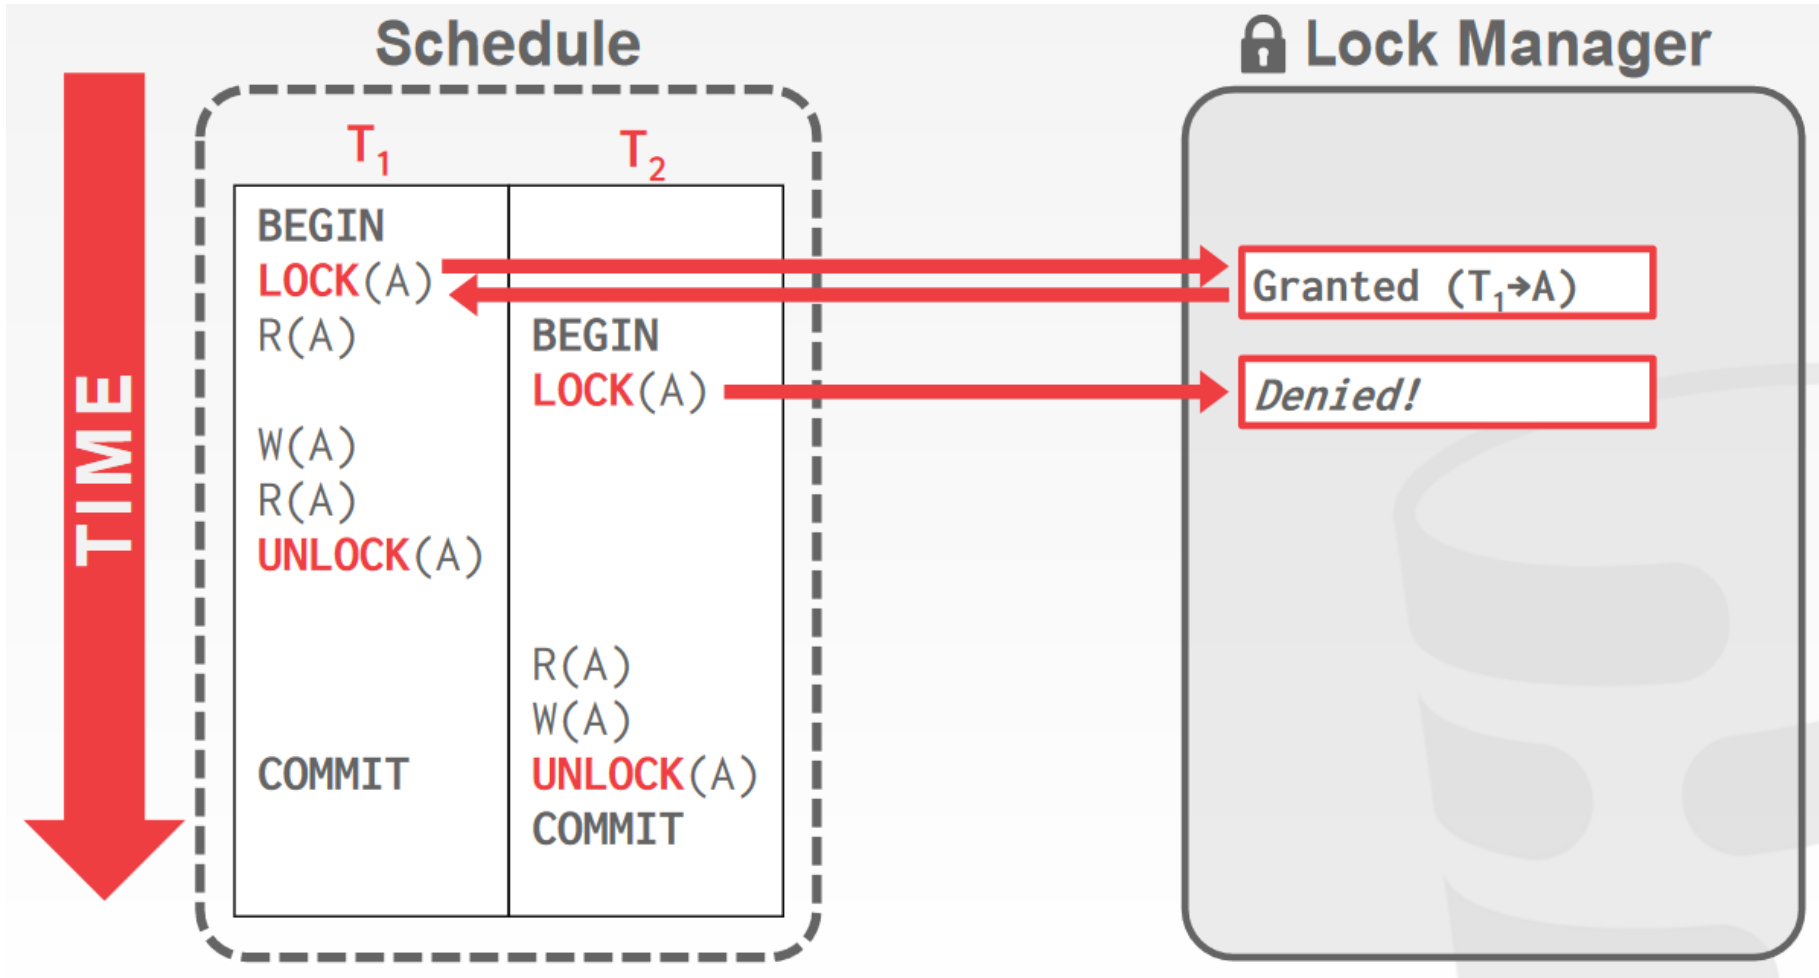


# Intuition



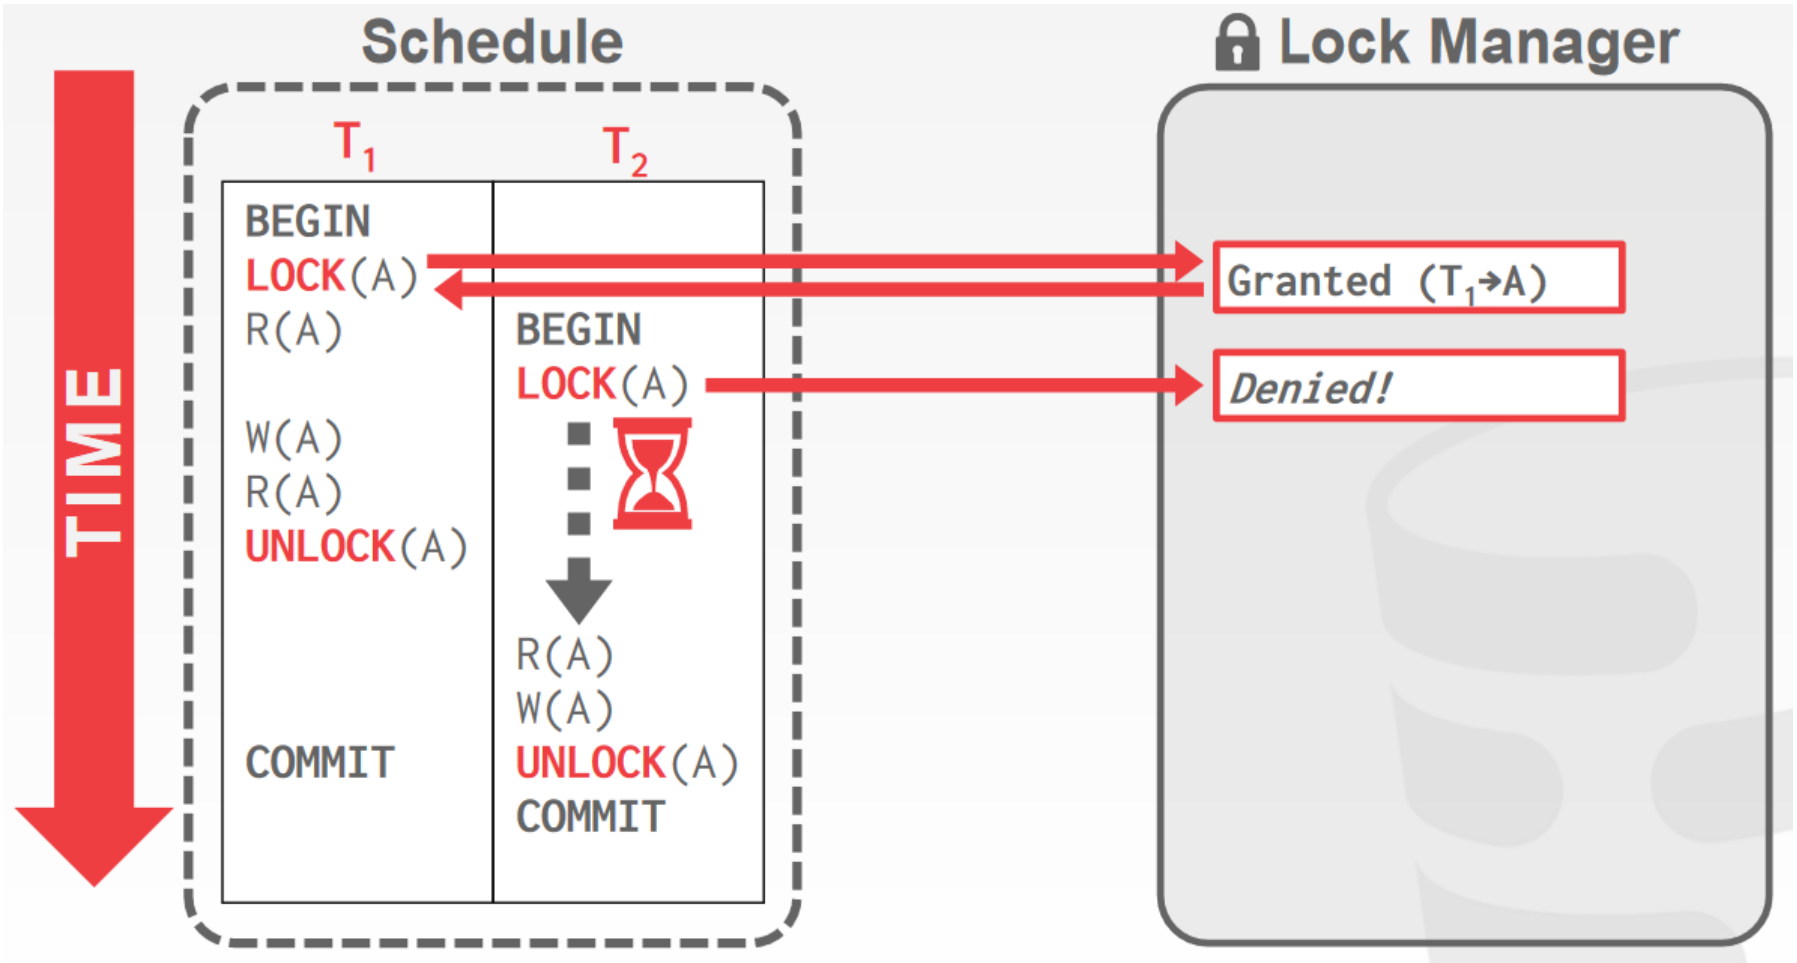


# Intuition



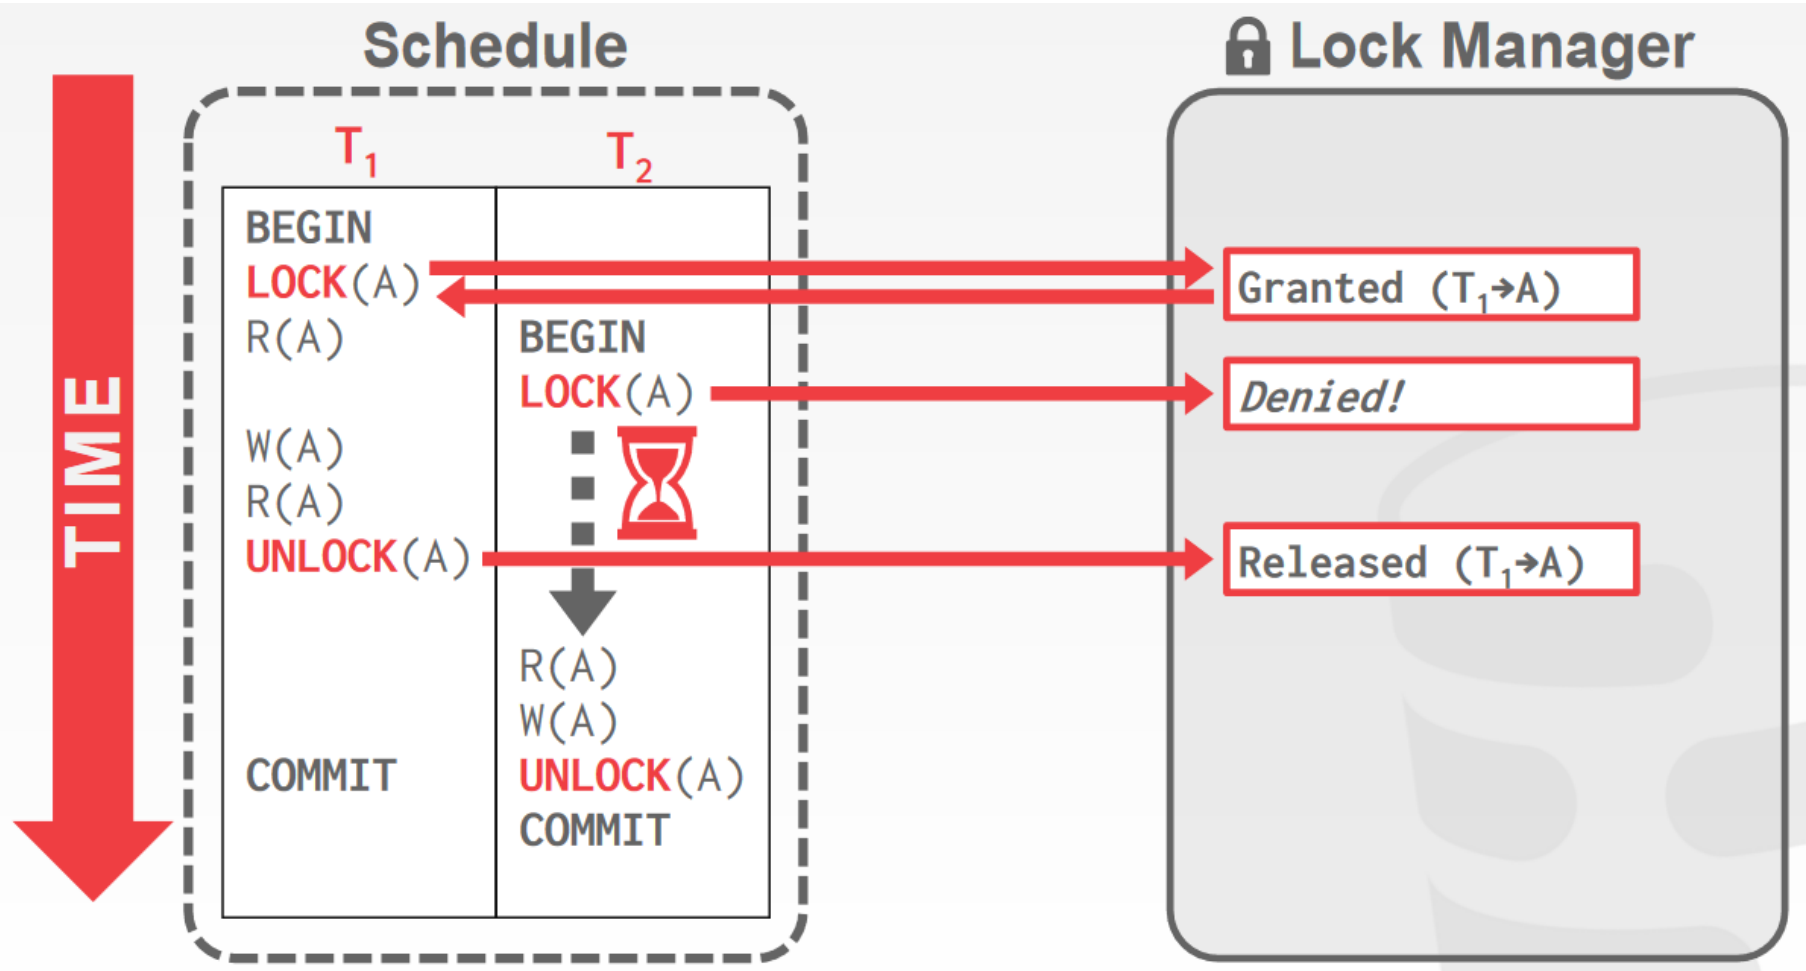


# Intuition



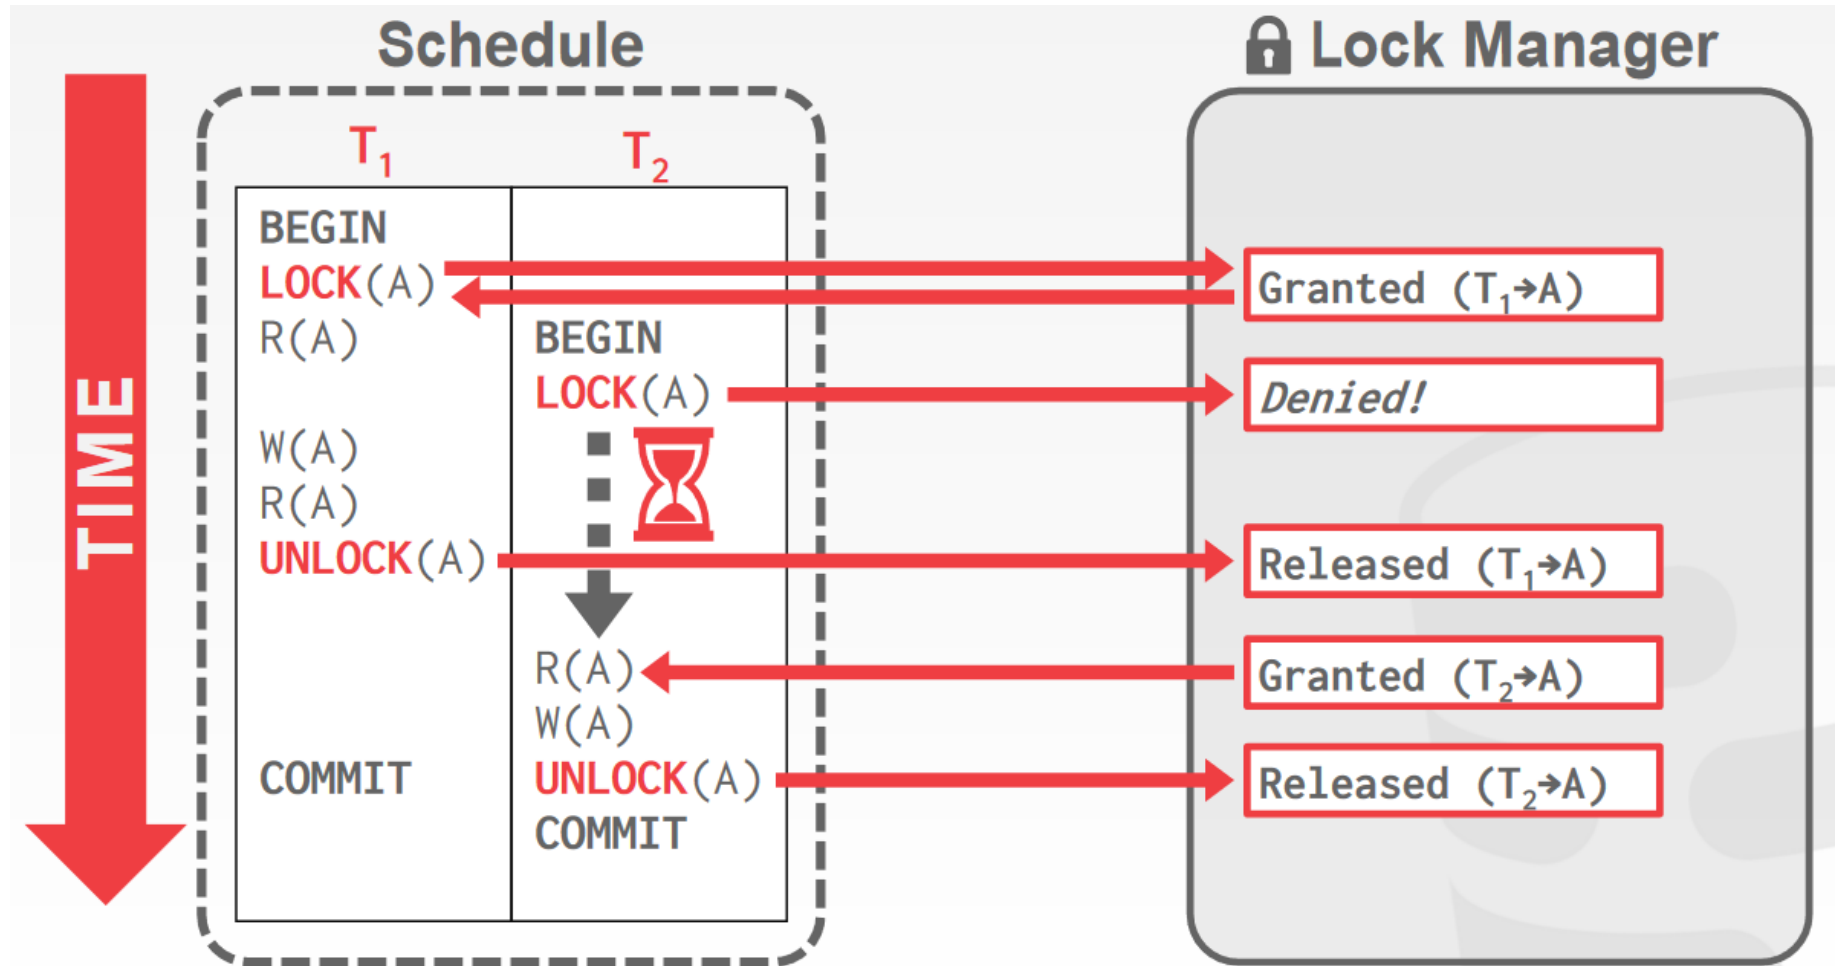


# Intuition





# Intuition







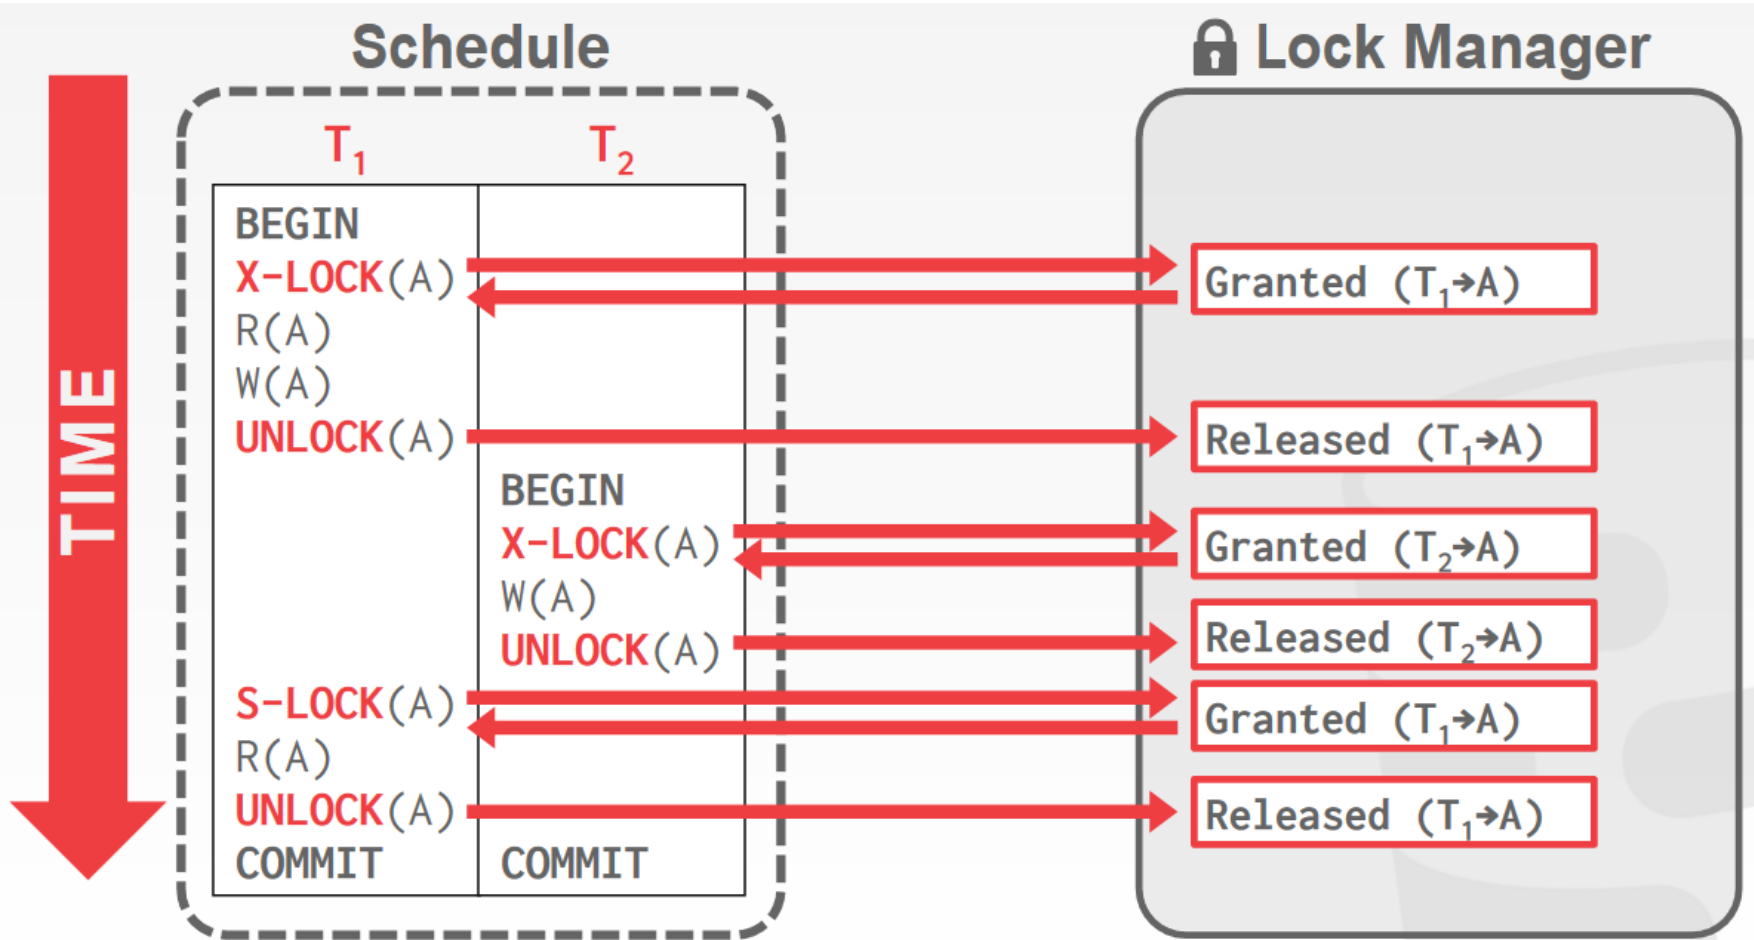
# Lock Types

- Two types of locks:
  - Shared locks (S-lock):** for **reads**
  - Exclusive locks (X-lock):** for **writes**
- Why not use shared locks for writes?
- Why not use exclusive locks for reads?

Compatibility Matrix		
	Shared	Exclusive
Shared	✓	X
Exclusive	X	X



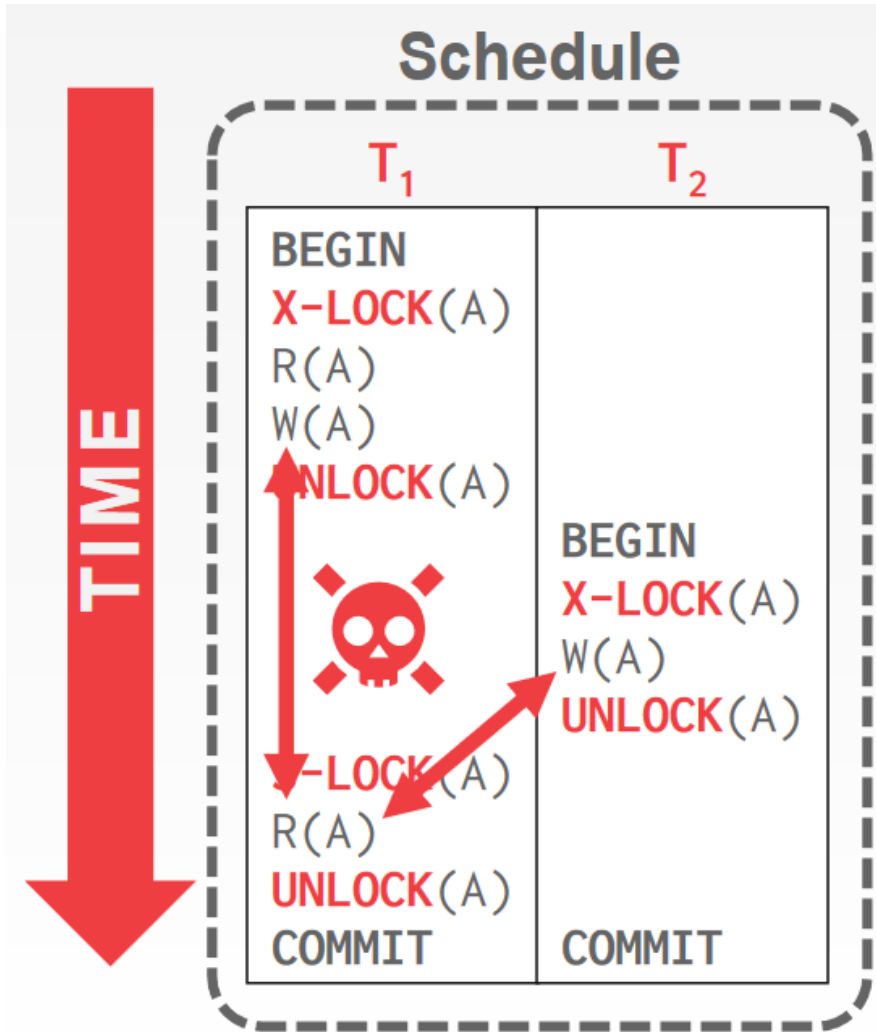
# Executing with Locks



Do you see a problem here?



# Executing with Locks



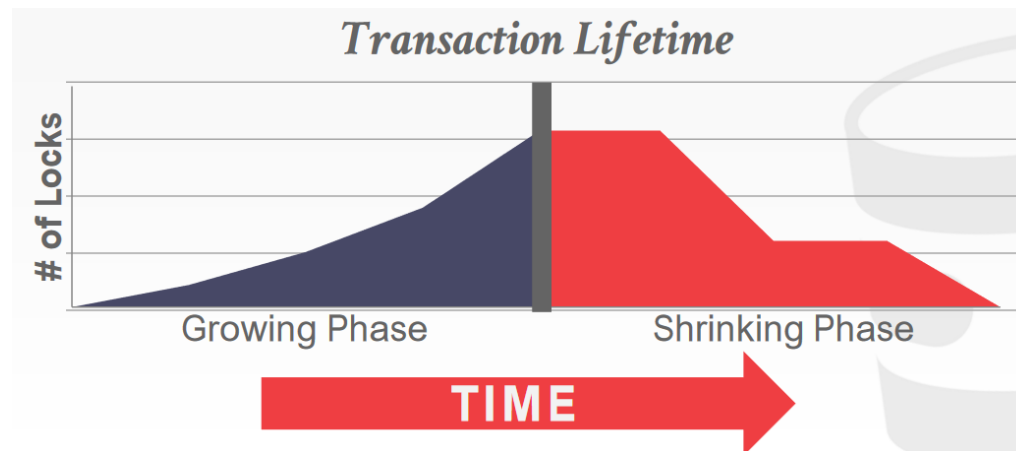
$T_1$  is expecting to read what it wrote but ends up reading what  $T_2$  wrote.

What if  $T_2$  decides to abort later?



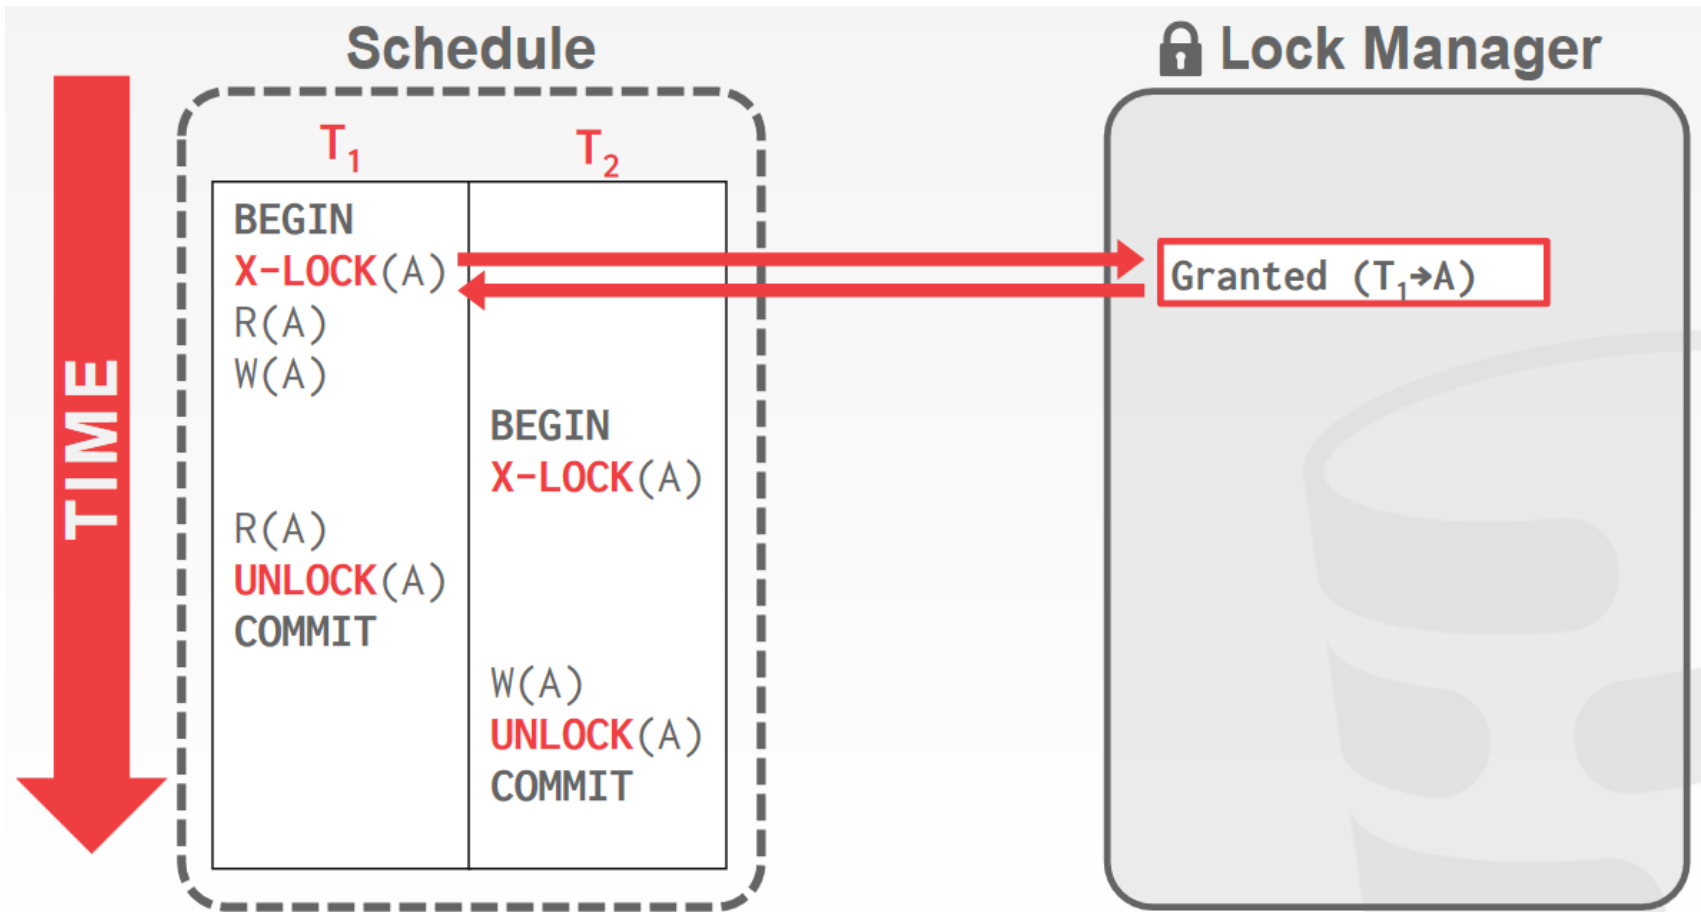
# Two-Phase Locking

- **Two-phase locking (2PL)** is a commonly used concurrency control protocol.
  - Phase #1: **Growing**
    - Each transaction requests locks from DBMS lock manager.
    - Lock manager grants or denies lock requests.
  - Phase #2: **Shrinking**
    - The transaction is allowed to release previously acquired locks, but it is not allowed to acquire new locks.





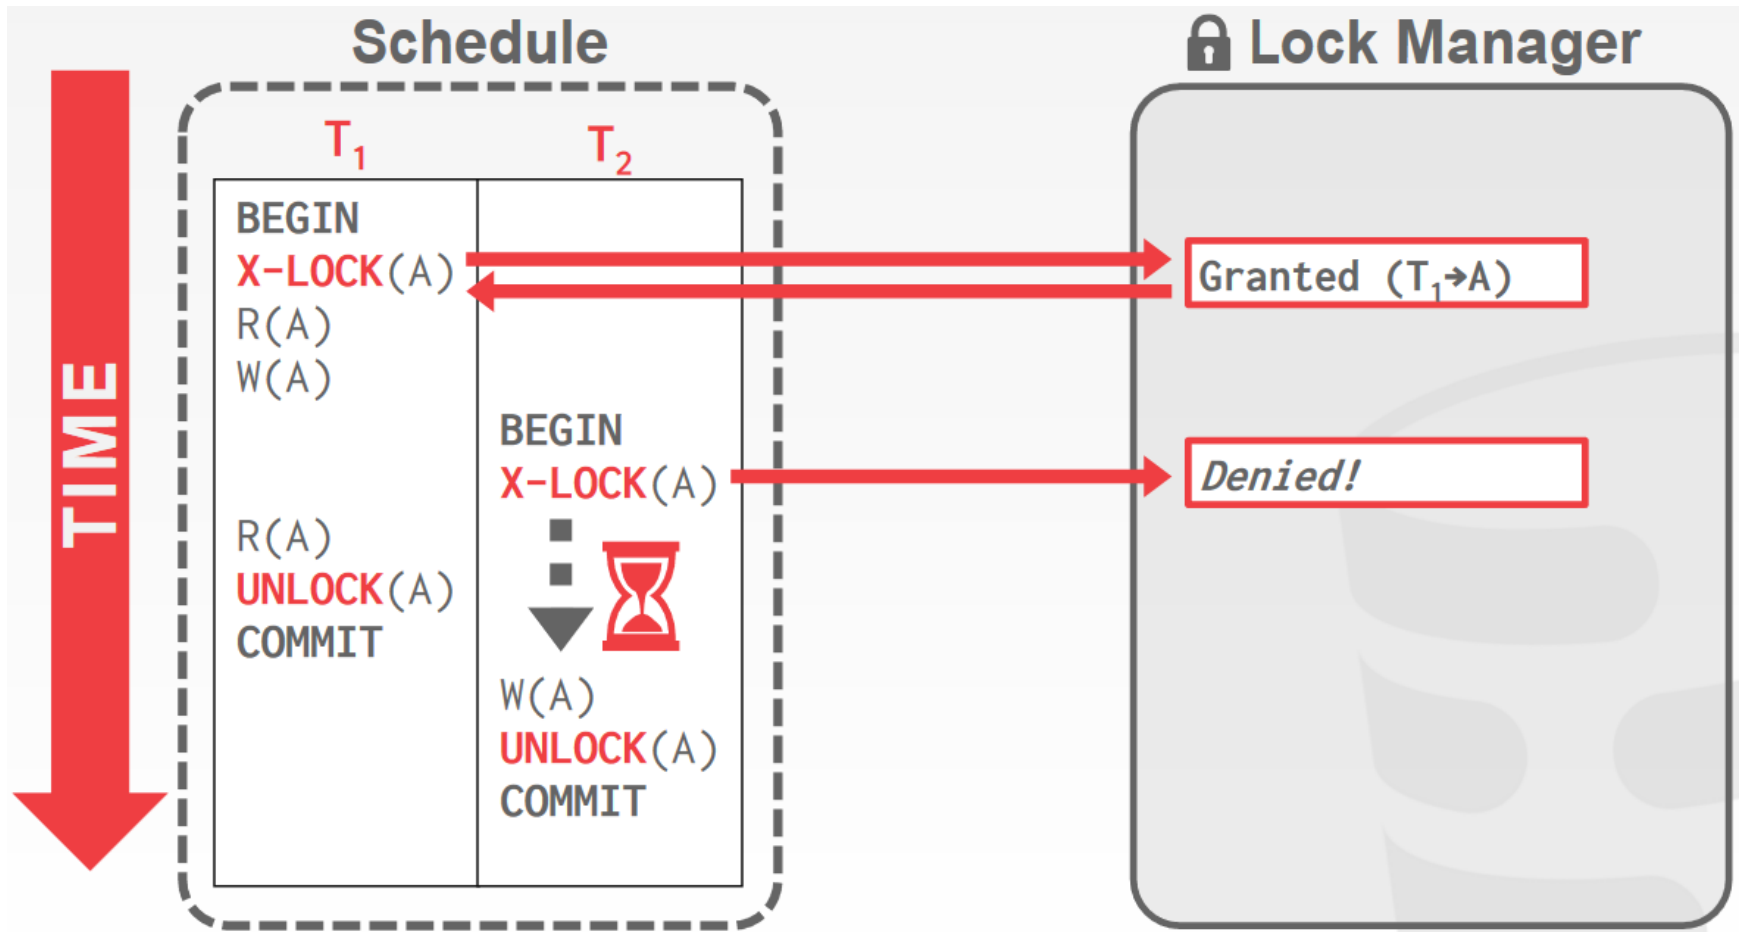
# Executing with 2PL



What is the main difference with the previous case?

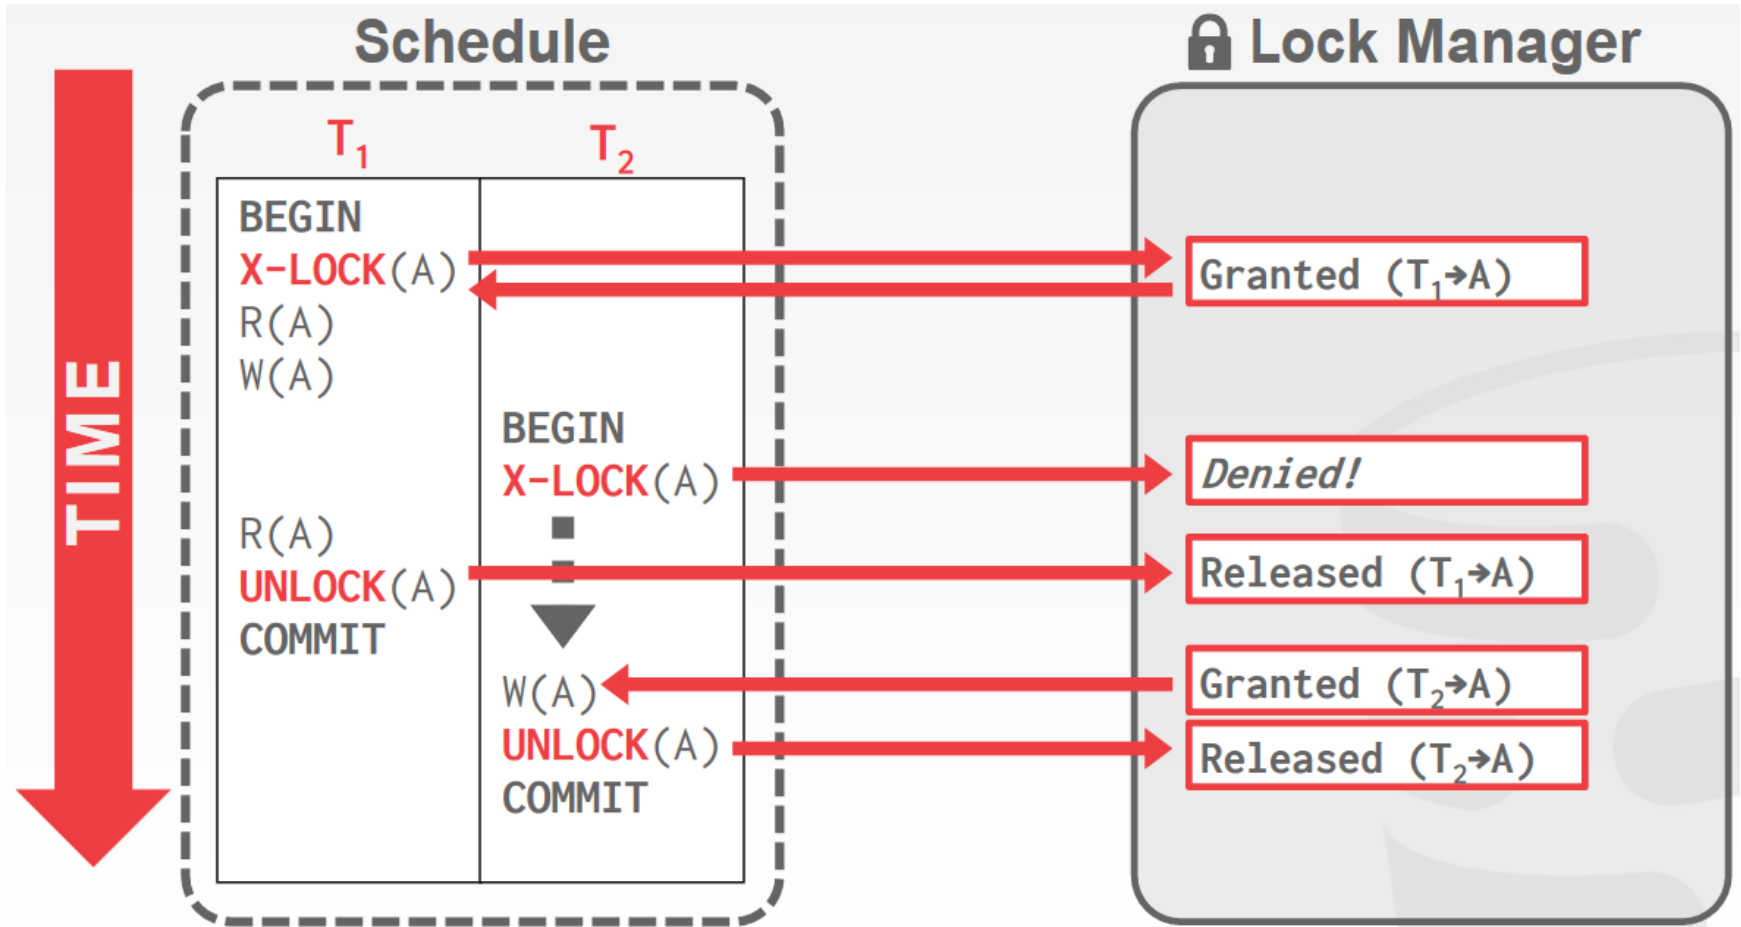


# Executing with 2PL





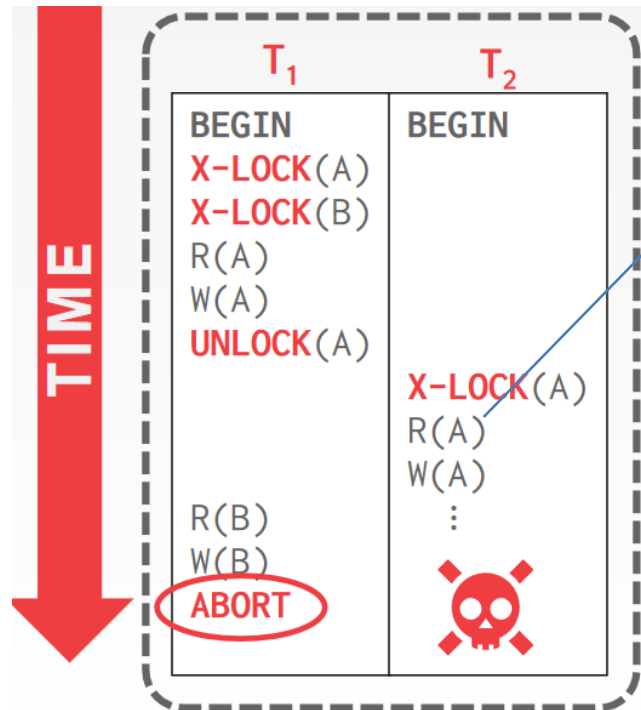
# Executing with 2PL





# Cascading Aborts

- 2PL achieves conflict serializability for transactions that commit, but it is subject to the **cascading aborts** problem.
  - When a transaction aborts, it causes other transactions to also have to abort.



But  $T_2$  had read what  $T_1$  wrote and assumed that it was true!


So, when  $T_1$  aborts we must also abort  $T_2$ .





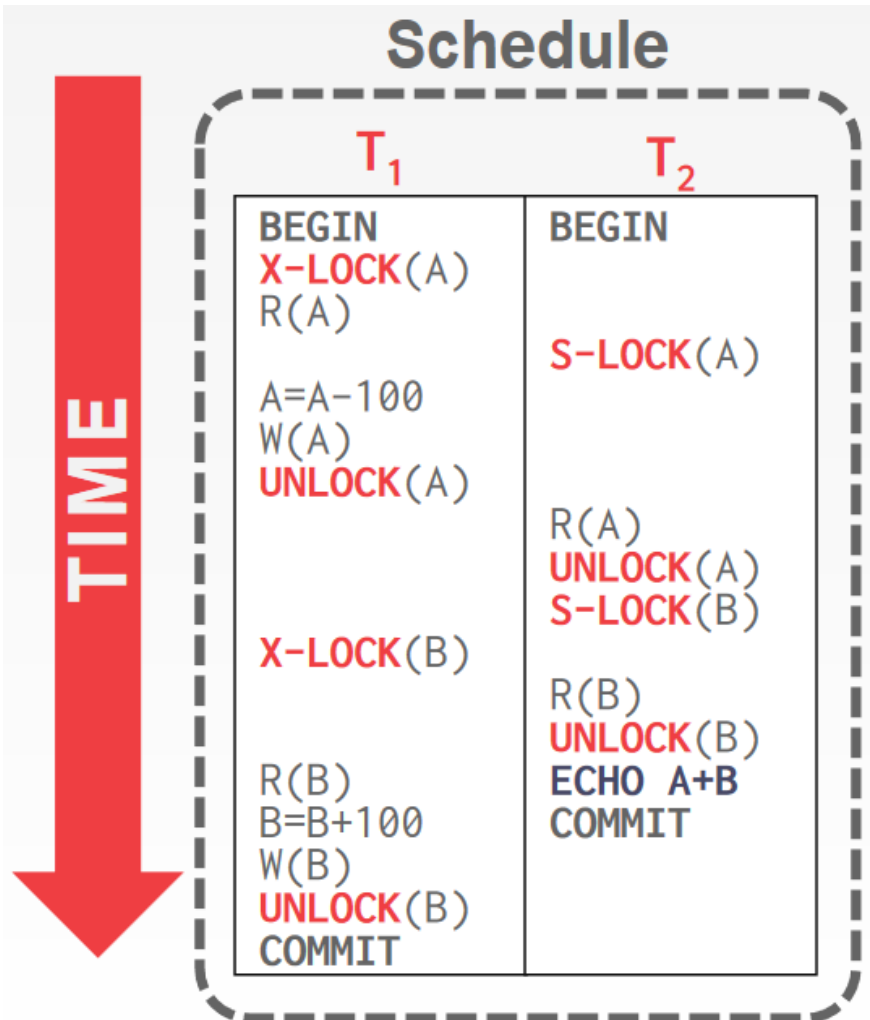
# Strict 2PL

- All locks held by a transaction are released when the transaction completes.
  - No **unlocking** in the middle of a transaction.
  - If I have an **X-lock** on **A**, no one else can get an **S-lock** or **X-lock** on **A** before I commit or abort.
  - Hence, no one can read/write on **A** before I am finished.
  - Hence, no cascading aborts.





# Example (Non-2PL)

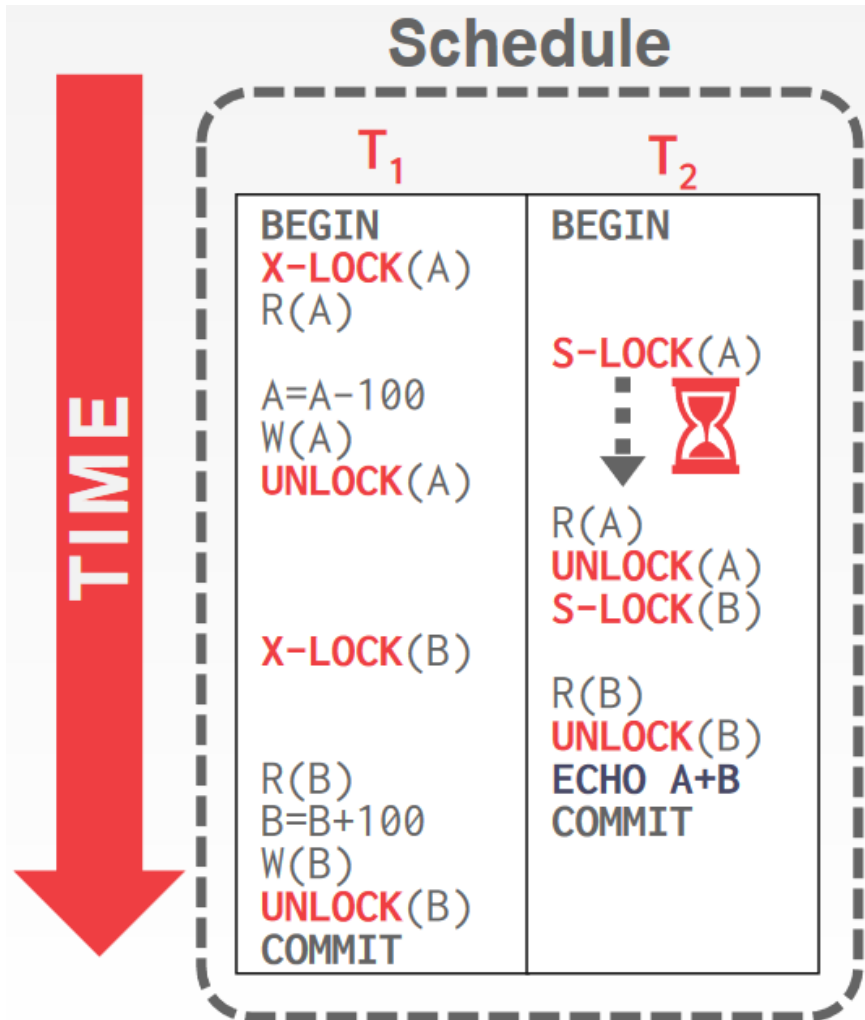


Assume  $A=1000$ ,  $B=1000$  initially.

What are the two transactions intending to do?

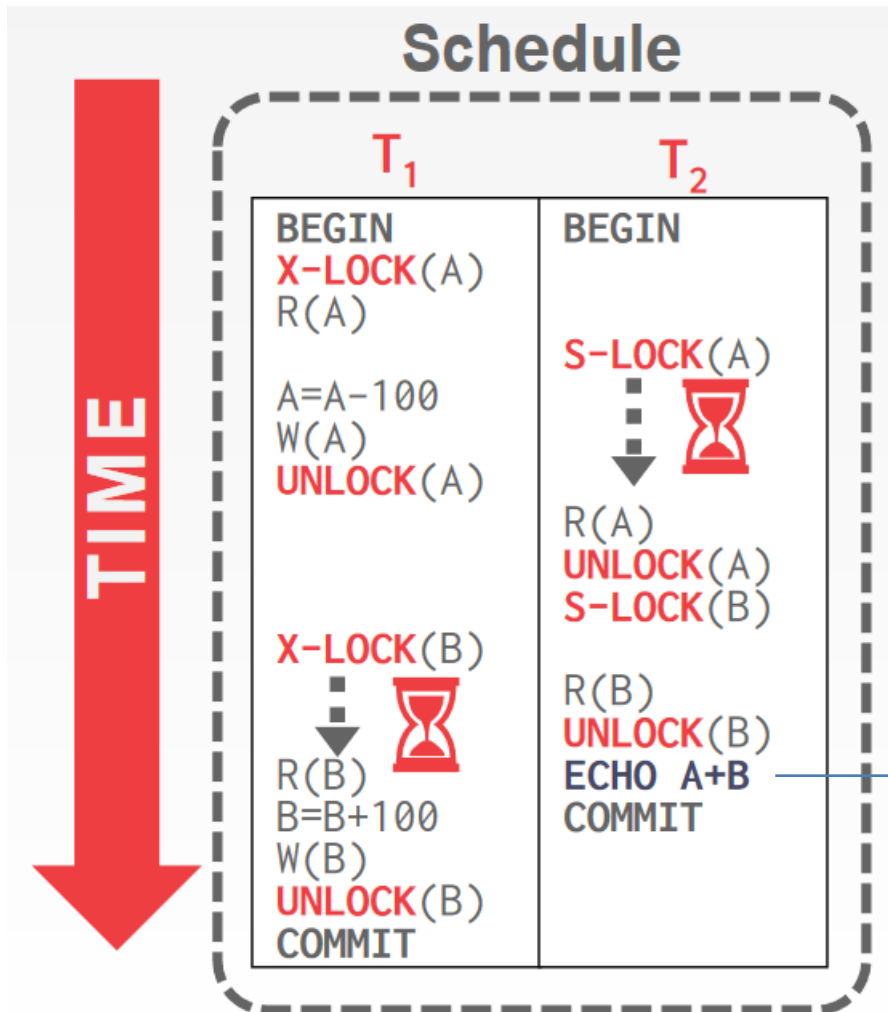


# Example (Non-2PL)





# Example (Non-2PL)



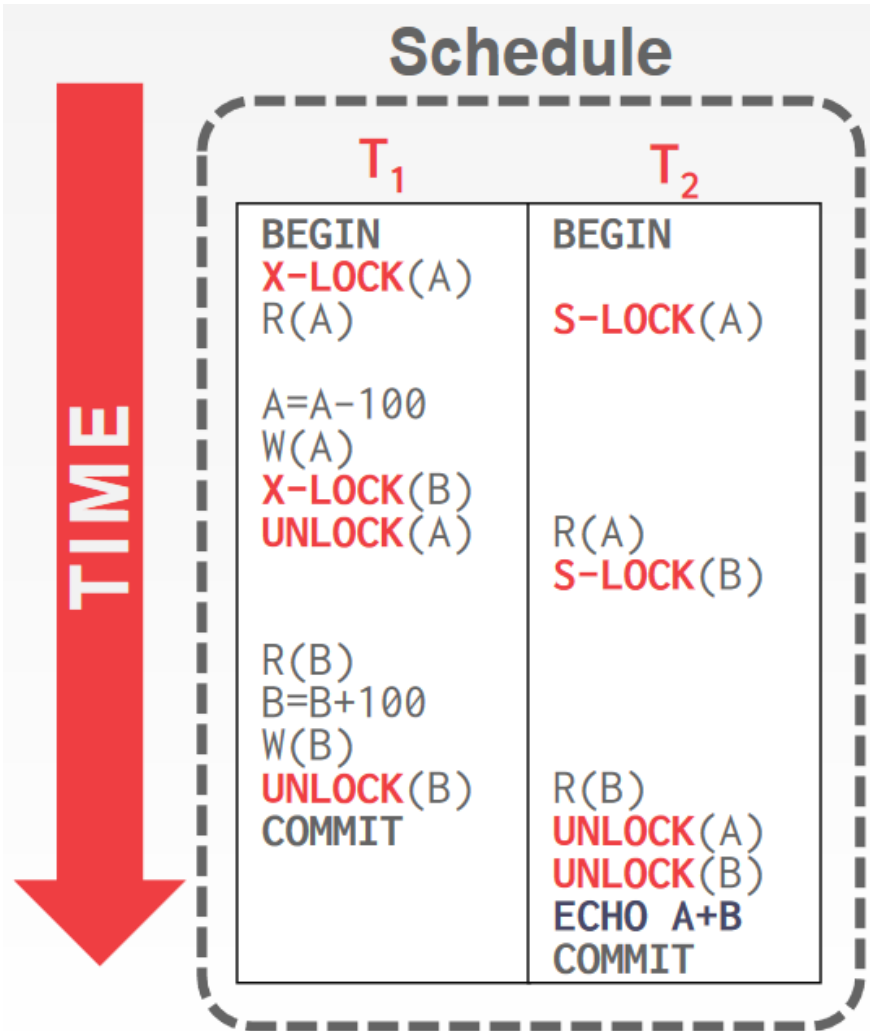
Assume  $A=1000$ ,  $B=1000$  initially.

What gets printed out?

What are the final values of  $A$  and  $B$ ?



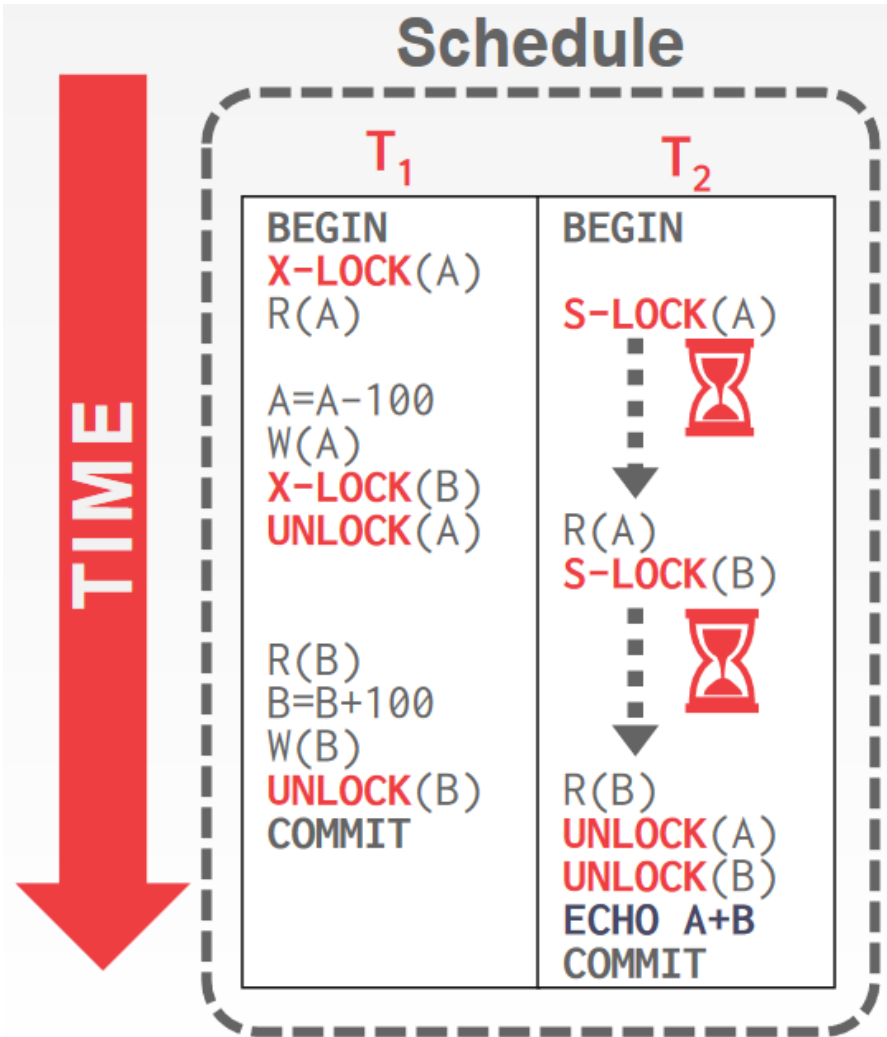
# Example (2PL)



Assume  $A=1000$ ,  $B=1000$  initially.




# Example (2PL)





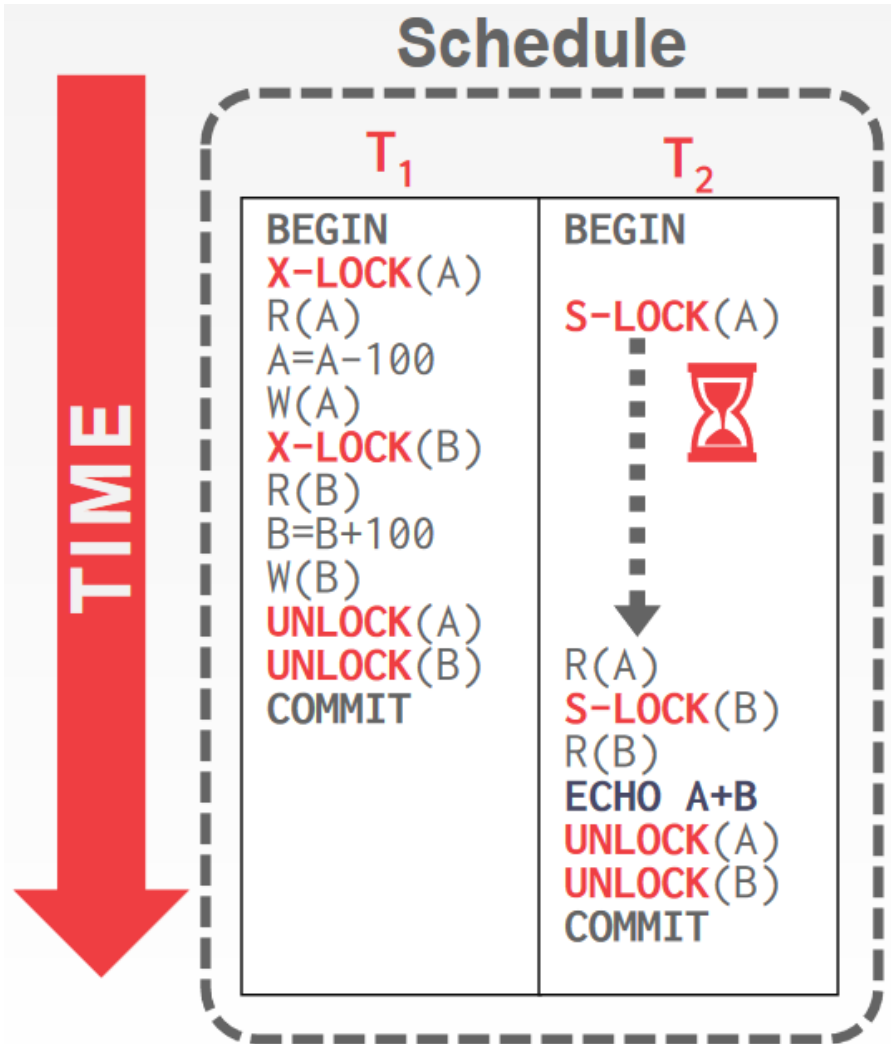
# Example (Strict 2PL)



Assume  $A=1000$ ,  $B=1000$  initially.



# Example (Strict 2PL)





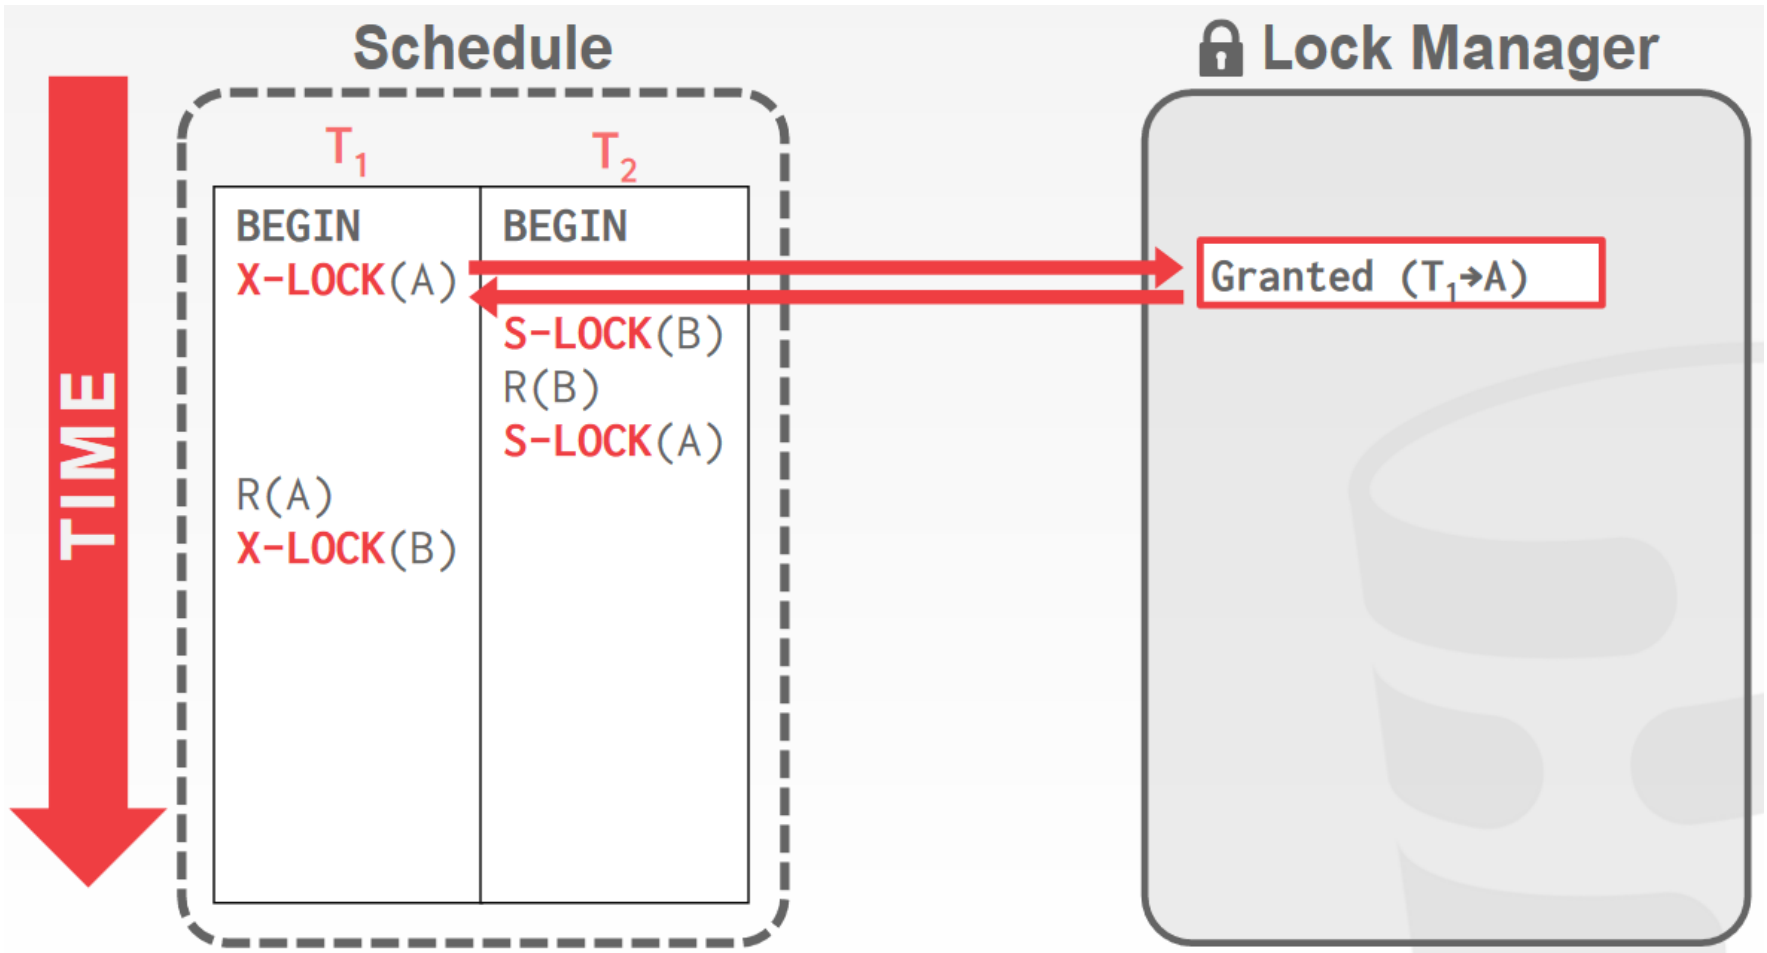


# Deadlocks

- 2PL can lead to deadlocks.
- A **deadlock** is a cycle of transactions waiting for locks to be released by each other.
  - $T_1$  locks **A**, needs access to **B** in order to continue.
  - $T_2$  locks **B**, needs access to **A** in order to continue.
- Two issues related to deadlocks:
  - Deadlock detection
  - Deadlock prevention

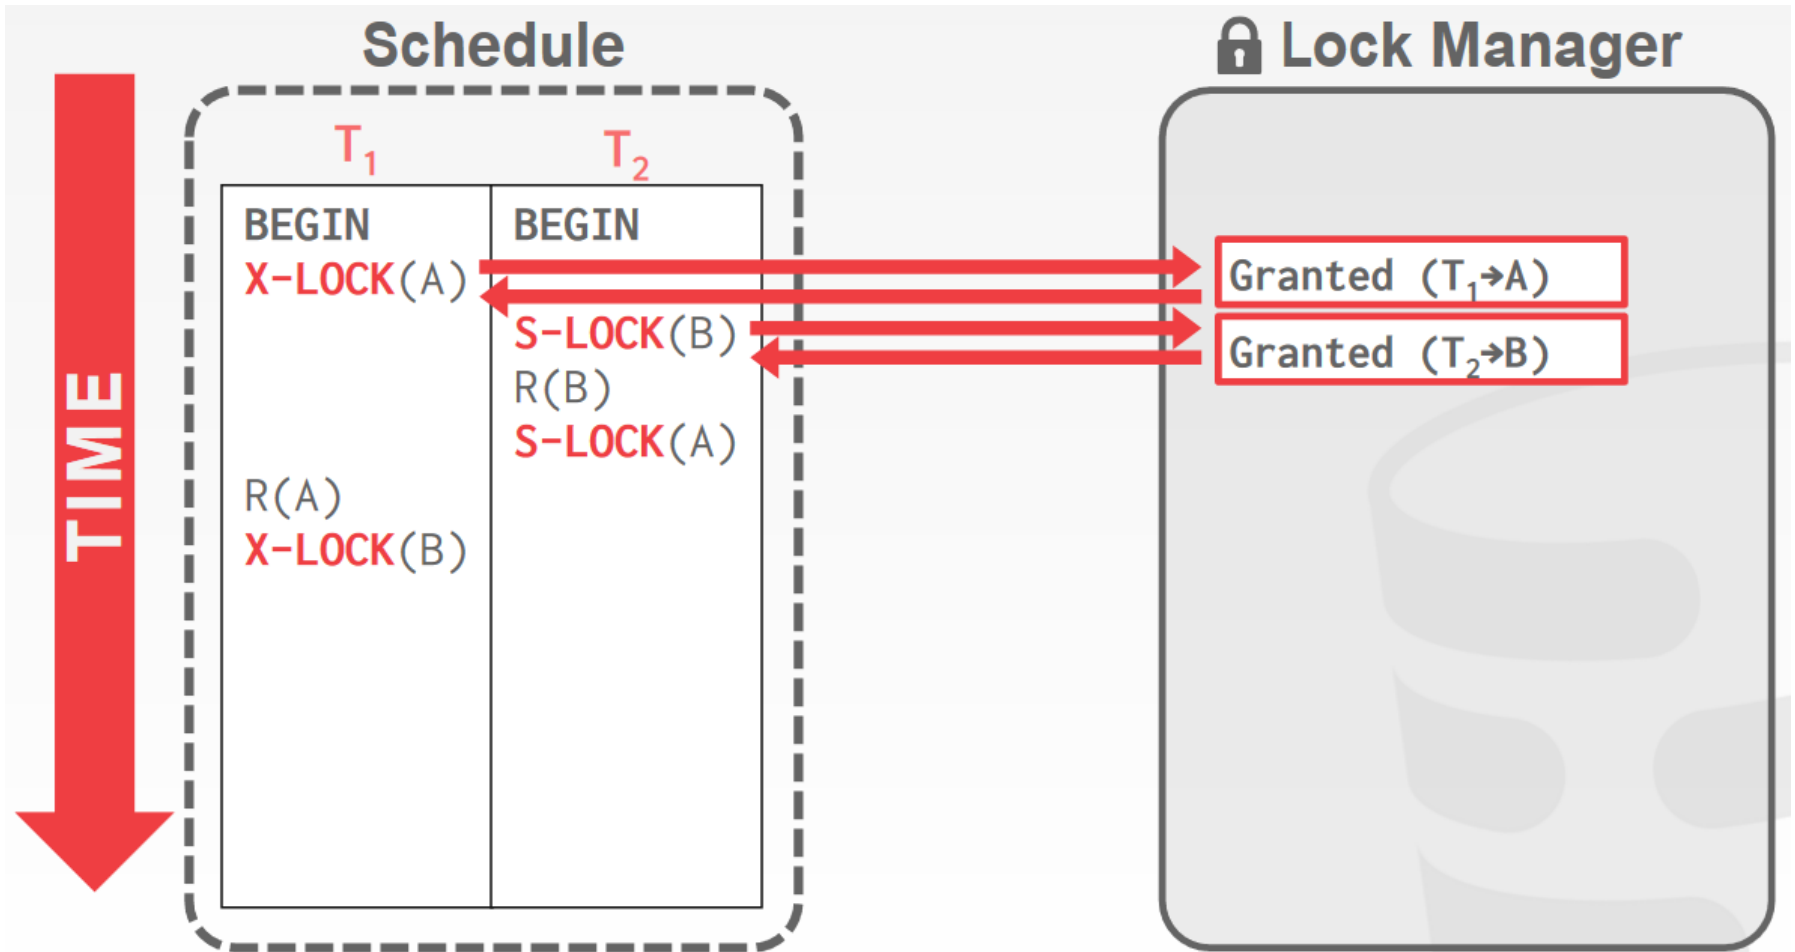


# Deadlocks



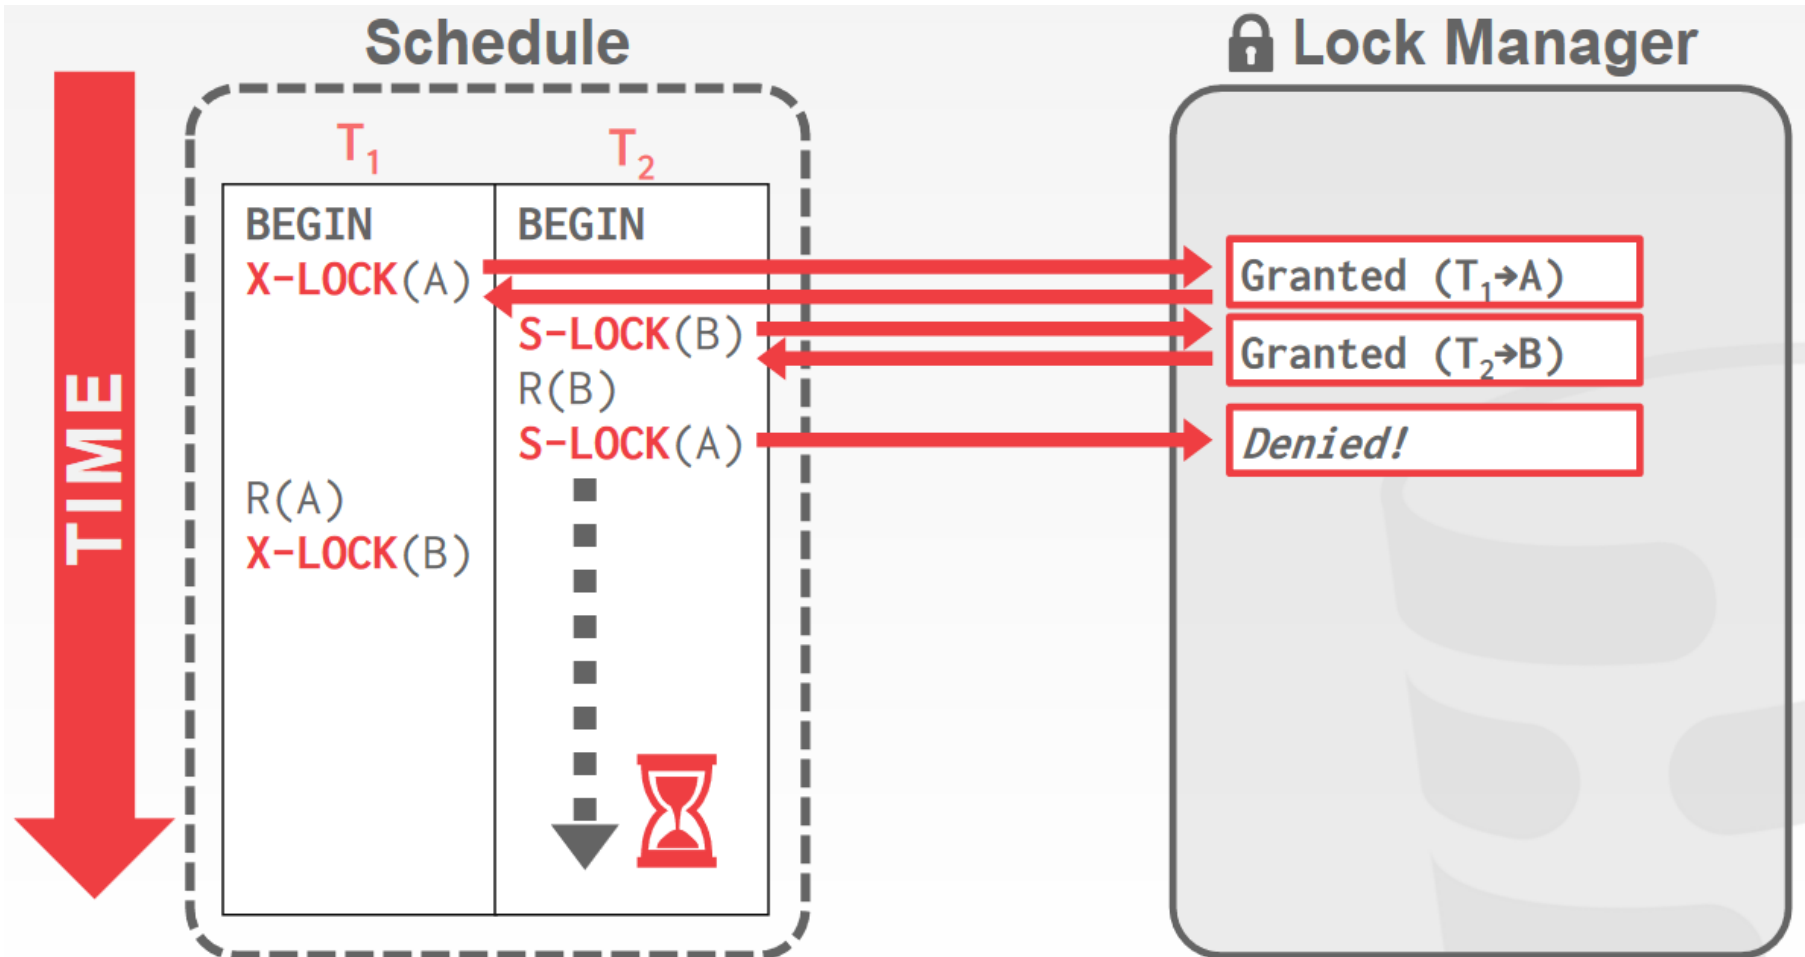


# Deadlocks



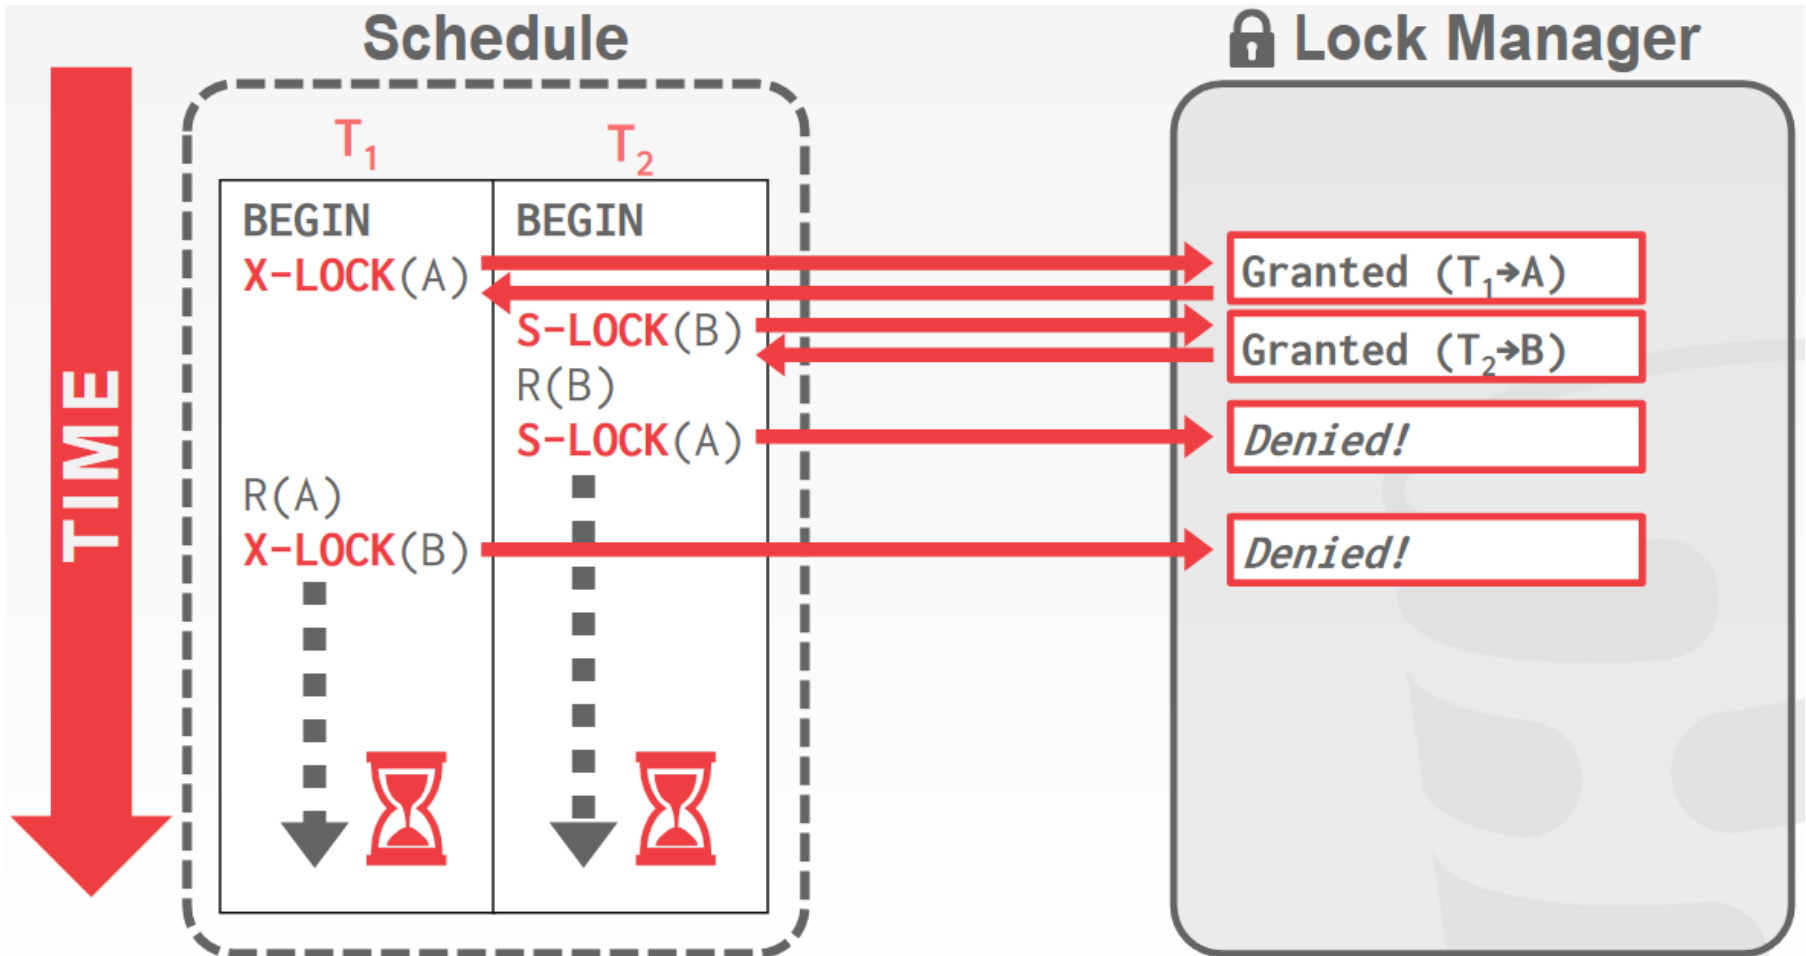


# Deadlocks






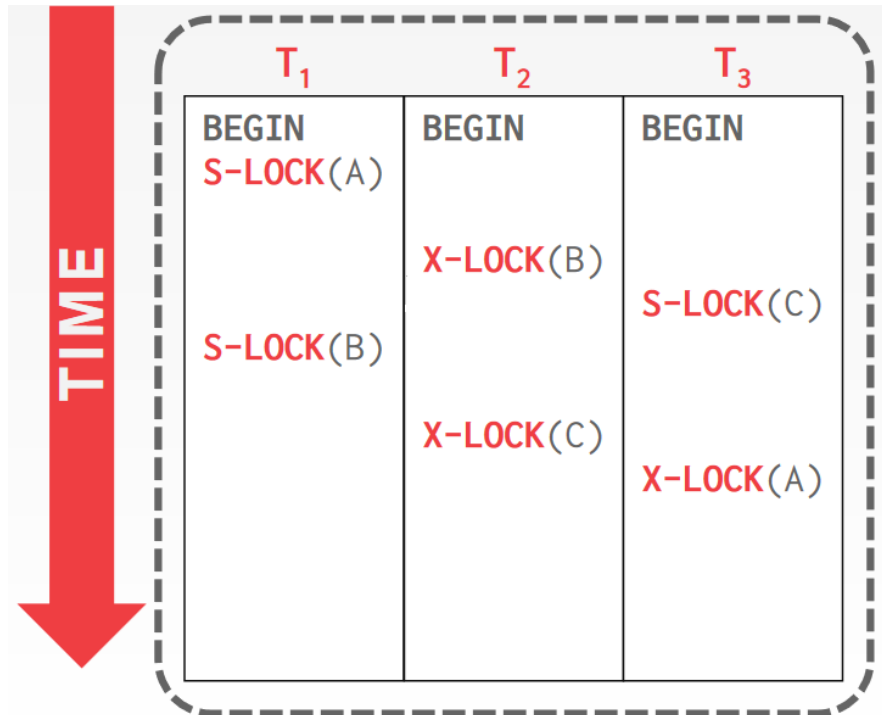
# Deadlocks





# Deadlock Detection

- We create a **waits-for graph**:
  - Each node is a transaction
  - Edge from  $T_i$  to  $T_j$  if  $T_i$  is waiting for  $T_j$  to release a lock
- Cycle in the graph means there's a deadlock





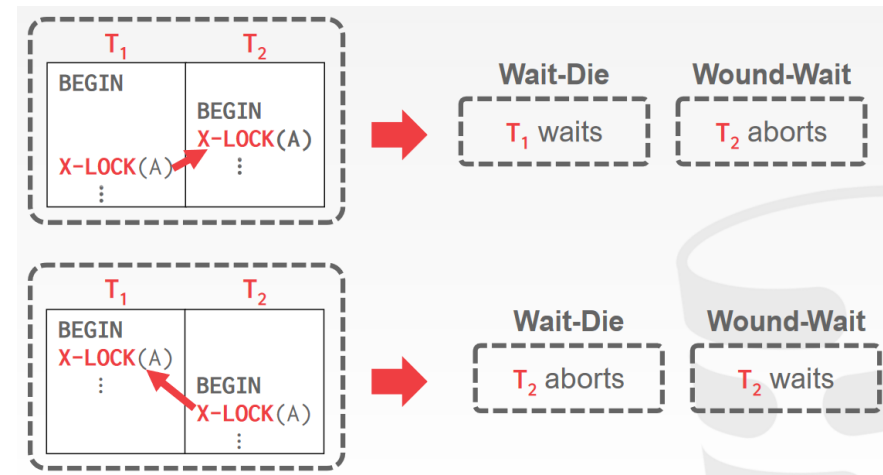
# Deadlock Detection

- When the DBMS detects a deadlock using the waits-for graph, it will select a **victim transaction** to break the cycle.
  - The victim is forced to restart or abort.
- How to select the proper victim? Possible heuristics:
  - By age (e.g., lowest timestamp)
  - By progress (e.g., fewest queries executed)
  - # of objects locked
  - ...
- We should keep in mind whether a transaction has been selected as a victim before, to avoid **starvation**.



# Deadlock Prevention

- When  $T_i$  tries to acquire a lock that is held by  $T_j$ , the DBMS "stops" one of them to prevent a deadlock.
  - Doesn't require a waits-for graph
- How to determine which transaction to stop?
  - Wait-Die**: If  $T_i$  started earlier than  $T_j$ , then  $T_i$  waits for  $T_j$ . Otherwise  $T_i$  aborts.
  - Wound-Wait**: If  $T_i$  started earlier than  $T_j$ , then  $T_j$  aborts and releases lock. Otherwise  $T_i$  waits.

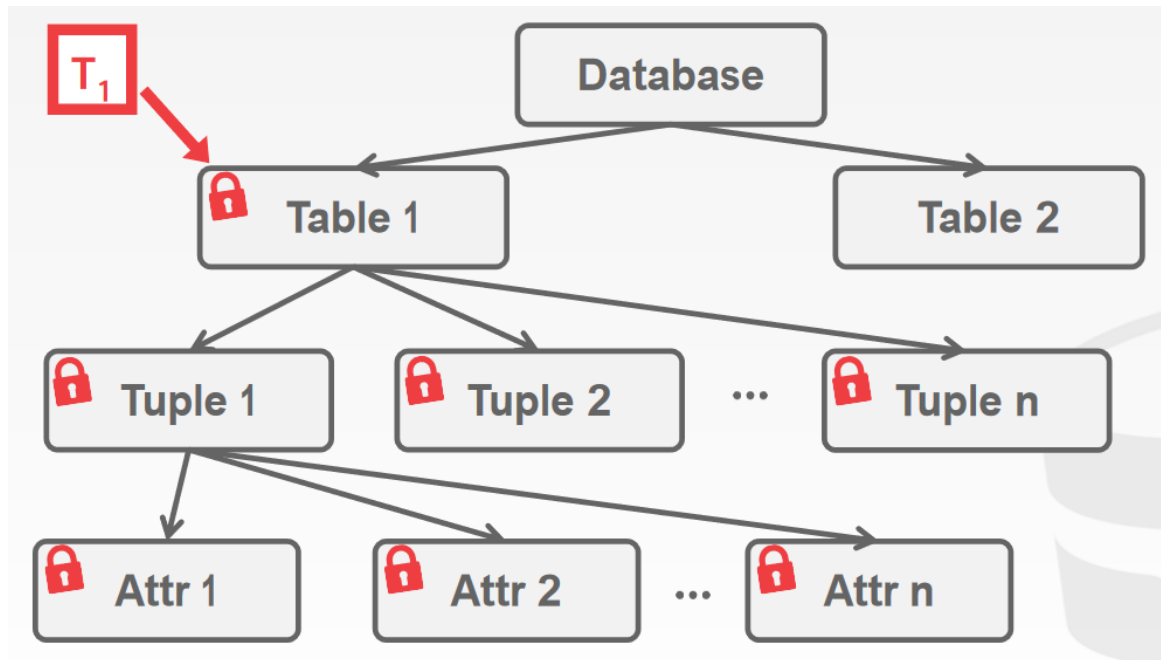






# Multi-Granularity Locking

- What are these **database objects** we have been locking?
  - Tables? Tuples? Attributes?
  - Locking too coarse: many conflicts, deadlocks
  - Locking too granular: can be inefficient






# Multi-Granularity Locking

- An **intention lock** allows a higher level node to be locked without having to check all descendant nodes.
  - If a node is intention locked, then explicit locking is being done at a lower level in the tree.
- Five types of locks:
  - **S** lock: shared lock (same as before)
  - **X** lock: exclusive lock (same as before)
  - **IS** lock: intention-shared lock
    - You get IS at the parent, then S at one or more descendants
  - **IX** lock: intention-exclusive lock
    - You get IX at the parent, then X or S at one or more descendants
  - **SIX** lock: S + IX together



# Example


Read Andy's record in **R**.





# Example


Read Andy's record in **R**.





# Example


Update Matt's record in R.





# Example

Update Matt's record in R.





# Compatibility


		$T_2$ Wants				
		IS	IX	S	SIX	X
$T_1$ Holds	IS	✓	✓	✓	✓	×
	IX	✓	✓	×	×	×
	S	✓	×	✓	×	×
	SIX	✓	×	×	×	×
	X	×	×	×	×	×

- If transactions' locks are not compatible, they cannot execute concurrently!



# Example

Scan **R** and update a few tuples. **T<sub>1</sub>**








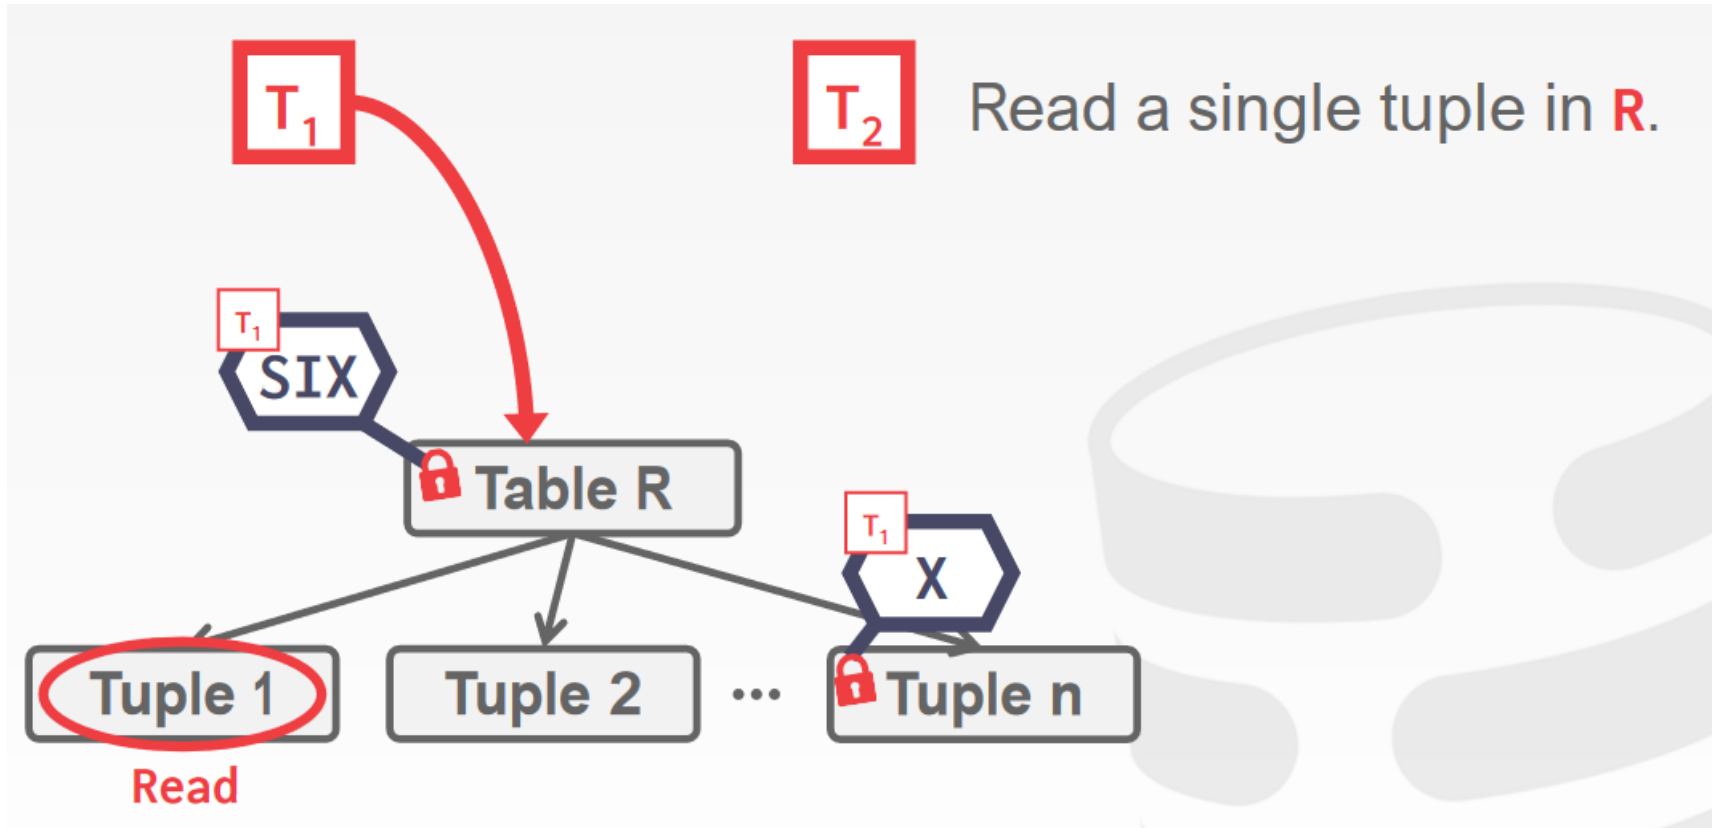
# Example

Scan **R** and update a few tuples.



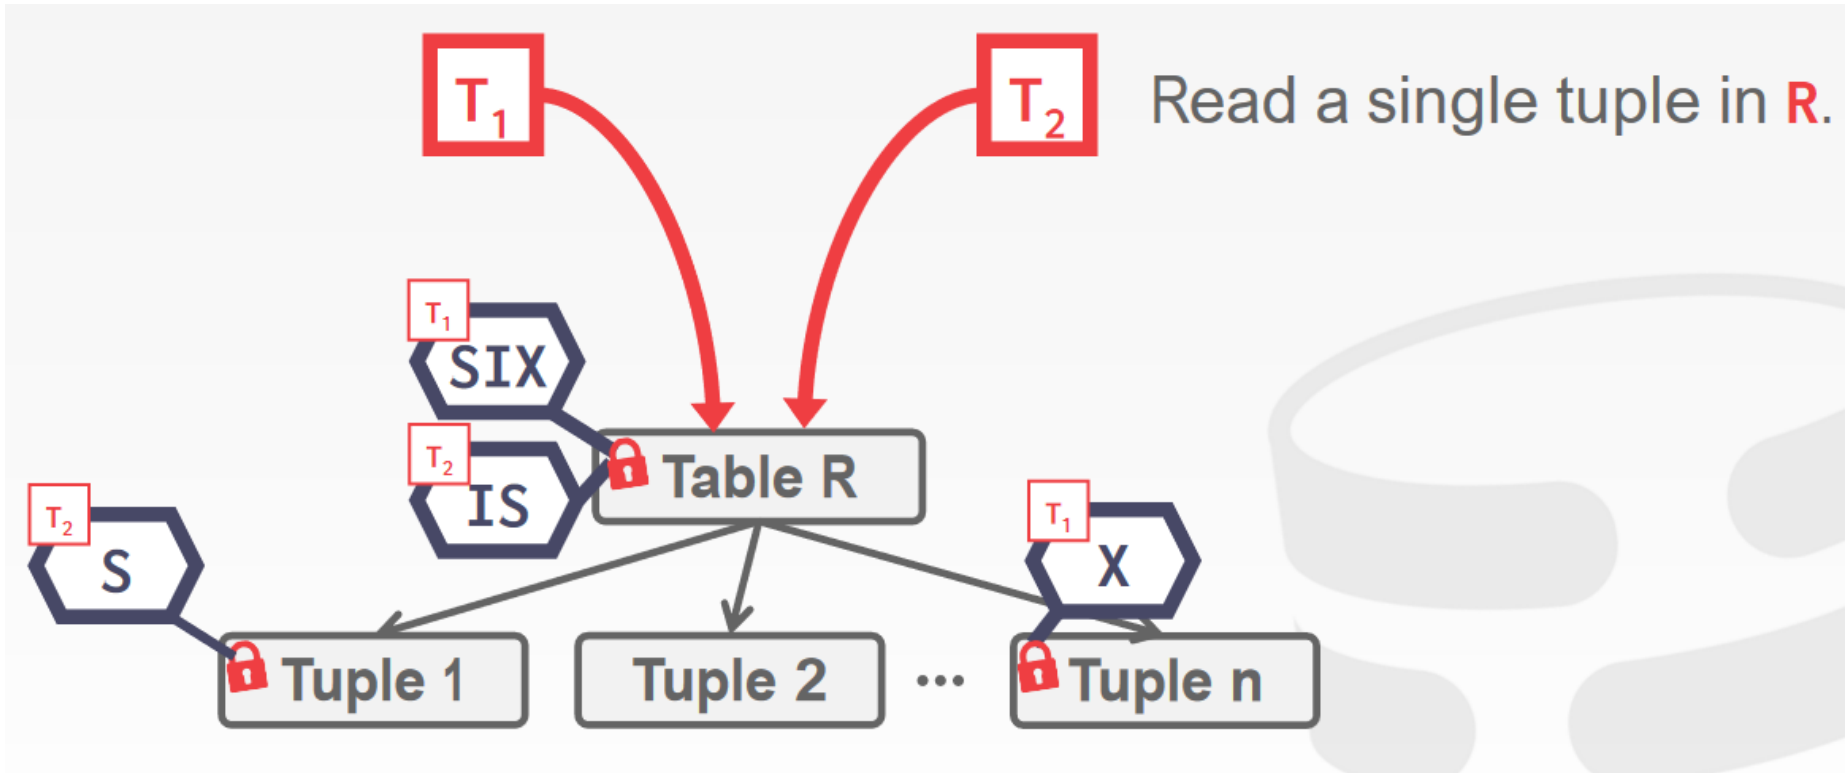


# Example






# Example






# Example





# Example

Scan all tuples in **R**.



**Are the locks of these transactions compatible?**



# Locking in Practice

- Different DBMSs use different locking methods
- The DBMS can automatically take care of locking, unlocking, etc.
- But there **are** methods for explicit, manual locking
  - Depends on which DBMS you are using
    - **PostgreSQL**: Locking in SHARE, EXCLUSIVE, and other modes
    - **MySQL**: Locking in READ,WRITE modes

 PostgreSQL

**ORACLE**

**IBM** **DB2**

 Microsoft  
SQL Server

 **MySQL**

```
LOCK TABLE <table> IN <mode> MODE;
```

```
SELECT 1 FROM <table> WITH (TABLOCK, <mode>);
```

```
LOCK TABLE <table> <mode>;
```