



**KOÇ
UNIVERSITY**

Database Management Systems

Hash-Based Indexing

M. Emre Gürsoy

Assistant Professor
Department of Computer Engineering

www.memregursoy.com



Motivation

- When tables are large, it is inefficient to search all values to find matches (aka: **linear search** or **full table scan**).

```
SELECT  Lname  
FROM    Employee  
WHERE   Ssn = 123456789
```

- In databases, an **index** is a data structure that enables locating data quickly without having to do a full table scan.
- We'll cover two types of indexing:
 - Hash-based indexing
 - Good for **equality selections** (not so much for **range selections**)
 - Tree-based indexing



“Create Index”

- Almost all DBMS have a **CREATE INDEX** command that allows you to create indexes on desired attributes.
 - Example from MySQL manual (on the right).
- Indexes are also used in DBMS internals.

```
CREATE [UNIQUE | FULLTEXT | SPATIAL] INDEX index_name
    [index_type]
    ON tbl_name (key_part,...)
    [index_option]
    [algorithm_option | lock_option] ...

key_part: {col_name [(length)] | (expr)} [ASC | DESC]

index_option: {
    KEY_BLOCK_SIZE [=] value
    | index_type
    | WITH PARSER parser_name
    | COMMENT 'string'
    | {VISIBLE | INVISIBLE}
    | ENGINE_ATTRIBUTE [=] 'string'
    | SECONDARY_ENGINE_ATTRIBUTE [=] 'string'
}

index_type:
    USING {BTREE | HASH}

algorithm_option:
    ALGORITHM [=] {DEFAULT | INPLACE | COPY}

lock_option:
    LOCK [=] {DEFAULT | NONE | SHARED | EXCLUSIVE}
```



Table Indexes

- Advantages:
 - Faster search and retrieval (**key reason!**)
- Disadvantages:
 - Indexes must also be updated at each insert/delete/upd.
 - Updating an index itself is work
 - Also, index must be locked before update, potentially affecting other transactions' throughput
 - Indexes cost space
- General advice: Create indexes when they are useful, but don't create redundant indexes!




Outline

- Hashing (basics)
- Static hashing with extensions
 - Overflow chaining
 - Linear probing
- Extendible hashing
- Linear hashing



Hashing


- A **hash table** is an associative array that maps **keys** to **values**.
 - Keys here do not mean keys in the relational model!
- A **hash function** is used to map data of arbitrary size (e.g., long strings) into fixed-size values (e.g., integers).
- A **hash collision** occurs when two different keys are mapped to the same hash value.
 - $\text{key1} \neq \text{key2}$, but $h(\text{key1}) = h(\text{key2})$
- Desired properties of hash functions:
 - Fast (not necessarily cryptographically secure)
 - Low collision rate





Static Hashing

- Allocate a large array that has one slot for each record
 - Potential hash fn: $h(k) \bmod N$, where N : # of records
 - Each slot holds one record

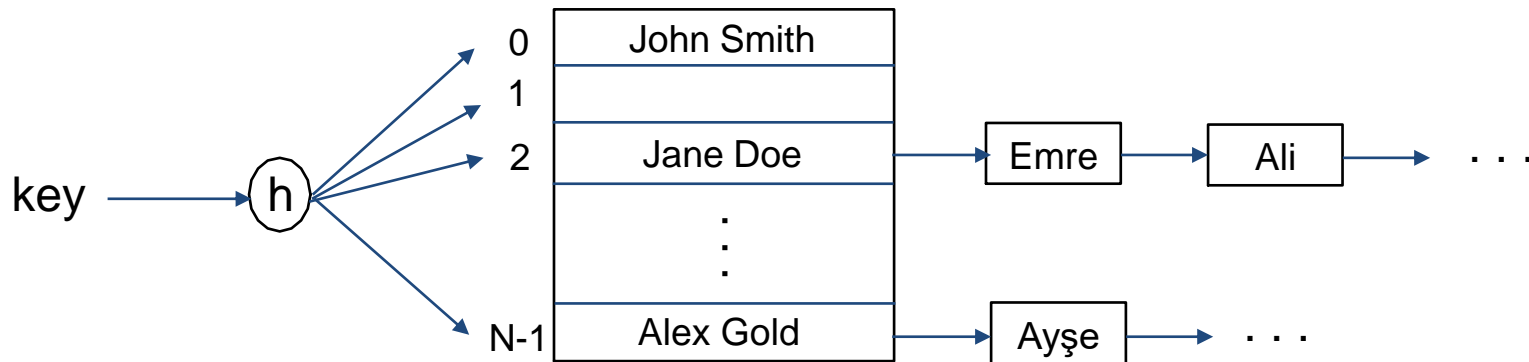


- Works when:
 - You know the # of records (N) ahead of time
 - You do not need to grow/shrink the hash table size!
 - Hash function is perfect, i.e., no hash collisions



Overflow Chaining

- A straightforward extension of static hashing to handle hash collisions
- If there is a hash collision, create **overflow buckets**
 - Overflow buckets become a **chain** (can be treated like a **linked list**)




- How do you **Search?** **Insert?** **Delete?**
 - Best case complexity? Worst case complexity?



Linear Probing

- Another extension of static hashing to handle hash collisions
- Resolve collisions by **linearly searching (probing) for the next free slot** in the hash table
 - Hash to the key's location and start scanning
 - When you find an empty slot, insert the record there





Linear Probing

hash(key)

A


B

C

D

E


F



<key> | <value>




Linear Probing





Linear Probing





Linear Probing

hash(key)

A

B

C

D

E

F

B | val

A | val

C | val

D | val






Linear Probing

hash(key)

A
B
C
D
E
F



B | val

A | val

C | val

D | val

E | val





Linear Probing

hash(key)

A

B

C

D

E

F



B | val

A | val

C | val

D | val

E | val

F | val

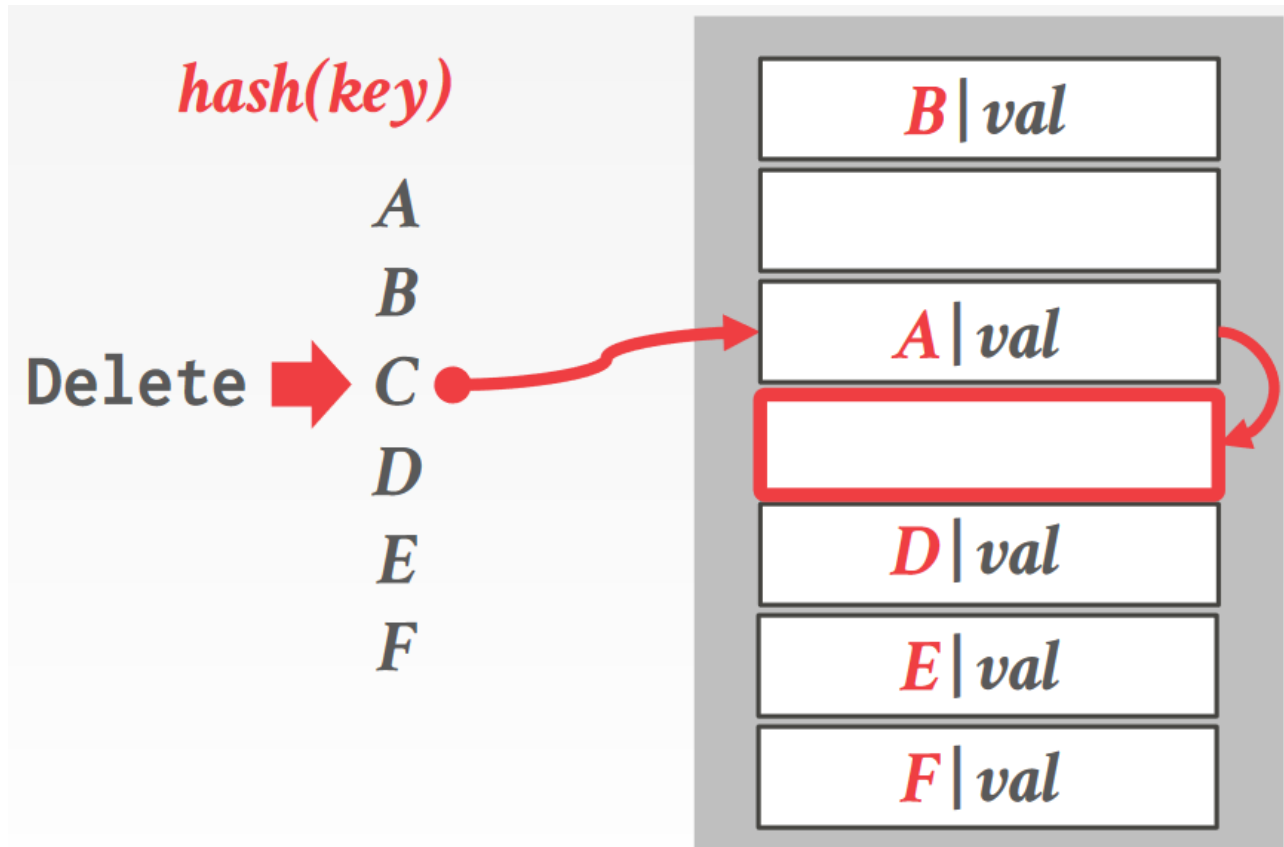


How to perform **Search**, e.g., search for G?



Linear Probing

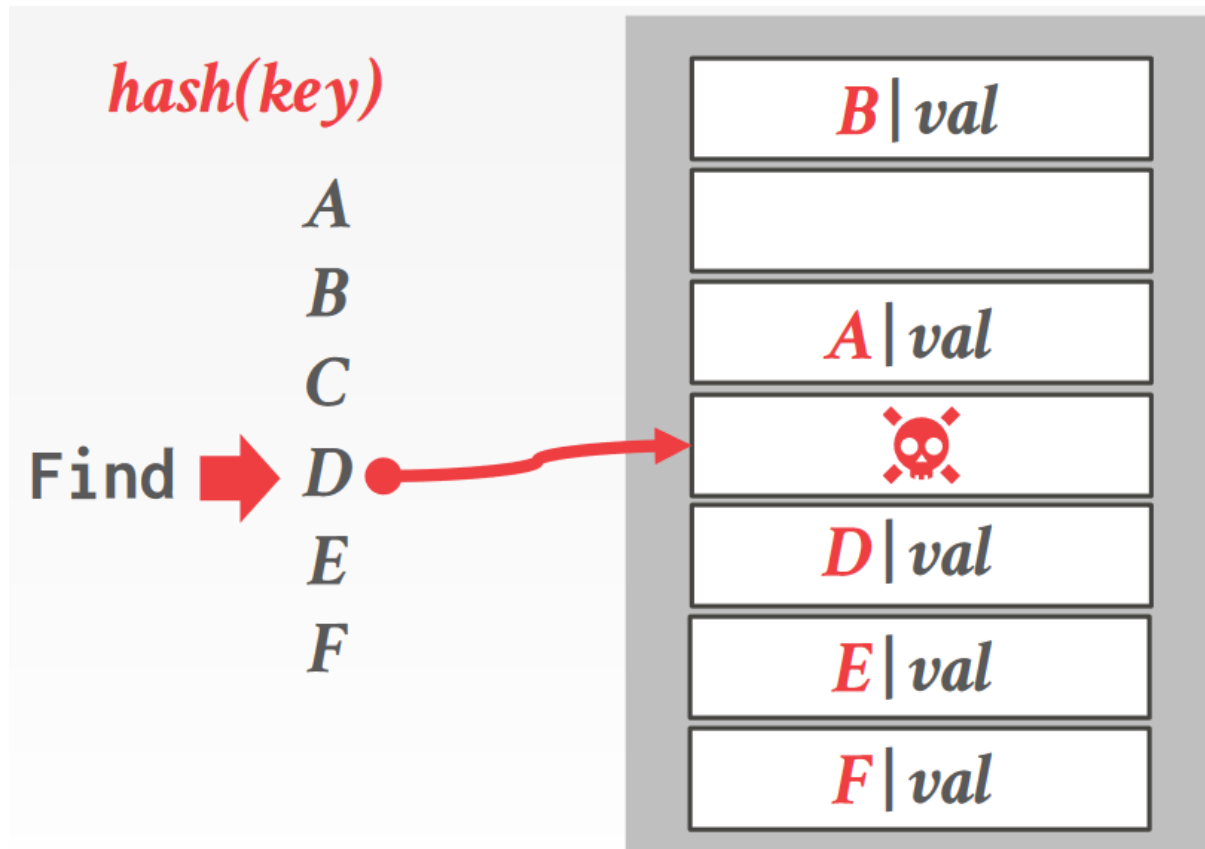
- Potential problem with **Delete**:





Linear Probing

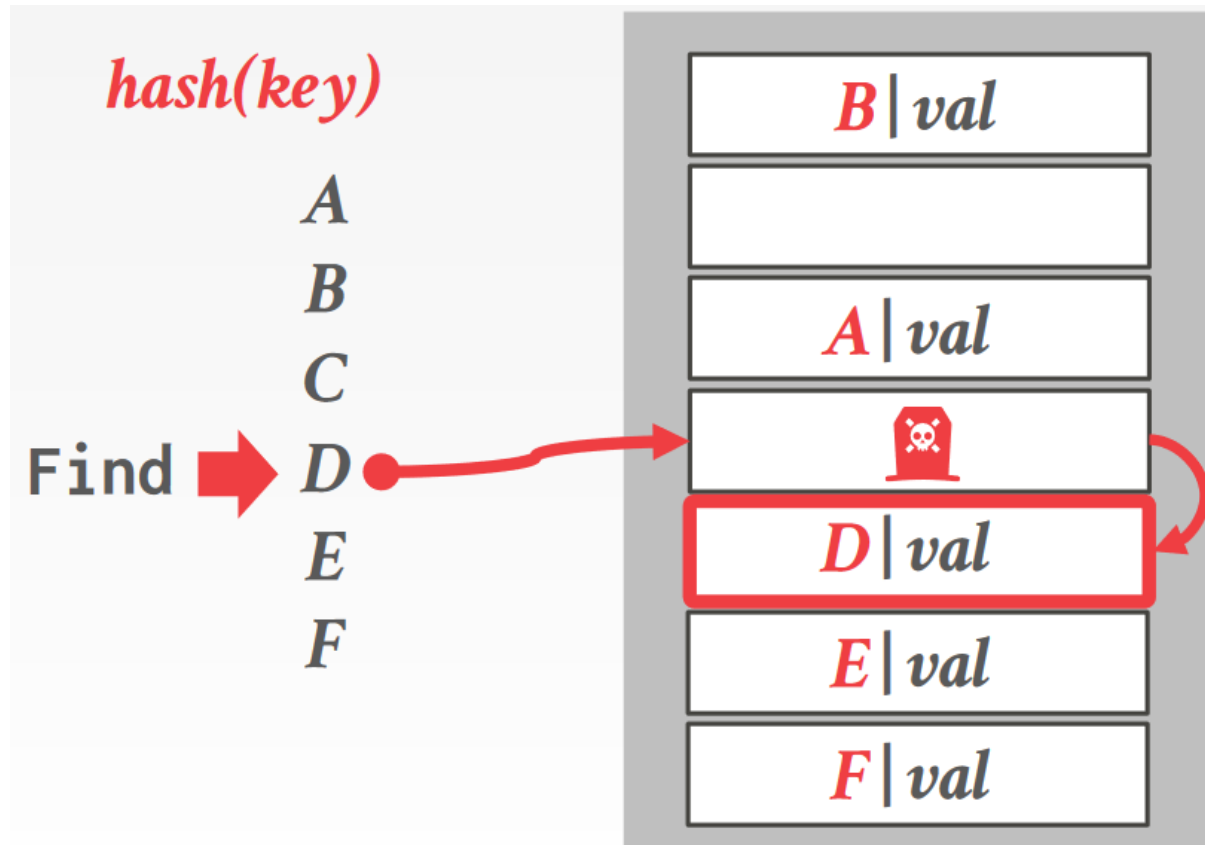
- Potential problem with **Delete**:





Linear Probing

- The “**Tombstone**” approach to solving this problem:





Non-Unique Hash Keys

- For simplicity, we'll make the assumption that our **hash keys are unique**.
 - Students(Sid, Fname, Lname, DoB, Address)
 - If **hash key = Sid**, then our assumption holds
 - If **hash key = Lname**, then we have non-unique keys
- **Handling non-unique keys:** Allow for redundancy
 - Store duplicate keys' entries separately in the hash table
 - Same idea can be used in other hashing schemes as well (not just linear probing)

XYZ value1
ABC value1
XYZ value2
XYZ value3
ABC value2



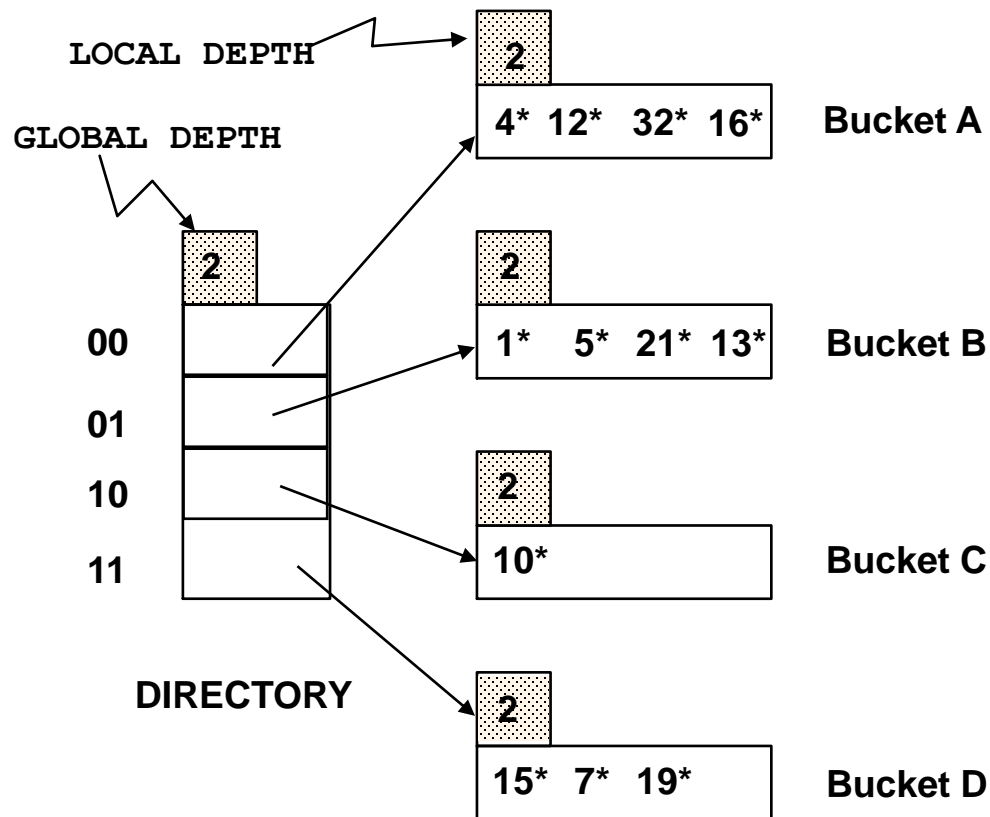
Outline

- Hashing (basics)
- Static hashing with extensions
 - Overflow chaining
 - Linear probing
- Extendible hashing
- Linear hashing



Extendible Hashing

- Dynamic hashing approach
- Directory contains pointers to buckets
- To find bucket for r , take last global depth # of bits of $h(r)$
 - $h(r) = 5 = \text{binary } 101$, so it is placed in bucket pointed to by 01

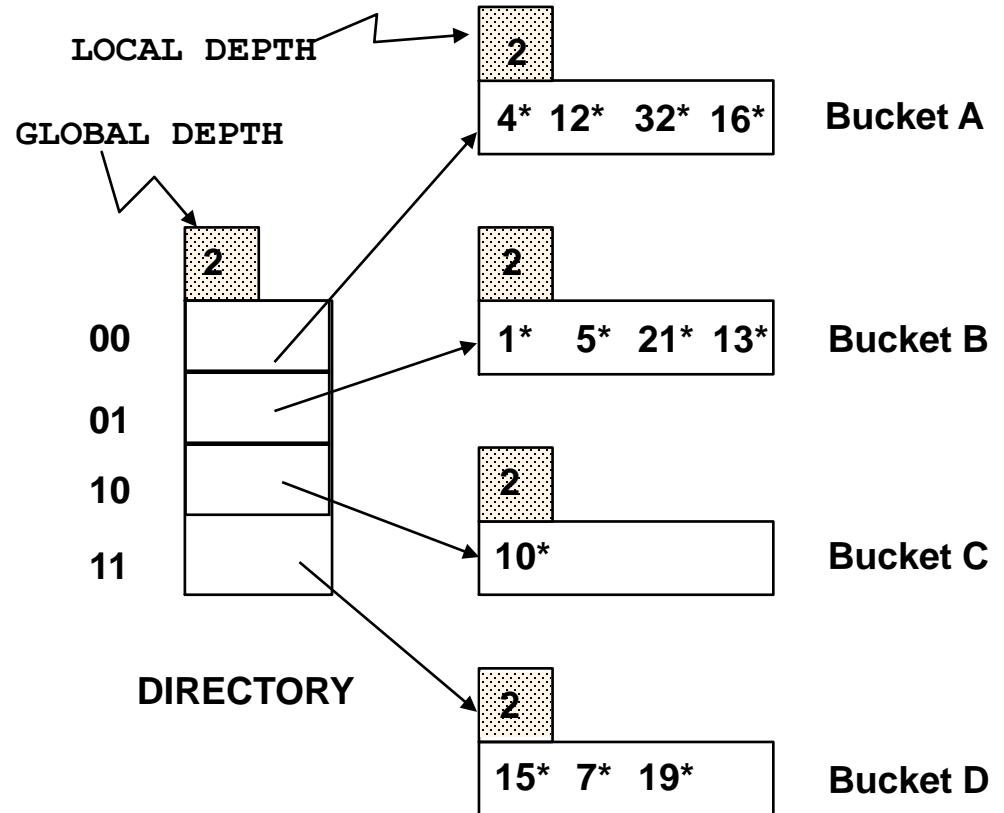


- We represent r by $h(r)$ in the figures for clarity.



Extendible Hashing

- Buckets have a **capacity** (here, 4).
- When a bucket is full and a new insertion arrives, the bucket must **split**.
- The directory may or may not have to **double** when a split occurs.
 - Compare global depth vs local depth

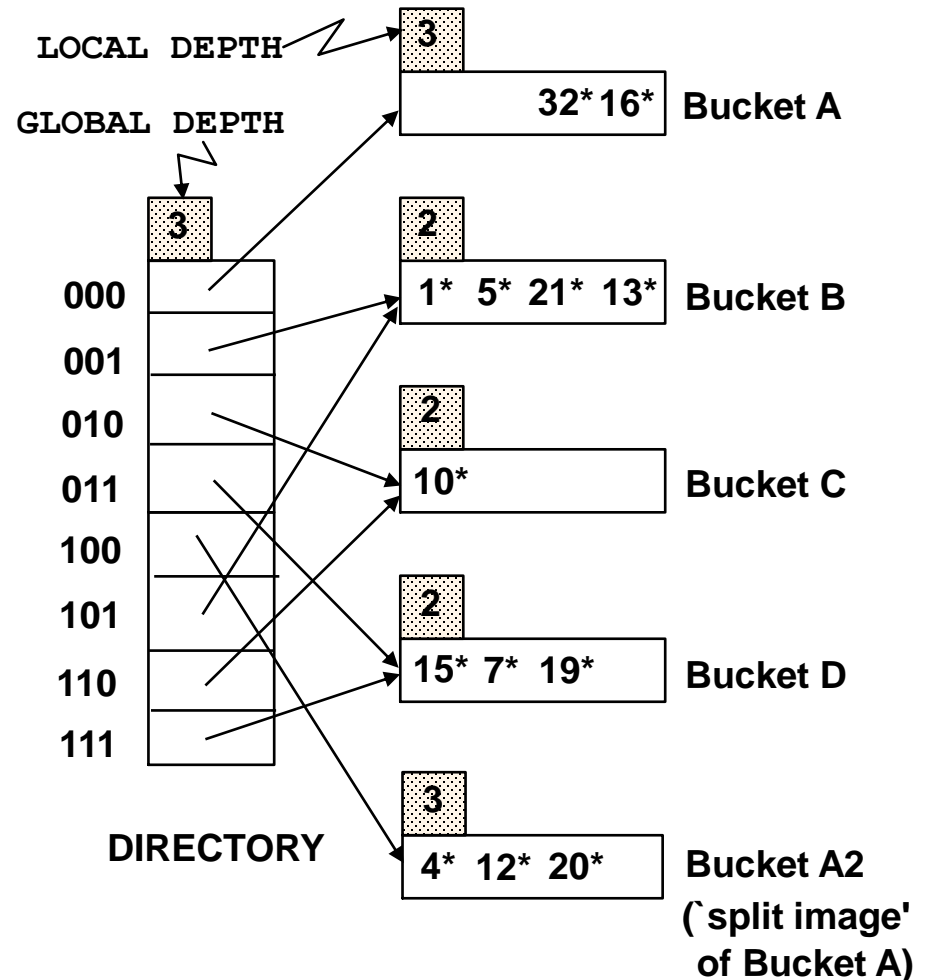


Insert $h(r)=20$ (Causes Splitting and Doubling)
20 is 10100 in binary format



Extendible Hashing

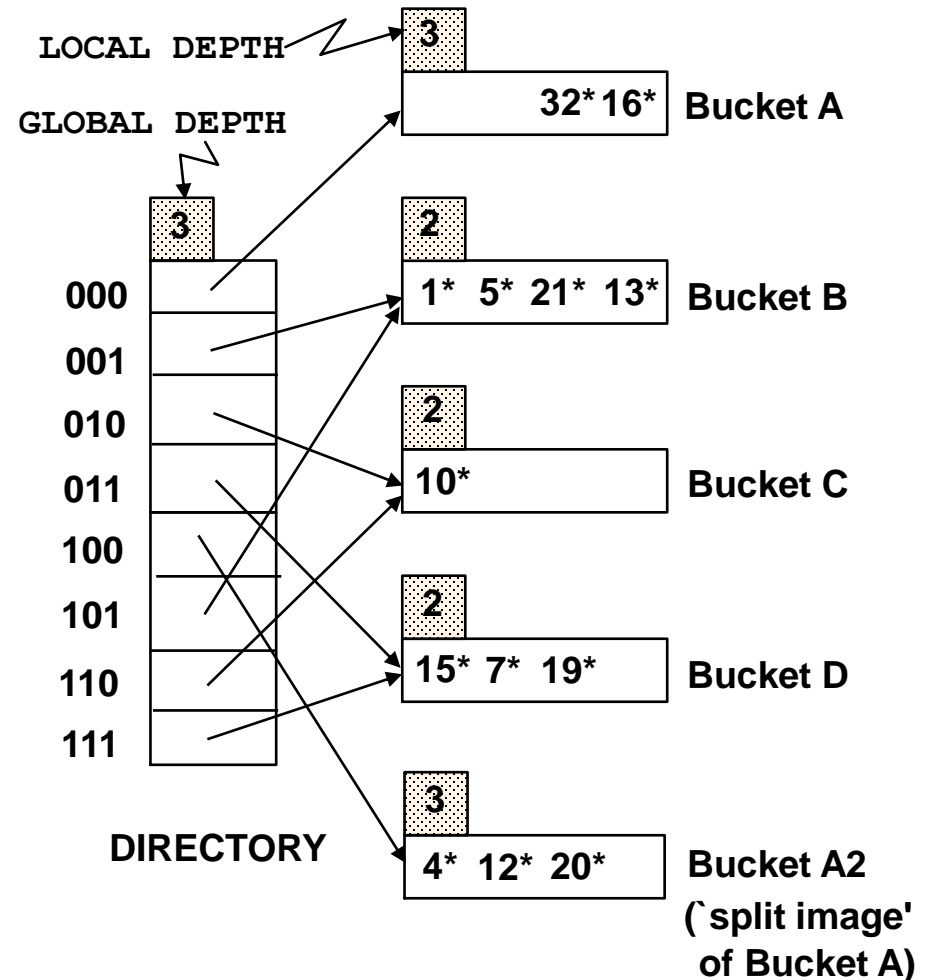
- When does **splitting** cause directory **doubling**?
 - Local depth > global depth
- Notice that buckets B, C, D are not modified by insertion of 20
- Will inserting to bucket C cause splitting? Or directory doubling?
 - How about bucket B?





Extendible Hashing

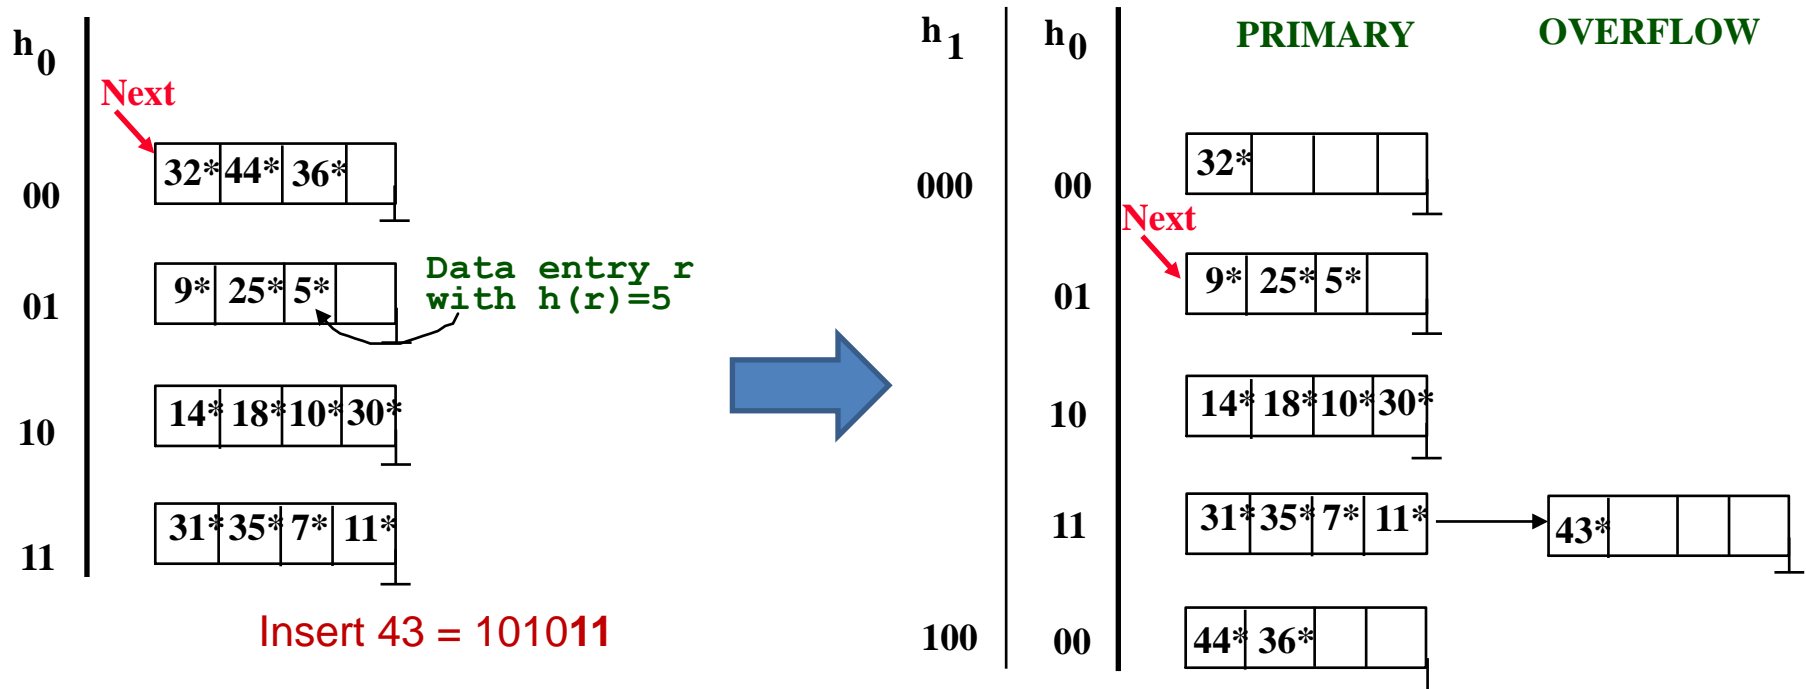
- How about **deletions**?
 - If bucket becomes empty and has a split image, they **merge**
 - If each directory element points to same bucket as its split image, **halve** the directory





Linear Hashing

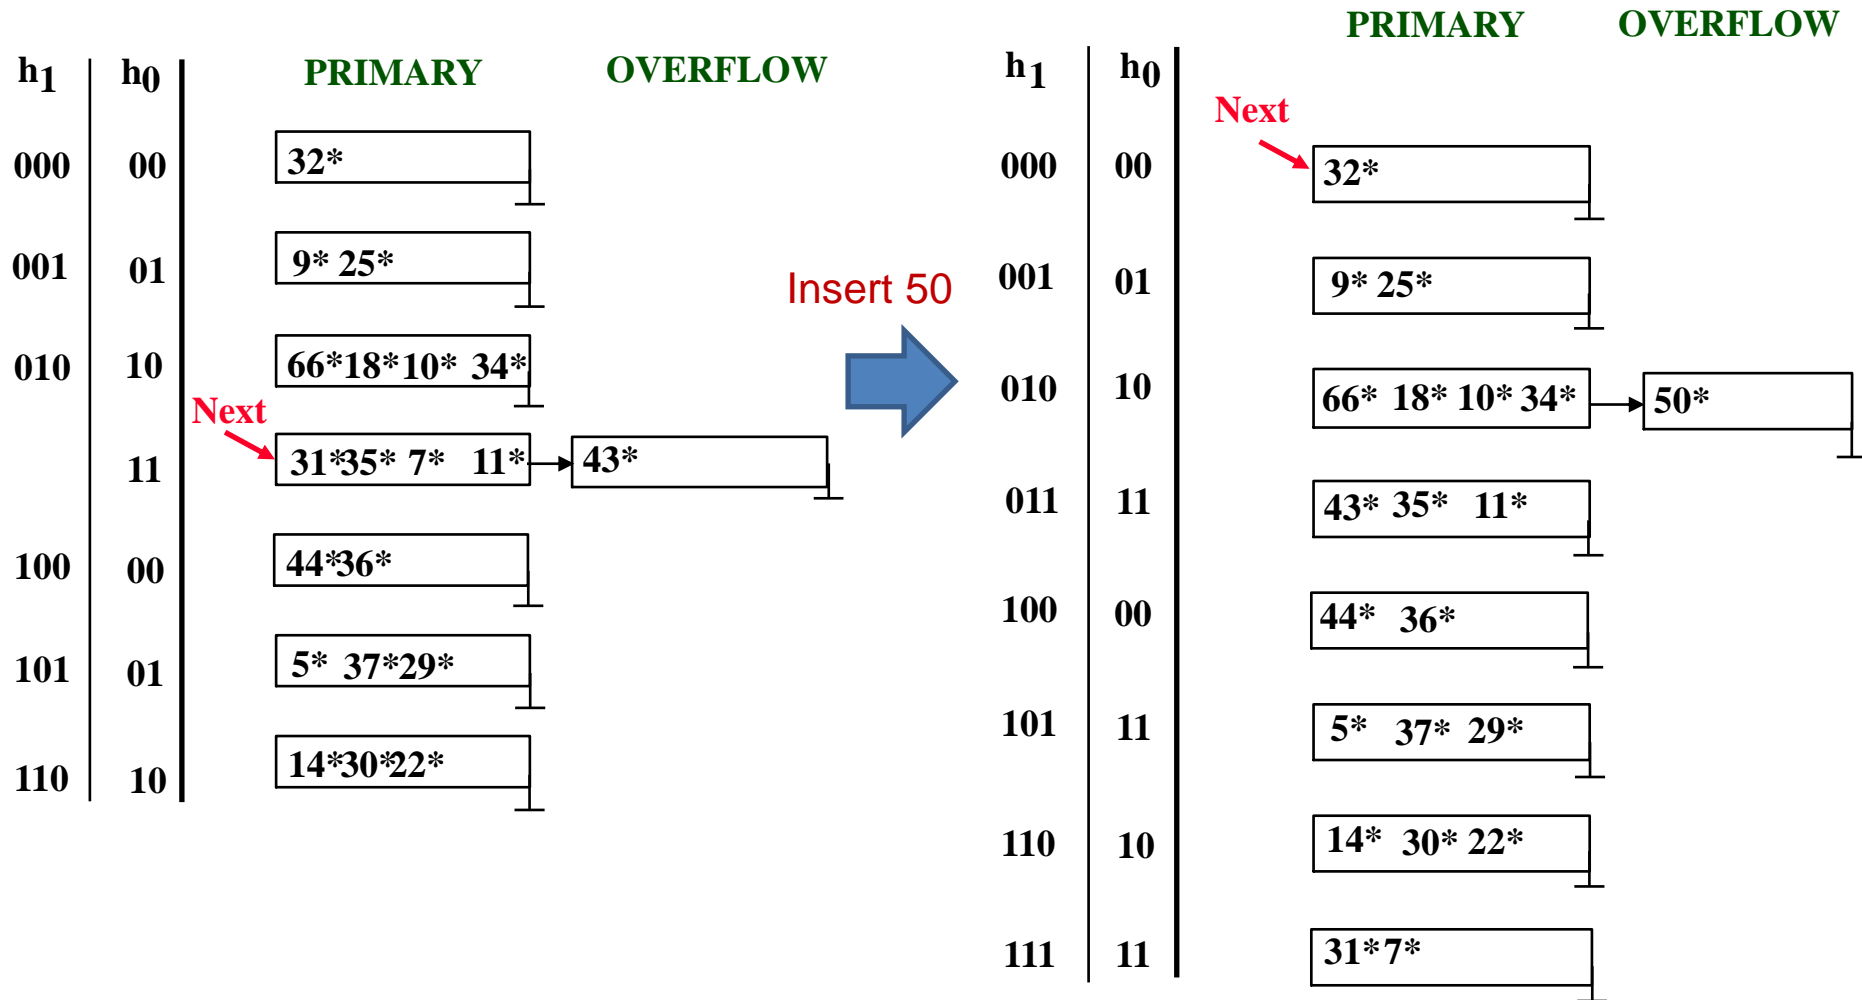
- As opposed to extendible hashing, always split the bucket that is pointed to by the “Next” pointer.
 - Splitting is **not guaranteed** to fix overflows
 - Next** pointer is linear – will eventually get to overflows
- If no overflow, no need to split or advance **Next** pointer





Linear Hashing

- End of a round:





Linear Hashing

- How to handle empty buckets resulting from **deletions**?
 - Delete 7, then delete 31
- **Strategy #1:** Merge buckets, decrement **Next** pointer
 - When there's a new overflow eventually, you'll need to re-do the work and the splitting
 - You may again end up with an empty bucket
- **Strategy #2:** Do nothing (let the bucket remain empty)
 - Cannot reclaim previously allocated memory
 - (This is the strategy we will use.)



Summary

- Hash-based indexes are fast data structures that provide $O(1)$ average time search.
 - Hash functions and hash-based indexing are also used in DBMS internals and for optimizing storage.
- Different hash-based indexing methods:
 - Static hashing with extensions: Linear probing, overflow chaining
 - Extendible hashing
 - Linear hashing
- Hash-based indexes are good for equality searches but not for range searches.
 - Next topic: Tree-based indexes, which better support range searches