



**KOÇ
UNIVERSITY**

Database Management Systems

Tree-Based Indexing

M. Emre Gürsoy

Assistant Professor
Department of Computer Engineering

www.memregursoy.com



Motivation


- **Range selection** (range search):
 - Different from **equality selection**
 - Common in practice
 - Not supported effectively by hash-based indexing
- **Tree-based indexing** techniques support both **range selection** and **equality selection** effectively.
 - Often used for table indexing in DBMS
- We'll cover two types of tree-based index structures:
 - ISAM Trees
 - Static, not preferred
 - B+ Trees
 - Also known as: "the coolest data structure ever"

```
SELECT *  
FROM   Students  
WHERE  Gpa > 3.0
```




ISAM Trees

- Tree-based data structure (non-leaf pages are static)
 - Leaf pages store actual data entries
 - Non-leaf pages act as the “guide”






Example ISAM Tree




- Structure of leaf pages: only store data entries
- Structure of non-leaf pages: **pointers** and **keys**





Insertion

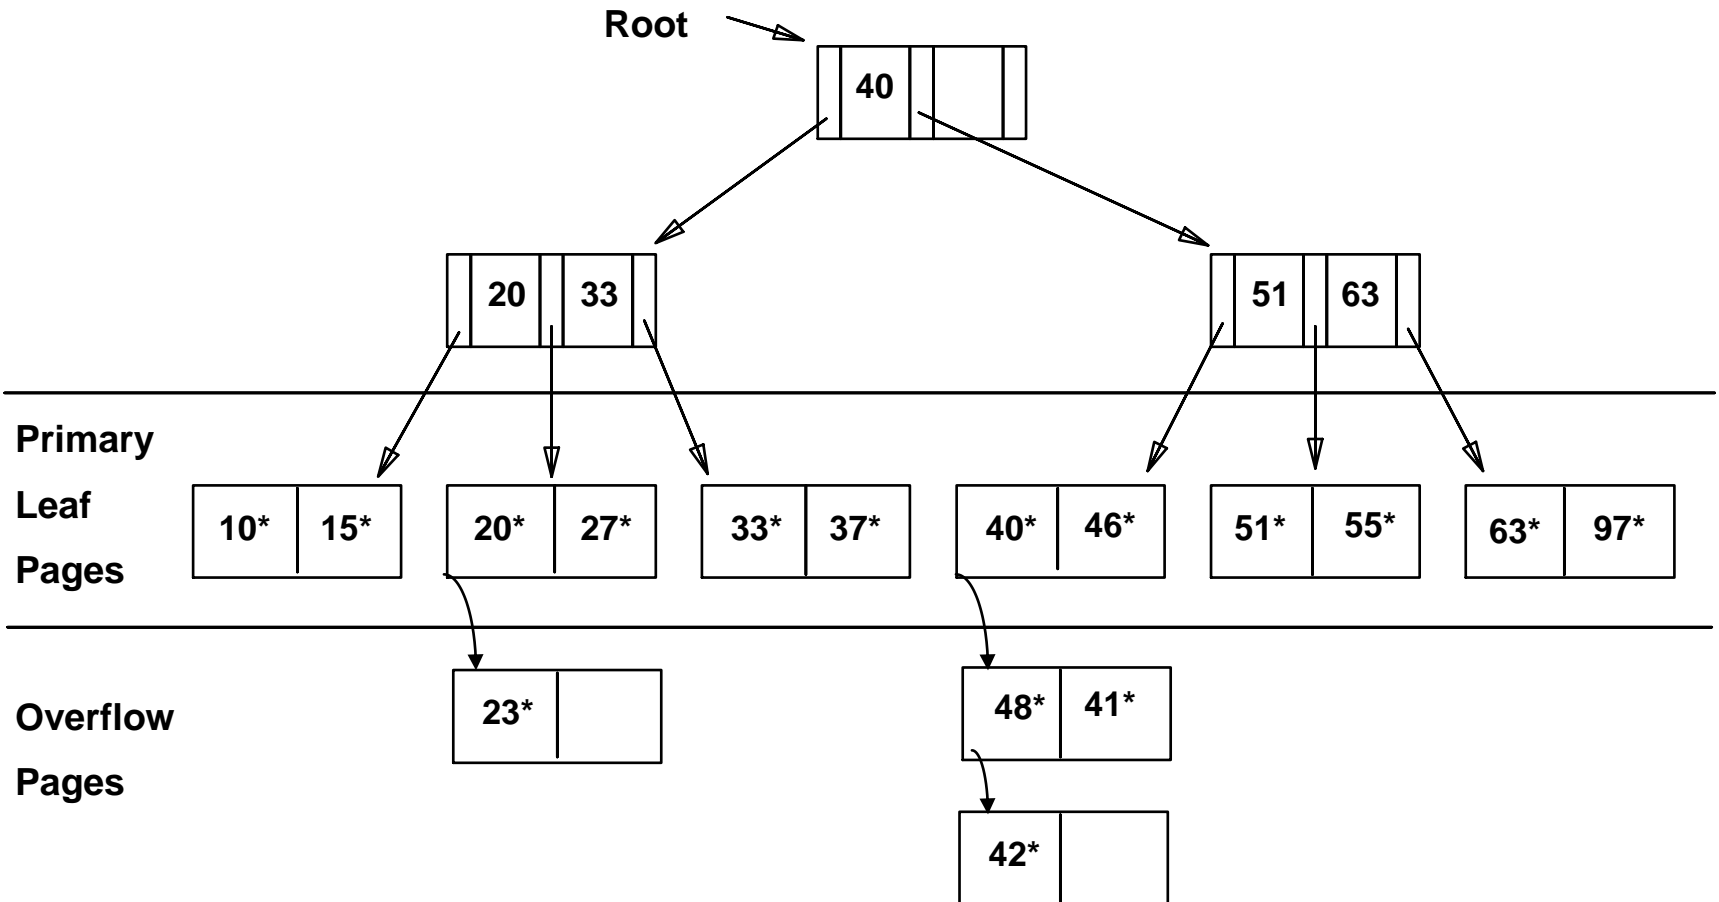
- After Inserting **23***:





Insertion

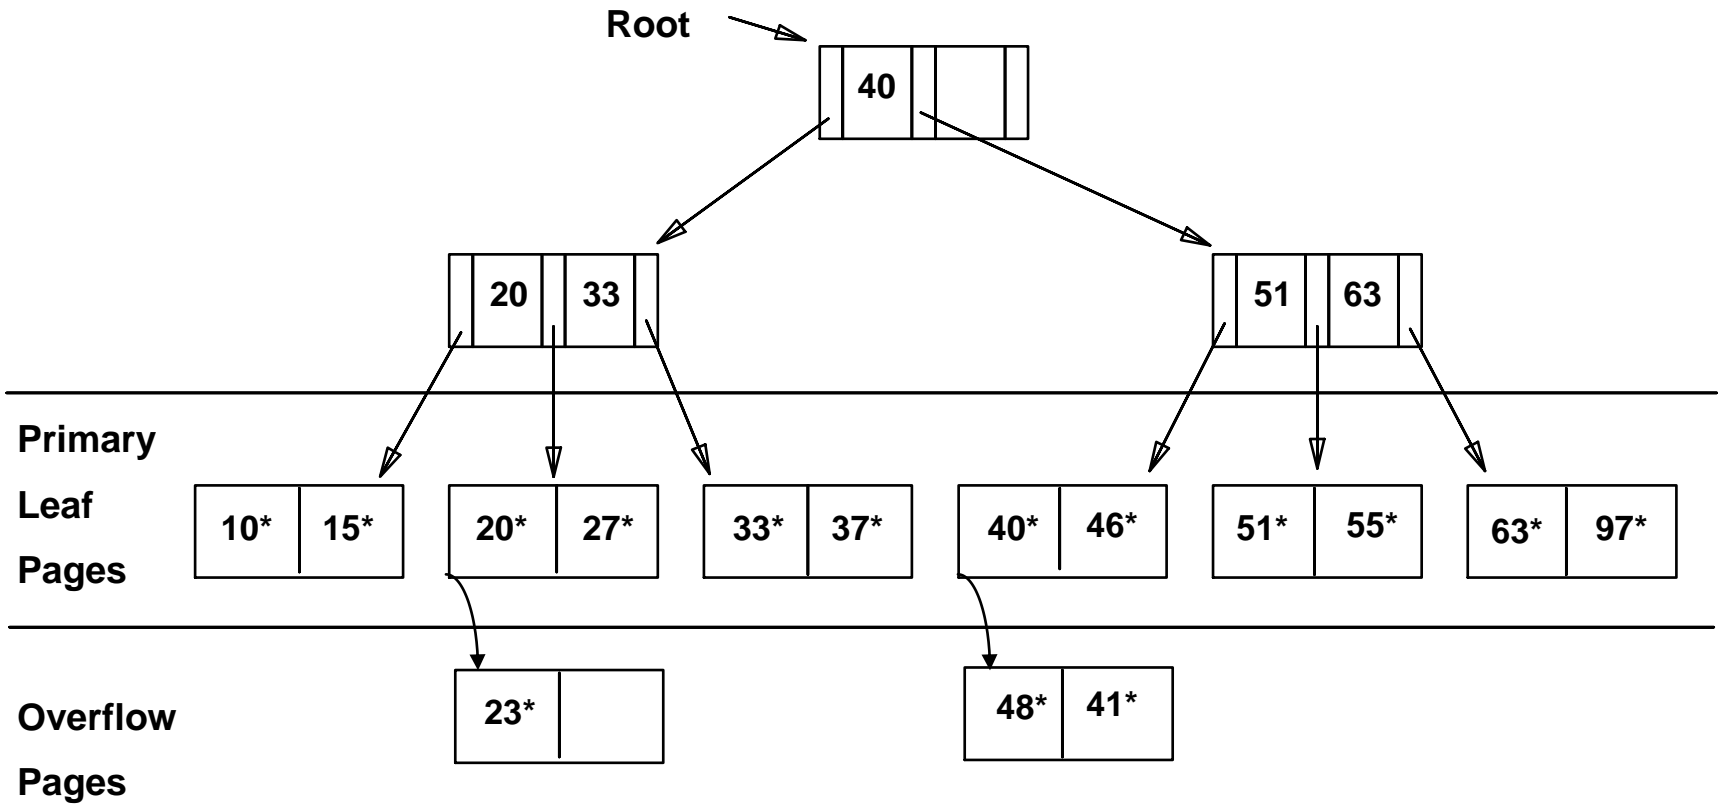
- After Inserting **23***, **48***, **41***, **42***:





Deletion

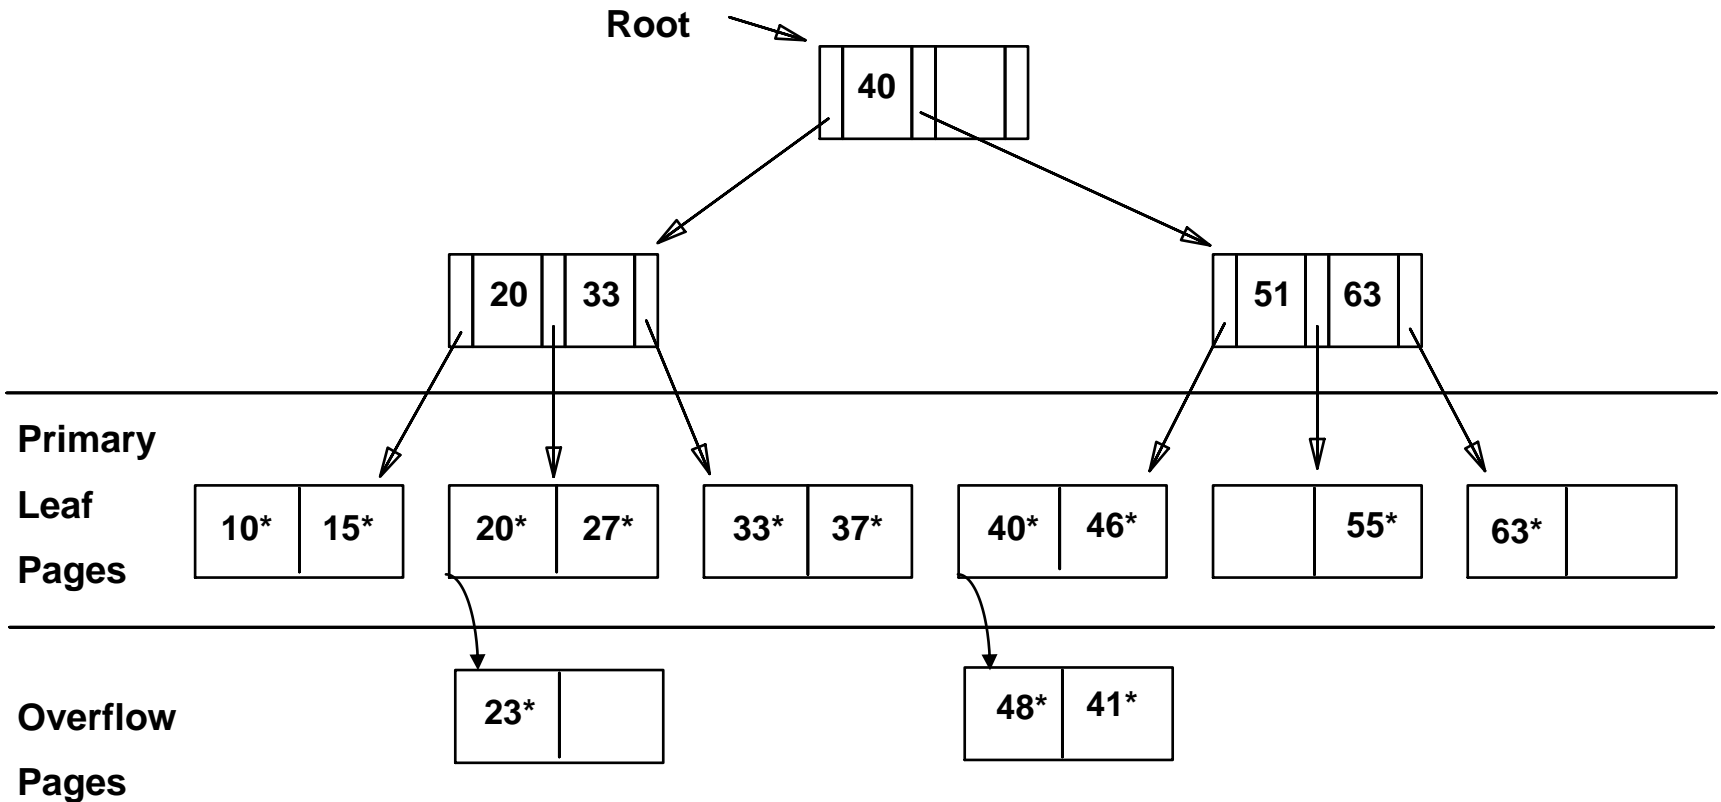
- After Deleting 42*:





Deletion

- After Deleting 51*, 97*:



Note that 51 can appear in the non-leaf pages but not in the leaf pages!



ISAM Shortcomings

- ISAM tree is a **static** data structure
 - Leaf pages and overflow pages may change since they contain actual data
 - Non-leaf pages (constituting the index) **do not change**
- ISAM tree cannot balance itself after inserts/deletes
 - Long overflow pages are possible
 - What if data distribution changes over time?
- **What are the best-case and worst-case complexities of searching for a record in an ISAM tree?**




B+ Trees

- A **self-balancing** tree data structure that keeps data sorted and allows search, sequential access, insertion, and deletion in **$O(\log N)$** time.
 - Belonging to the “family” of **B-Trees**
 - B-tree, B+ tree, B* tree, B^{link} tree, ...
 - **Very commonly** used in DBMS
- Properties:
 - Perfectly height-balanced at all times
 - Every node other than the root must be at least half full
 - Order = **m** means nodes have at most **m** children
 - Thus, each node holds at most **m-1** data entries



Example B+ Tree

- Search for 5^*
- Search for 15^*
- Search for all entries ≥ 24 and ≤ 35
 - Use pointers among the leaves






Insertion

- Start at root, find the leaf node *L* to insert to.
- If *L* has enough space:
 - Insert into *L*, ensure the leaf is sorted, done!
- Otherwise:
 - Must split *L* into two nodes
 - Redistribute entries evenly, copy up the middle key
 - This split may cause recursive splits in upper levels
- To handle splits in upper levels:
 - Redistribute entries evenly, push up the middle key
 - Always check if the push up causes further splits





Insertion

- Inserting **15*** is straightforward
- How about inserting **8***?






Insertion





Insertion

- After inserting 8*
 - Notice how the root was split, causing increased height
 - The end tree is balanced
 - “At least half full” constraint remains satisfied (how?)






Deletion

- Start at root, find leaf L where entry resides.
- Remove the entry from L .
 - After removal, if L is at least half full, done!
 - Otherwise, try to re-distribute by borrowing from L 's sibling (adjacent node with same parent as L).
 - If re-distribution fails, merge L and its sibling.
- If a merge occurs, we must delete the entry (pointing to L or its sibling) from the parent of L .
- Merge could propagate all the way up to the root and cause decreased tree height.



Deletion


- Delete 19^* , then delete 20^*
 - Deleting 19^* is easy
 - Deleting 20^* requires re-distribution





Deletion


- Result after **19*** and **20*** are deleted
 - Notice how **24** in **L**'s parent changes to **27**
- Now, let's delete **24***






Deletion

- Must merge leaves and throw away **27** in the parent:



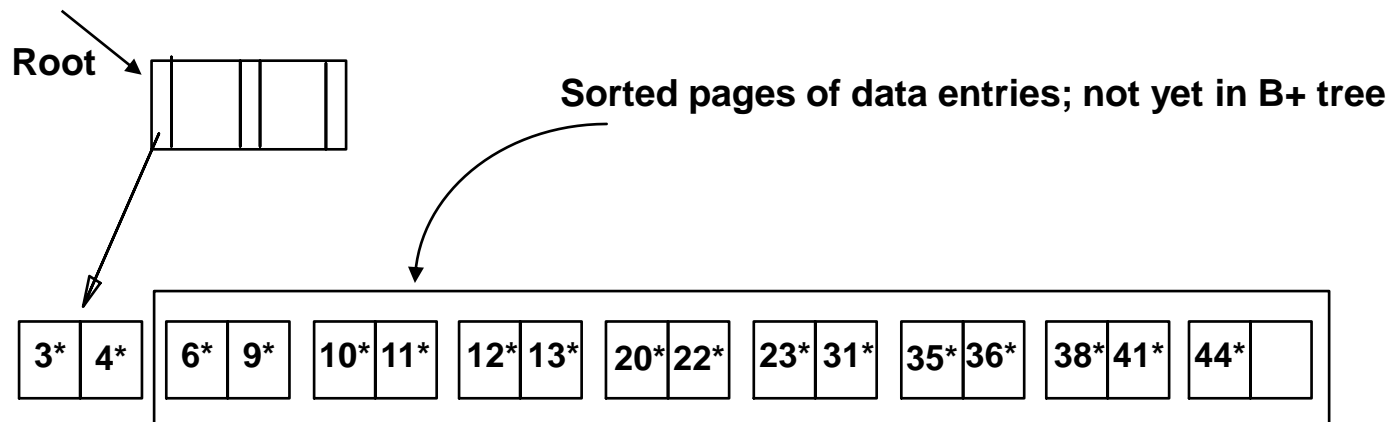
- But this time, the parent is no longer half full!
 - Merge propagates up one more level







Bulk Loading

- If we have a large collection of records and we want to create a B+ tree with them, doing so by **repeatedly inserting** records in **random order** can be slow.
 - Can be solved by **bulk loading** of the B+ tree.
- Sort all records, initialize empty root and start adding.






Bulk Loading






Bulk Loading





Clustered vs Unclustered Index

- In a **clustered index**, the physical storage layout of the tuples in pages will match the order in which they appear in the index.
 - How does this become useful in physical disk I/O?
 - In what ways could this be difficult? (Insert/delete)
- **Note:** Clustered indexes are not limited to B+ trees.





Summary

- **Tree-structured indexes** are good fit for indexing tables
 - Efficiently support **range selection** and **equality selection**
- **ISAM tree** is a static structure
 - Only leaf pages are modified
 - Insertions and deletions are straightforward
 - Long overflows can degrade performance
- **B+ tree** is a dynamic structure
 - Very popular and commonly used
 - Insertions and deletions are more involved than ISAM, but search remains $O(\log N)$
 - Use bulk loading for the initial construction of B+ tree
- **Clustered indexes** determine physical storage layout