



**KOÇ
UNIVERSITY**

Database Management Systems

Query Execution and Optimization

M. Emre Gürsoy

Assistant Professor
Department of Computer Engineering

www.memregursoy.com



Introduction

- SQL is a **declarative** language, not procedural
 - You tell the DBMS **what** you want, but you don't tell the DBMS **how** to compute it
 - The DBMS must decide how to compute it
 - And the DBMS must find a way to compute it **efficiently**

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating =  
      (SELECT MAX (S2.rating)  
       FROM Sailors S2)
```

What are the alternative ways to compute the answer to this query? Which way is fastest?



Overview

Queries

```
Select *  
From Blah B  
Where B.blah = blah
```

Query Parser

Query Optimizer

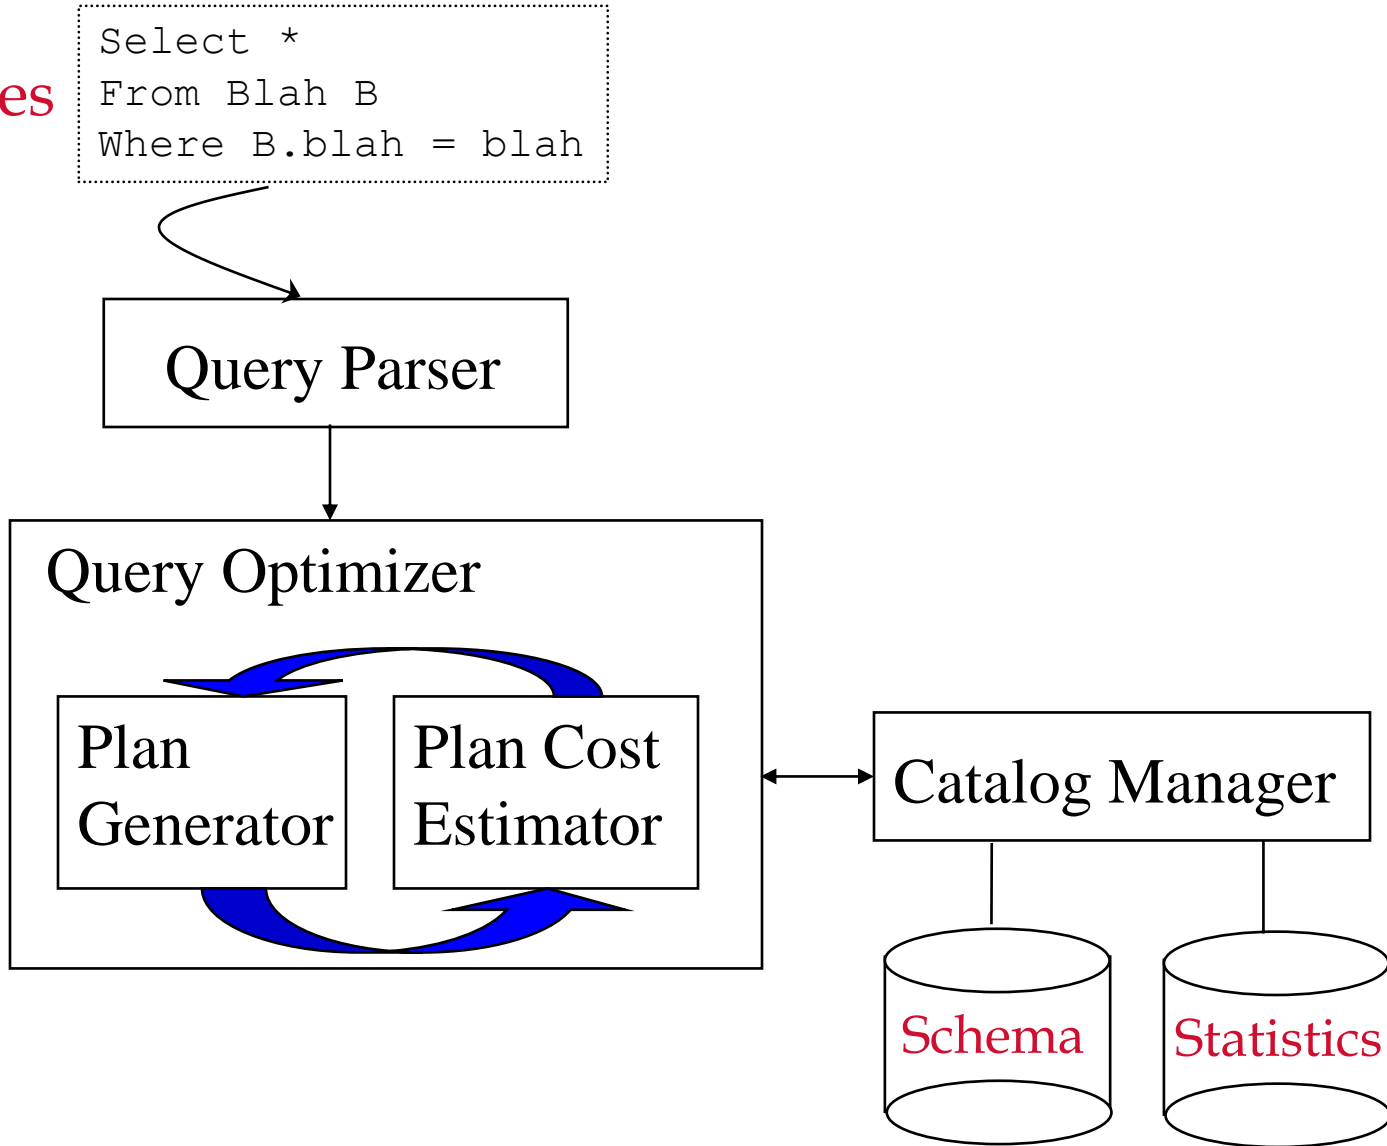
Plan
Generator

Plan Cost
Estimator

Catalog Manager

Schema

Statistics





Steps of Query Optimization

- **Step 1:** SQL query is parsed and divided into **blocks**.
- **Step 2:** Each block is converted to its “relational algebra equivalent”, which is used to construct its **Relational Algebra tree (RA Tree)**. This is the initial plan.
 - The DBMS may not internally draw an RA tree, but this is how we (humans) reason about it.
- **Step 3:** A subset of alternative plans are generated.
- **Step 4:** Costs of the different plans are estimated.
 - The DBMS picks and executes the plan with the lowest estimated cost.



SQL Blocks

- SQL queries are optimized by **decomposing** them into a collection of smaller units, called **blocks**.
 - Typically, query optimizer concentrates on optimizing a single block at a time.
- A **block** is an SQL query with:
 - No nesting
 - Exactly 1 SELECT and 1 FROM clause
 - At most 1 WHERE, 1 GROUP BY and 1 HAVING clause

SELECT	S.name
FROM	Student S, Takes T
WHERE	T.cid='415' and
	S.ssn=T.ssn



Nested Queries

- If the query is nested (but not correlated):
 - The inner query is the first block to be processed
 - Its output is finalized
 - The output is incorporated into the outer query as if it had been part of the original statement
 - The outer query is processed as the second block

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating = (SELECT MAX (S2.rating)  
                  FROM Sailors S2)
```



Correlated Nested Queries

- If the query is a **correlated** nested query, we can't evaluate the inner query once and be done with it
- In this case, the typical strategy is to evaluate the inner query for each different tuple of Sailors
 - Repeat the inner block many times

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS
  (SELECT *
   FROM Reserves R
   WHERE R.bid=103
   AND R.sid=S.sid)
```



RA Tree

- Say that I have an SQL block, how do I convert it into a query evaluation plan?
 - **Relational Algebra tree (RA tree)**
 - The RA tree is **procedural**, i.e., it tells you what to do step-by-step

```
SELECT S.sname  
FROM Reserves R, Sailors S  
WHERE R.sid=S.sid AND  
      R.bid=100 AND S.rating>5
```

What is the corresponding
RA tree of this query?


- First express the query in relational algebra:

$$\pi_{sname}(\sigma_{bid=100 \wedge rating > 5}(Reserves \bowtie_{sid=sid} Sailors))$$



RA Tree



$\pi_{sname}(\sigma_{bid=100 \wedge rating > 5}(Reserves \bowtie_{sid=sid} Sailors))$





RA Tree

```
SELECT S.rating, COUNT (*)  
FROM Sailors S  
WHERE S.rating > 5 AND S.age = 20  
GROUP BY S.rating
```

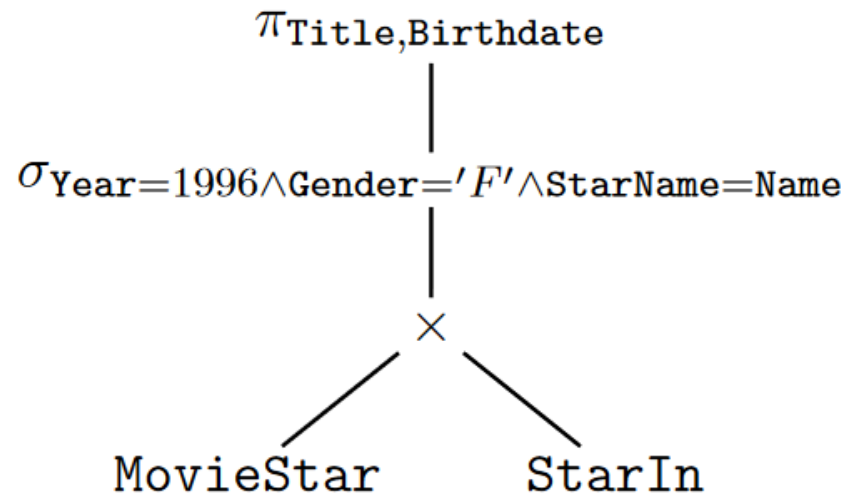




RA Tree

- “Find the birthdate and movie title for those female stars who appeared in some movie in 1996.”

```
SELECT Title, Birthdate  
FROM    MovieStar, StarIn  
WHERE Year=1996 AND Gender='F' AND Name=StarName
```





RA Equivalences

- Note that given one SQL query, there may be **multiple equivalent ways** to write it in relational algebra.
 - Cartesian product, then selection
 - Join MovieStar and StarIn wrt Name=StarName

```
SELECT Title, Birthdate  
FROM   MovieStar, StarIn  
WHERE  Year=1996 AND Gender='F' AND Name=StarName
```

- There may be optimizations you can perform on the RA query while maintaining equivalence:

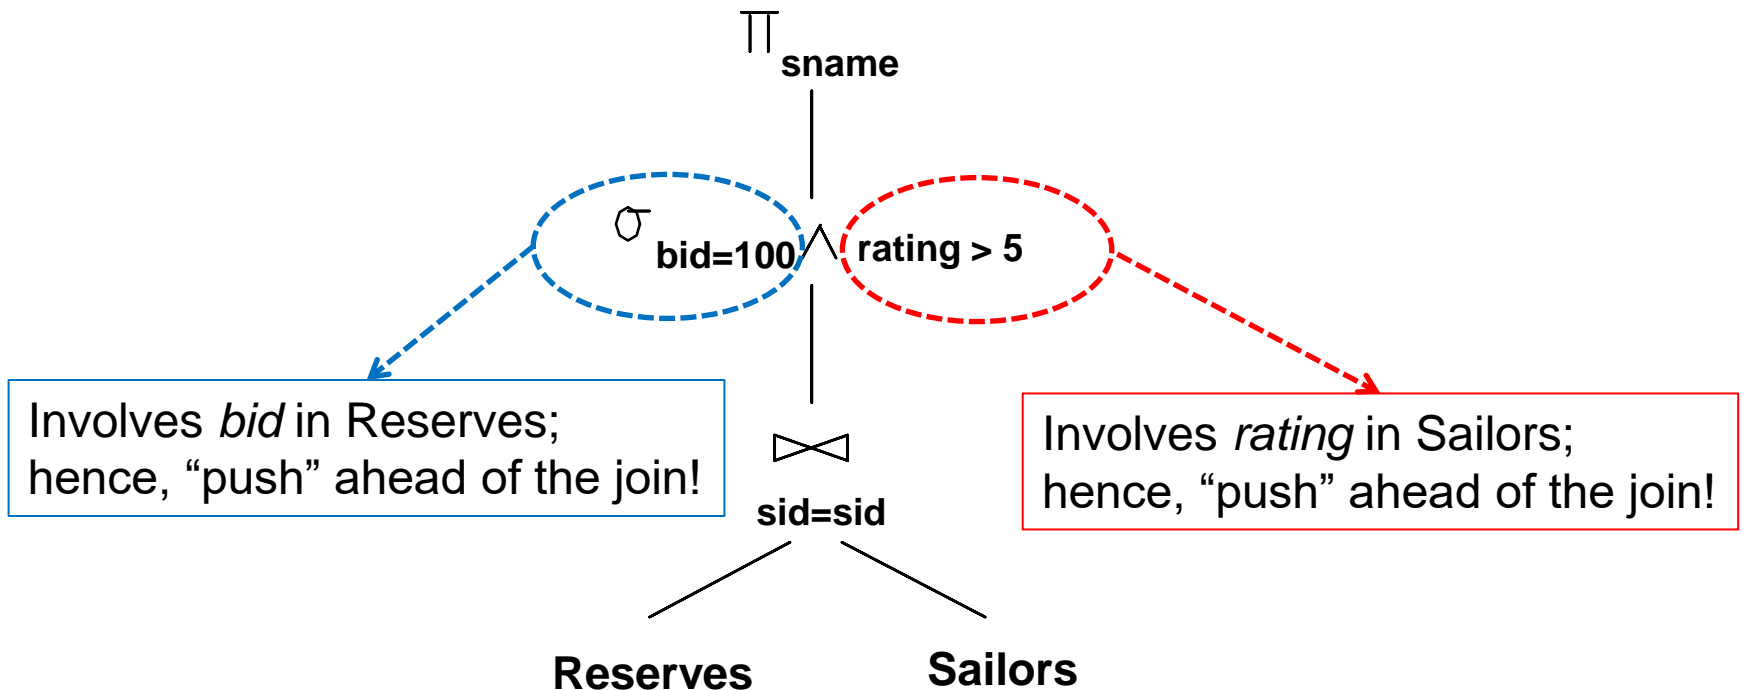
$$\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$$

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$






Operator Pushdown

- You can decrease the cost of an RA tree (aka: query execution plan) by **pushing some operators down**
 - Be careful – **do not hurt equivalence!**








Operator Pushdown





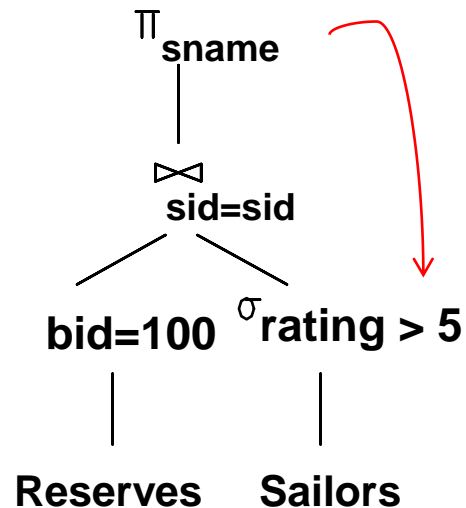

Operator Pushdown





Operator Pushdown

Replacing a {Cross Product + Select} with a Join can also be considered an instance of operator pushdown.


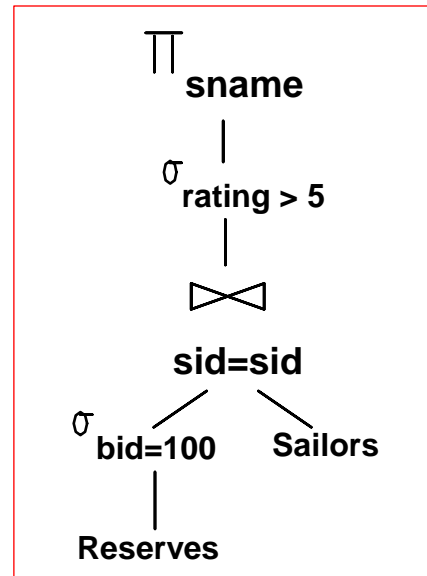
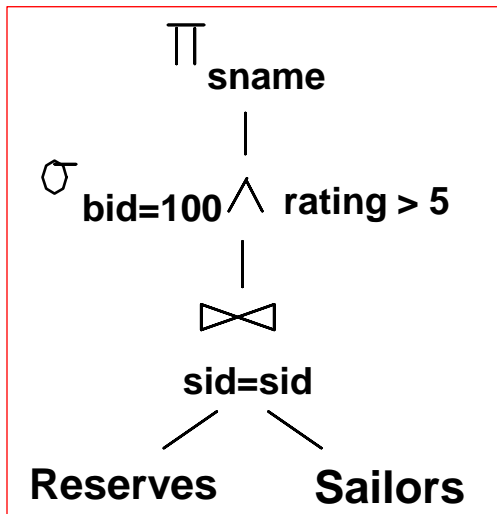


Can you push this projection down?



Ensuring Correctness

- Different RA trees can be created this way
 - Each tree yields a different query execution plan
- Some pushdowns are correct, some are not
 - Humans can distinguish, but DBMSs?
 - Then, how does a DBMS ensure correctness?





Equivalence Rules

- Query optimizer uses **equivalence rules** of RA to find equivalent expressions for a given query
- Two RA expressions are called **equivalent** if they produce the same result on all possible instances of relations
- Equivalence rules allow us to:
 - Push selections and projections ahead of joins
 - Combine selections and cross products into joins
 - Choose different operation/join orders, etc.
- There are many equivalence rules, let's see a small subset of them as examples



Equivalence Rules

- Cascading of selections

- Allows us to combine several selections into one
- Or replace one selection with several smaller selections

$$\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots \sigma_{cn}(R))$$

- Commutation of selections

- Allows us to test selection conditions in different order

$$\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$$



Equivalence Rules

- Joins and cross products are commutative

$$(R \times S) \equiv (S \times R)$$

$$(R \bowtie S) \equiv (S \bowtie R)$$

- Joins and cross products are also associative

$$R \times (S \times T) \equiv (R \times S) \times T$$

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$$

It follows: $R \bowtie (S \bowtie T) \equiv (T \bowtie R) \bowtie S$

This means the query optimizer can choose any order it prefers to perform the joins!



Equivalence Rules

- Selection + cross product -> Join

$$R \bowtie_c T \equiv \sigma_c(R \times S)$$

- Commutation of selection with cross product and join:
 - If **c** appears only in **R** but not **S**

$$\sigma_c(R \times S) \equiv \sigma_c(R) \times S$$

$$\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S$$



Equivalence Rules

- Following commutation of selection w/ cross product:

$$\begin{aligned}\sigma_c(R \times S) &\equiv \sigma_{c1 \wedge c2 \wedge c3}(R \times S) \\ &\equiv \sigma_{c1}(\sigma_{c2}(\sigma_{c3}(R \times S))) \\ &\equiv \sigma_{c1}(\sigma_{c2}(R) \times \sigma_{c3}(S))\end{aligned}$$


This says we can push part of the selection condition **c** ahead of the cross-product

It holds for joins as well!

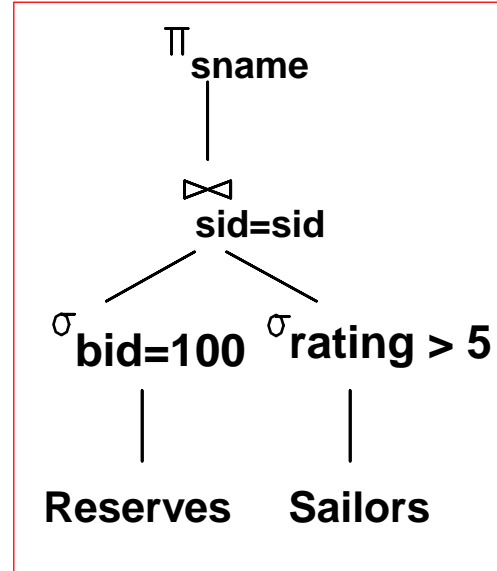


Cost Estimation

- Say that we have a bunch of different query plans (expressed as different RA trees). We want to pick the one with the **lowest cost**, i.e., the one that is **fastest**.
 - How can you **estimate** cost without executing the plan?



One large join +
one selection



Two selections + one smaller join

Is the version with
operator pushdown
always faster?



Cost Estimation

- Cost estimation depends on mainly two factors:
 - How fast can you implement an operator?
 - How fast is selection? Join? Projection?
 - What are the sizes of the operator inputs and outputs?
 - Early selections will reduce how much of the data? 10% or 95%?
 - This is called the **reduction factor** (or **selectivity**)
 - The DBMS can **estimate** reduction factors according to available indexes or summary statistics (e.g., histograms)



Implementing Operators

- We'll study the implementation of the following operators:
 - *Selection* (σ)
 - *Projection* (π)
 - *Join* (\bowtie)
 - *Set operators (union, intersect, difference)*
- Each operator returns a **relation**, therefore an RA tree is a composition of these operators
- Other SQL functionalities (e.g., SUM/MAX/MIN, GROUP BY) can be implemented in linear time
 - We won't cover them – but it's easy to verify



Selection

- Consider the following, basic selection:
 - No conjunctions (**AND**), no disjunctions (**OR**)

```
SELECT *  
FROM Reserves R  
WHERE R.rname = 'Joe'
```

- If a hash/tree index is available, use it
- If not, perform a full-table scan
- How about the case with **multiple conjunctions**?

```
SELECT *  
FROM Reserves R  
WHERE day < 24/3/2015 AND bid=5 AND sid=3
```



Selection

```
SELECT *  
FROM Reserves R  
WHERE day < 24/3/2015 AND bid=5 AND sid=3
```

- If no index is available for **day**, **bid** or **sid**, you must perform a full-table scan
- If an index is available for one of them, evaluate its condition before the others
 - E.g.: B+ tree index on day is available
 - Use it to retrieve tuples which satisfy day < 24/3/2015
 - Among those tuples, check bid=5 and sid=3

Reduce the amount of linear scanning as much as possible



Selection

```
SELECT *  
FROM Reserves R  
WHERE day < 8/9/1994 OR rname="Alice_res"
```

- Case with **disjunctions** (OR):
 - Detailed algorithms exist, we won't cover them
- **Two intuitive remarks:**
 - Use indexes as much as possible
 - E.g., if there's an index on day and an index on rname, use them to individually compute **day < 8/9/1994** condition and **rname="Alice_res"** condition
 - Subsequently, union their results
 - Full-table scan may become unavoidable
 - E.g., index on day exists but no index on rname
 - We can't avoid full-table scan to evaluate **rname="Alice_res"**



Projection

- Consider the following query which implies a projection:

```
SELECT DISTINCT R.sid, R.bid  
FROM Reserves R
```

- How can we evaluate this query?
 - First remove unwanted attributes
 - Then eliminate any duplicate tuples
 - This is the more difficult step, naively, it is $O(N^2)$
- Two strategies for duplicate elimination:
 - Sorting
 - Hashing



Projection w/ Sorting

```
SELECT DISTINCT R.sid, R.bid  
FROM Reserves R
```

- Produce a set of tuples **S** which contains only **sid** and **bid** attributes
- Sort **S**
- Scan the sorted result, compare adjacent tuples, and discard duplicates

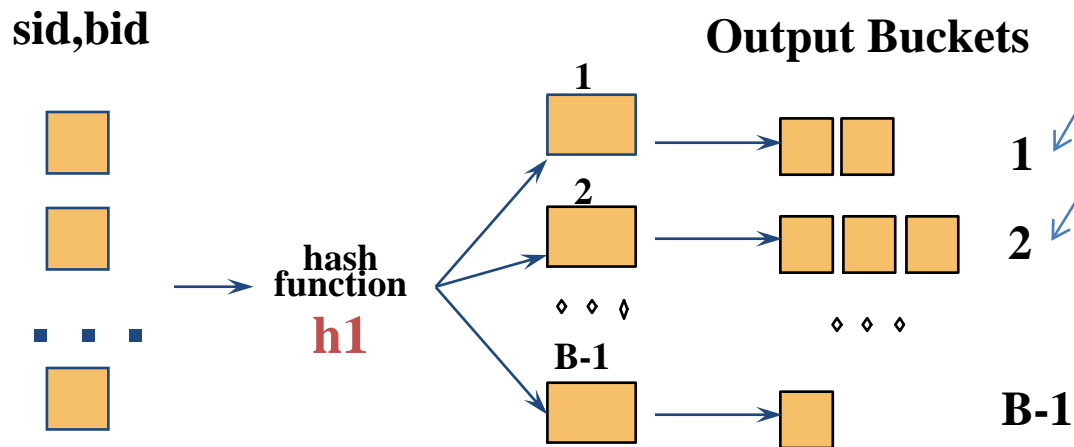
<u>sid</u>	<u>bid</u>
28	103
28	103
31	101
31	101
31	102
58	103



Projection w/ Hashing

```
SELECT DISTINCT R.sid, R.bid  
FROM Reserves R
```

Two tuples that belong to different buckets are guaranteed not to be duplicates





Sorting vs Hashing

- Which approach to use?
- In terms of average-case time complexity (wrt data size), **hashing** is better: **$O(N)$** vs **$O(N \log N)$**
- **Sorting** is better when:
 - There is a high frequency of duplicates
 - Hash function causes too many hash collisions
 - Side benefit of sorting: the result is **sorted**, which can benefit downstream operators (such as joins)



Join

- Consider the following query which implies a join:

```
SELECT *  
FROM Reserves R, Sailors S  
WHERE R.sid = S.sid
```

- We will study three join algorithms:
 - Nested Loop Join $O(|R| * |S|)$
 - Sort-Merge Join $O(|R| * \log|R| + |S| * \log|S|)$
 - If relations are pre-sorted, then complexity is: $O(|R| + |S|)$
 - Hash Join $O(|R| + |S|)$



Nested Loop Join


- Implement join as a nested loop:

```
for each tuple r in R
  for each tuple s in S
    if match: print (r, s)
```


Outer Relation

Inner Relation

R(A,..)



S(A,)





Sort-Merge Join

```
SELECT *  
FROM Reserves R, Sailors S  
WHERE R.sid = S.sid
```

- Individually sort both relations on the join attribute
 - Sort **R** according to **sid**, sort **S** according to **sid**
 - $O(|R| \cdot \log|R| + |S| \cdot \log|S|)$
 - This cost can be avoided if relations have been sorted beforehand (e.g., index or prior projection)
- Scan each relation and merge
 - $O(|R| + |S|)$
- Works only for equality join conditions

Sort-Merge Join: An Example

= ?

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= NO

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= ?

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= YES

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Output the two tuples

Sort-Merge Join: An Example

= ?

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= YES

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= YES

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Output the two tuples

Sort-Merge Join: An Example

= ?

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= NO

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= ?

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: An Example

= YES

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Output the two tuples


Continue the same way!



Hash Join

```
SELECT *  
FROM Reserves R, Sailors S  
WHERE R.sid = S.sid
```

- Create a hash index for **S** using **sid**
 - $O(|S|)$
- For each tuple r in R :
 - $\text{hash}(r)$ and check the index
 - If match, output the two tuples
 - $O(|R|) * O(1)$
- Total: **$O(|R| + |S|)$**





Set Operators

- UNION: $R \cup S$
 - Vertically concatenate R and S
 - Eliminate duplicates (re-use ideas from projection)
- INTERSECTION: $R \cap S$
 - Create hash index for S
 - For each tuple in R : check if it's in that hash index; if so, include it in the output
- DIFFERENCE (EXCEPT): $R - S$
 - Create hash index for S
 - For each tuple in R : check if it's in that hash index; if not, include it in the output