

# COMP 341 Intro to AI

## Making Sequential Decisions under Action Uncertainty



Asst. Prof. Barış Akgün  
Koç University

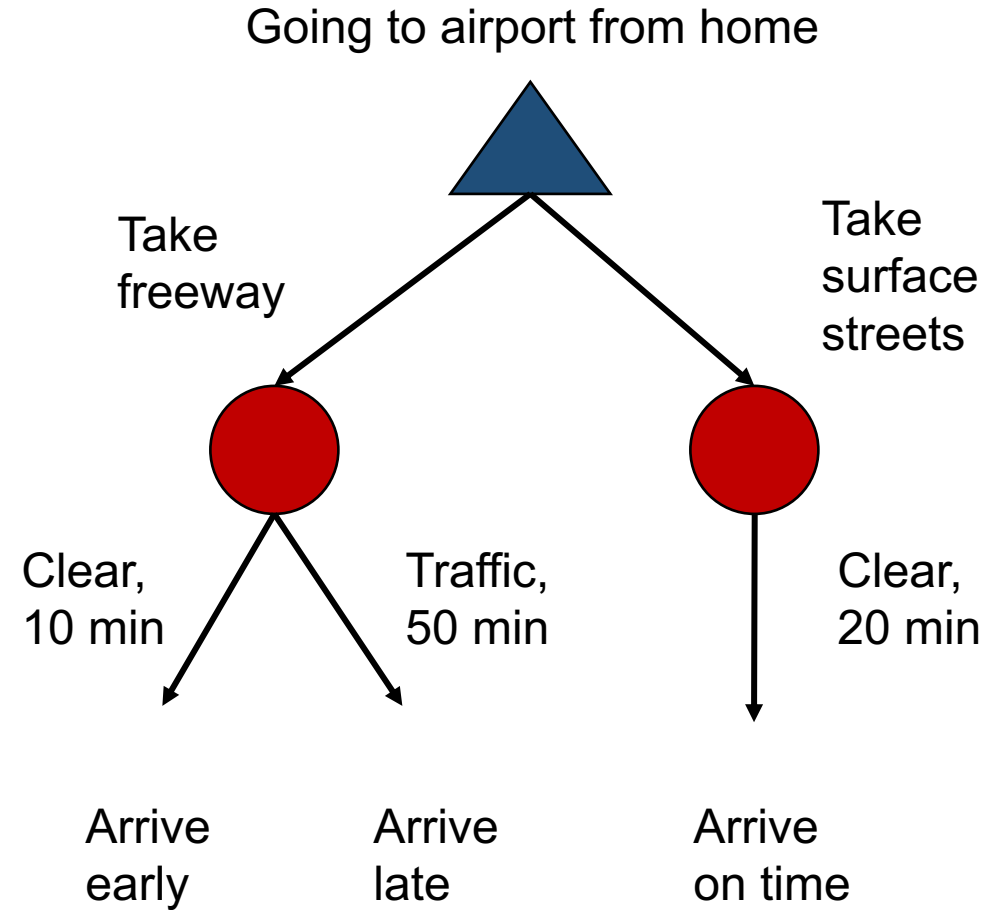
# Maximum Expected Utility

- A rational agent should choose the action which **maximizes its expected utility, given its knowledge**
- Questions:
  - Where do utilities come from?
  - What do utilities represent?
  - Why expected utility?

# Utilities and Unknown Outcomes

- One way has a chance to be better or worse
- How to decide?
- Which would you pick if you are catching a flight?
- Which if you are picking up a friend?

Assigning relative value to outcomes = Utilities




# Agent Rational Decisions

- Representing decisions and maximizing utility:
  - Receive feedback in the form of **rewards**
  - Agent's **utility** is defined by the **reward function**
  - Act to maximize expected rewards over time
  - Can learn how to maximize rewards via Reinforcement Learning
- Examples:
  - Playing a game, reward at the end for winning / losing
  - Vacuuming a house, reward for each piece of dirt picked up
  - Automated taxi, reward for each passenger delivered
- Decision networks: Single (or temporally unrelated) decisions
- RL: Sequential Decisions

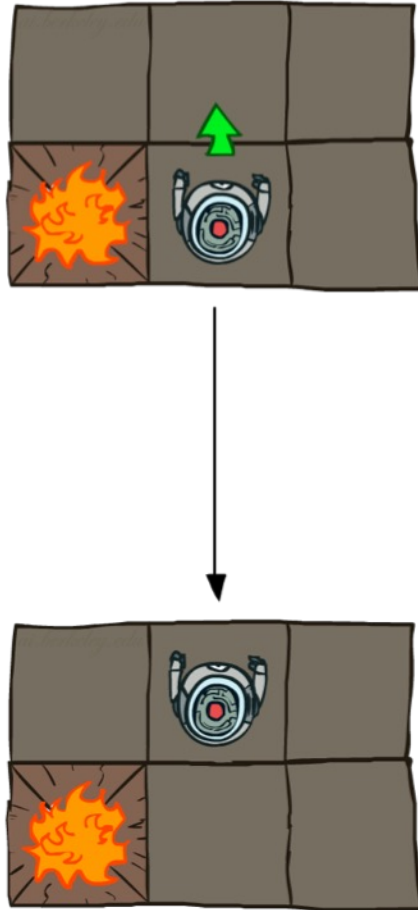
# Example: Grid World

- A maze-like problem
  - The agent lives in a grid
  - Walls block the agent's path
- Noisy movement: actions do not always go as planned
  - 80% of the time, the action North takes the agent North (if there is no wall there)
  - E.g., 10% of the time, North takes the agent West; 10% East
  - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives rewards each time step
  - Small "living" reward each step (can be negative)
  - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards

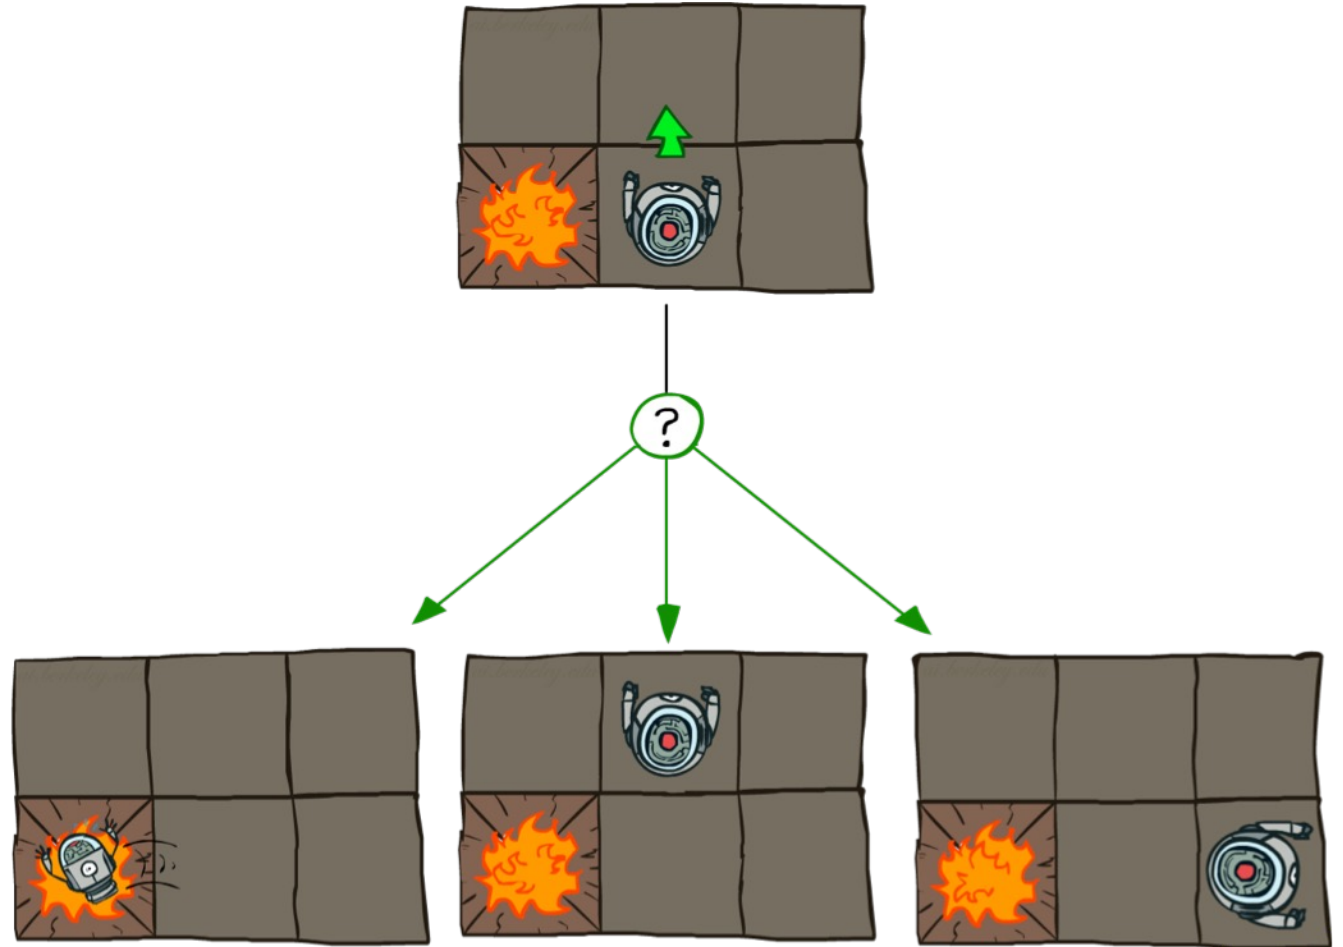


# Grid World Actions

Deterministic Grid World

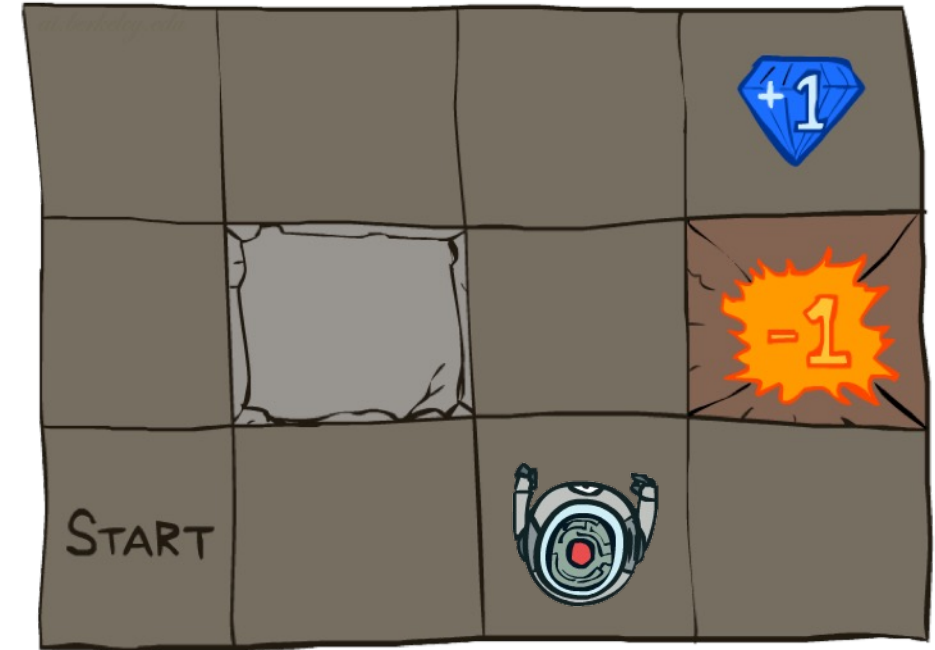


Stochastic Grid World



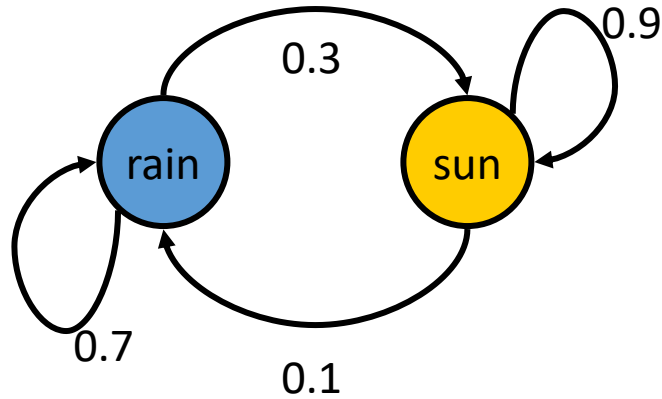
# Grid World

- Why not just search?
  - Stochastic Action Outcomes
- Why not use expectimax and re-plan at each state?
  - A valid idea but...
- Computational burden, repeated states, infinite search tree...
- Markov Decision Processes are a good general way to attack this problem
- The solution will be a sort of “search with memory”



# Recall Markov Chains

- States, transition model and initial distribution

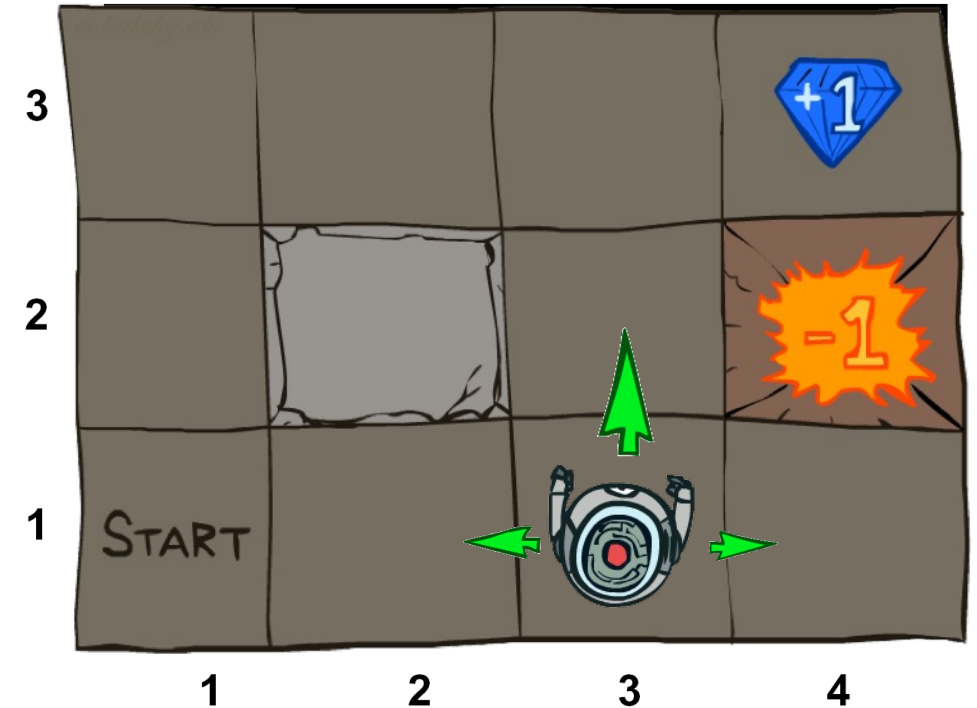


- Assume you have invented a weather machine
  - The states change, with some uncertainty, based on your actions
- You have energy costs but want to keep it sunny
  - You make action decisions based on rewards



# Markov Decision Processes

- An MDP is defined by:
  - A **set of states**  $s \in S$
  - A **set of actions**  $a \in A$
  - A **transition function**  $T(s, a, s')$ 
    - Can be stochastic ( $P(s' | s, a)$ ) or deterministic
    - Also called the model or the dynamics
  - A **reward function**  $R(s, a, s')$ 
    - Sometimes just  $R(s)$  or  $R(s, a)$  or  $R(s')$  (all are equivalent)
  - A **start state**
  - Maybe a **terminal state**
- MDPs are non-deterministic search problems
  - One way to solve them is with expectimax search
  - We'll have a new tool soon



# What is Markovian about MDPs?


- Recall: Markov property generally means that given the present state, past and the future are independent
- For MDPs, action outcomes only depend on the current state

$$\begin{aligned} P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1} = a_{t-1}, \dots, S_0 = s_0) \\ = P(S_{t+1} = s' | S_t = s_t, A_t = a_t) \end{aligned}$$

- Just like search where the successor function only depends on the current state

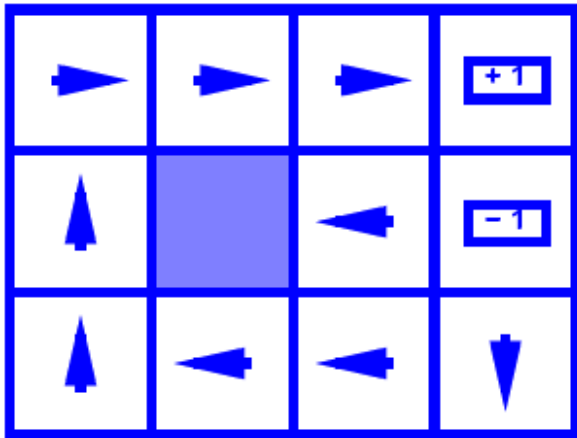
# Policies

- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an optimal policy  $\pi^*: S \rightarrow A$ 
  - A policy  $\pi$  gives an action for each state
  - An optimal policy is one that maximizes expected utility if followed
  - An explicit policy defines a reflex agent
- Expectimax didn't compute entire policies
  - It computed the action for a single state only
  - Doing it at each step would be inefficient and sometimes not possible

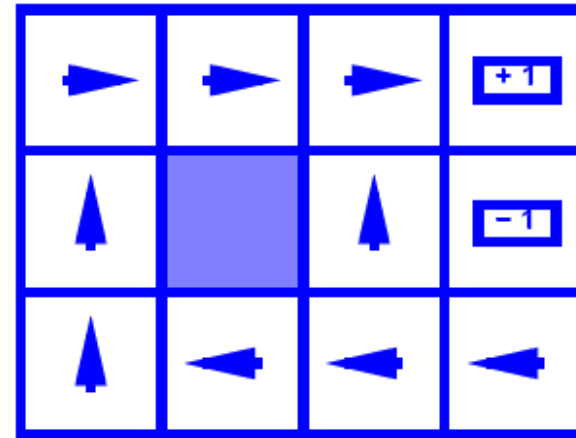


Optimal policy when  $R(s, a, s') = -0.03$   
for all non-terminals  $s$

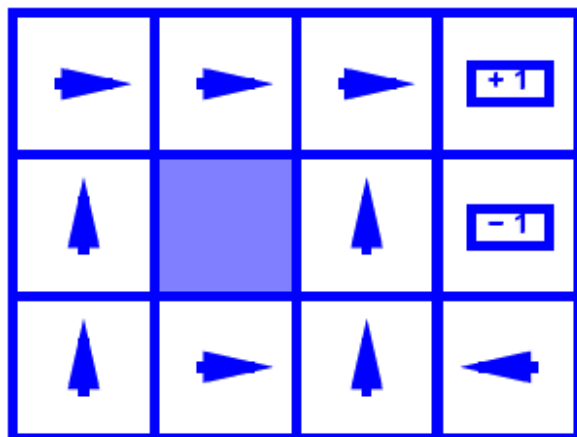
# Optimal Policies



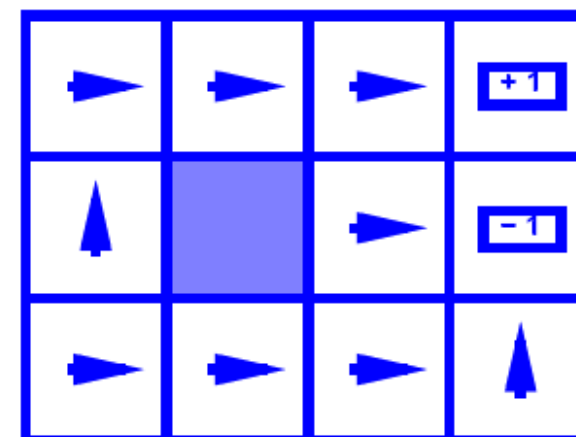
$R(s) = -0.01$



$R(s) = -0.03$



$R(s) = -0.4$



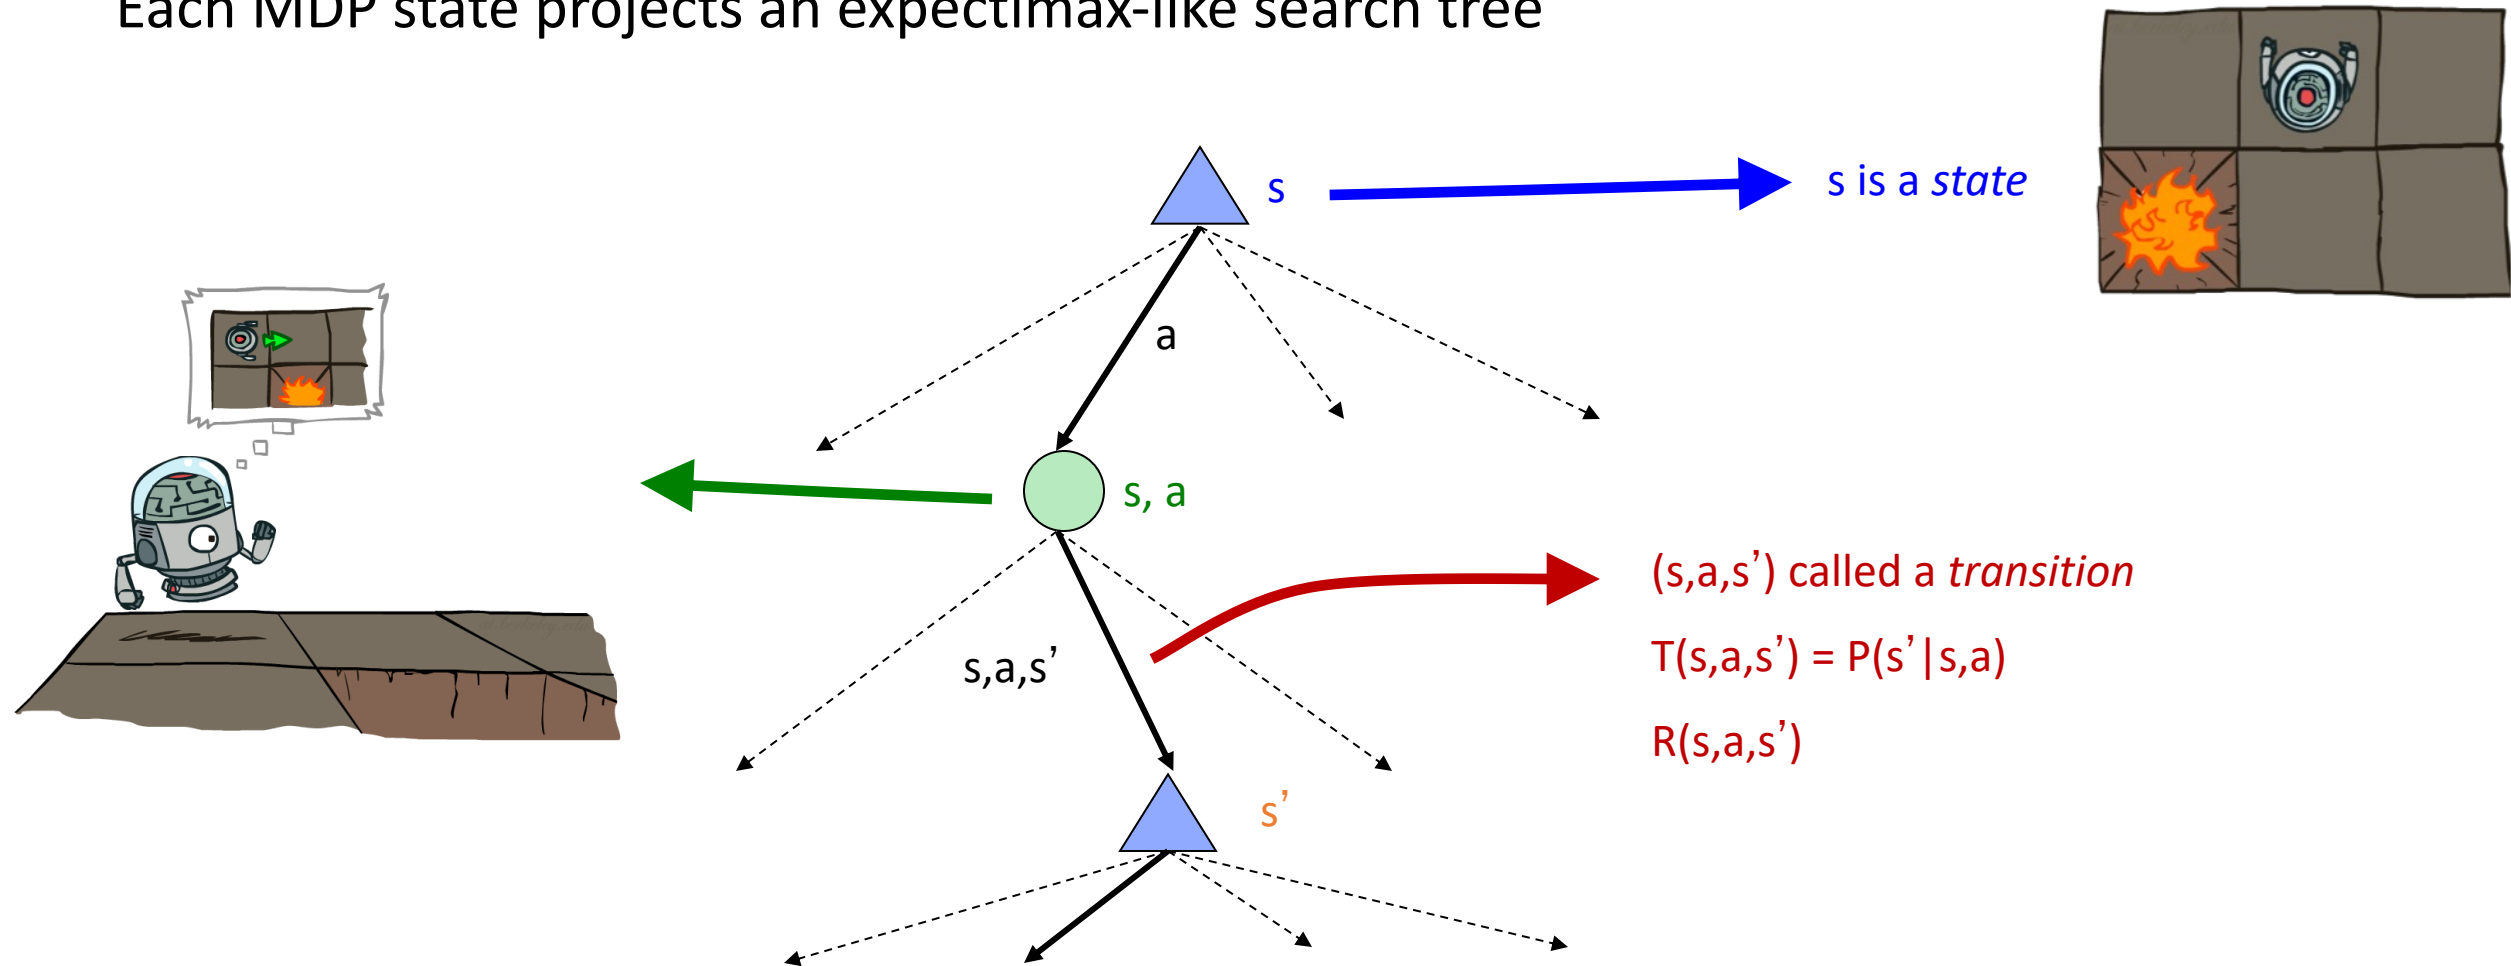
$R(s) = -2.0$

# Utilities and Policies

- Utility: Defined by the reward function
- Solving MDPs: Finding a policy
- Policy: What action to take in each state?
- Optimal Policy: Has the highest expected utility
  
- How to calculate the utility of a policy?

# MDP Search Trees

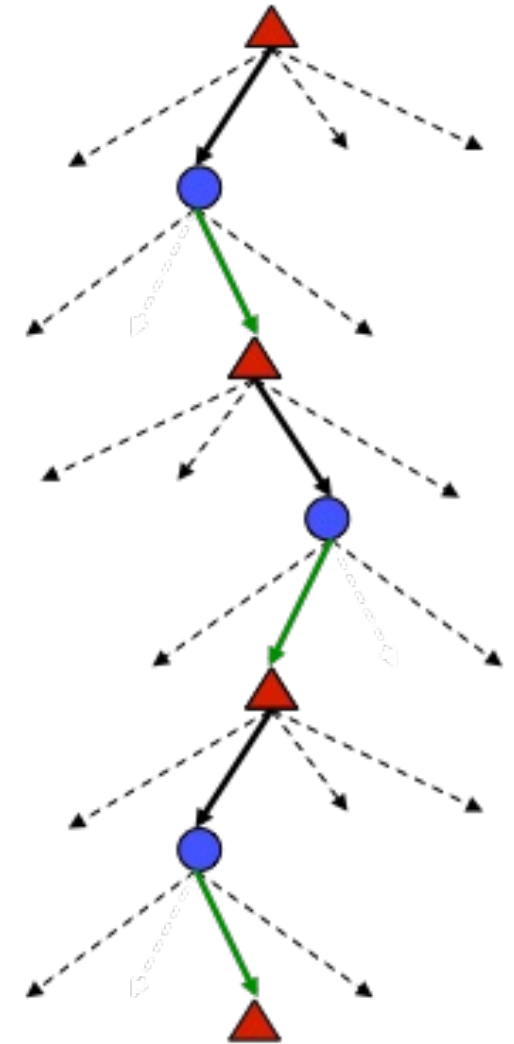
Each MDP state projects an expectimax-like search tree



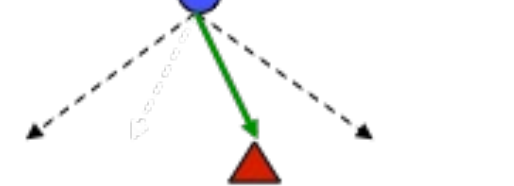
# Solution Horizon

- Finite:
  - Agent must solve the problem in a finite amount of time/steps
  - The state sequences must be finite
  - The right action in a state depends on how much time left
- Infinite:
  - Agent does not have a time/step limit
  - Optimal action depends only on the state

**Non-Stationary**




**Stationary**



# Recap: Defining MDPs


- Markov decision processes:
  - Set of states  $S$
  - Start state  $s_0$
  - Set of actions  $A$
  - Transitions  $P(s' | s, a)$  (or  $T(s, a, s')$ )
  - Rewards  $R(s, a, s')$  (and discount  $\gamma$ )
- MDP quantities so far:
  - Policy = Choice of action for each state
  - Utility = sum of (discounted) rewards or average rewards





# Utilities of State Sequences

- A path in a tree gives a sequence
- We receive rewards at each state!
- What is the utility of a sequence?
  - Sum of rewards?
  - Average rewards?
  - Reward now is better than later?
  - What about infinite sequences?
- Idea: Sum of discounted rewards




# Discounting

- How to discount?
  - Each time we descend a level, we multiply in the discount once
- Why discount?
  - Sooner rewards probably do have higher utility than later rewards
  - Also helps our algorithms converge

- Example: discount of 0.5

$$U([1,2,3]) = 1 \cdot 1 + 0.5 \cdot 2 + 0.25 \cdot 3$$


$$U([1,2,3]) < U([3,2,1])$$



# Stationary Preferences

- Theorem: if we assume **stationary preferences**:

$$\begin{aligned} [a_1, a_2, \dots] &\succ [b_1, b_2, \dots] \\ &\Updownarrow \\ [r, a_1, a_2, \dots] &\succ [r, b_1, b_2, \dots] \end{aligned}$$



- Then: there are only two ways to define utilities

- Additive utility:  $U([r_0, r_1, r_2, \dots]) = r_0 + r_1 + r_2 + \dots$
- Discounted utility:  $U([r_0, r_1, r_2, \dots]) = r_0 + \gamma r_1 + \gamma^2 r_2 \dots$

# Infinite Utilities?!


- Problem: What if the game lasts forever? Do we get infinite rewards?

- Solutions:

- Finite horizon: (like depth-limited search)
  - Terminate episodes after a fixed T steps (e.g. life)
  - Gives non-stationary policies ( $\pi$  depends on time left)
- Discounting: use  $0 < \gamma < 1$

$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max} / (1 - \gamma)$$

- Smaller  $\gamma$  means smaller “horizon” – shorter term focus
  - The discount factor favors “shorter” solutions (unless there is a high living reward!)
- Absorbing state: guarantee that for every policy, a terminal state will eventually be reached



# Calculating Policies

- MDPs end up being a good representation for many real-world problems.
- Given an MDP description, we will see several algorithms for solving for the optimal policy
- Two General Class of Algorithms:
  - Value Iteration
  - Policy Iteration

# Optimal Quantities

- The value (utility) of a state  $s$ :


$V^*(s)$  = expected utility starting in  $s$  and acting optimally

- The value (utility) of a q-state  $(s,a)$ :


$Q^*(s,a)$  = expected utility starting out having taken action  $a$  from state  $s$  and (thereafter) acting optimally

- The optimal policy:

$\pi^*(s)$  = optimal action from state  $s$



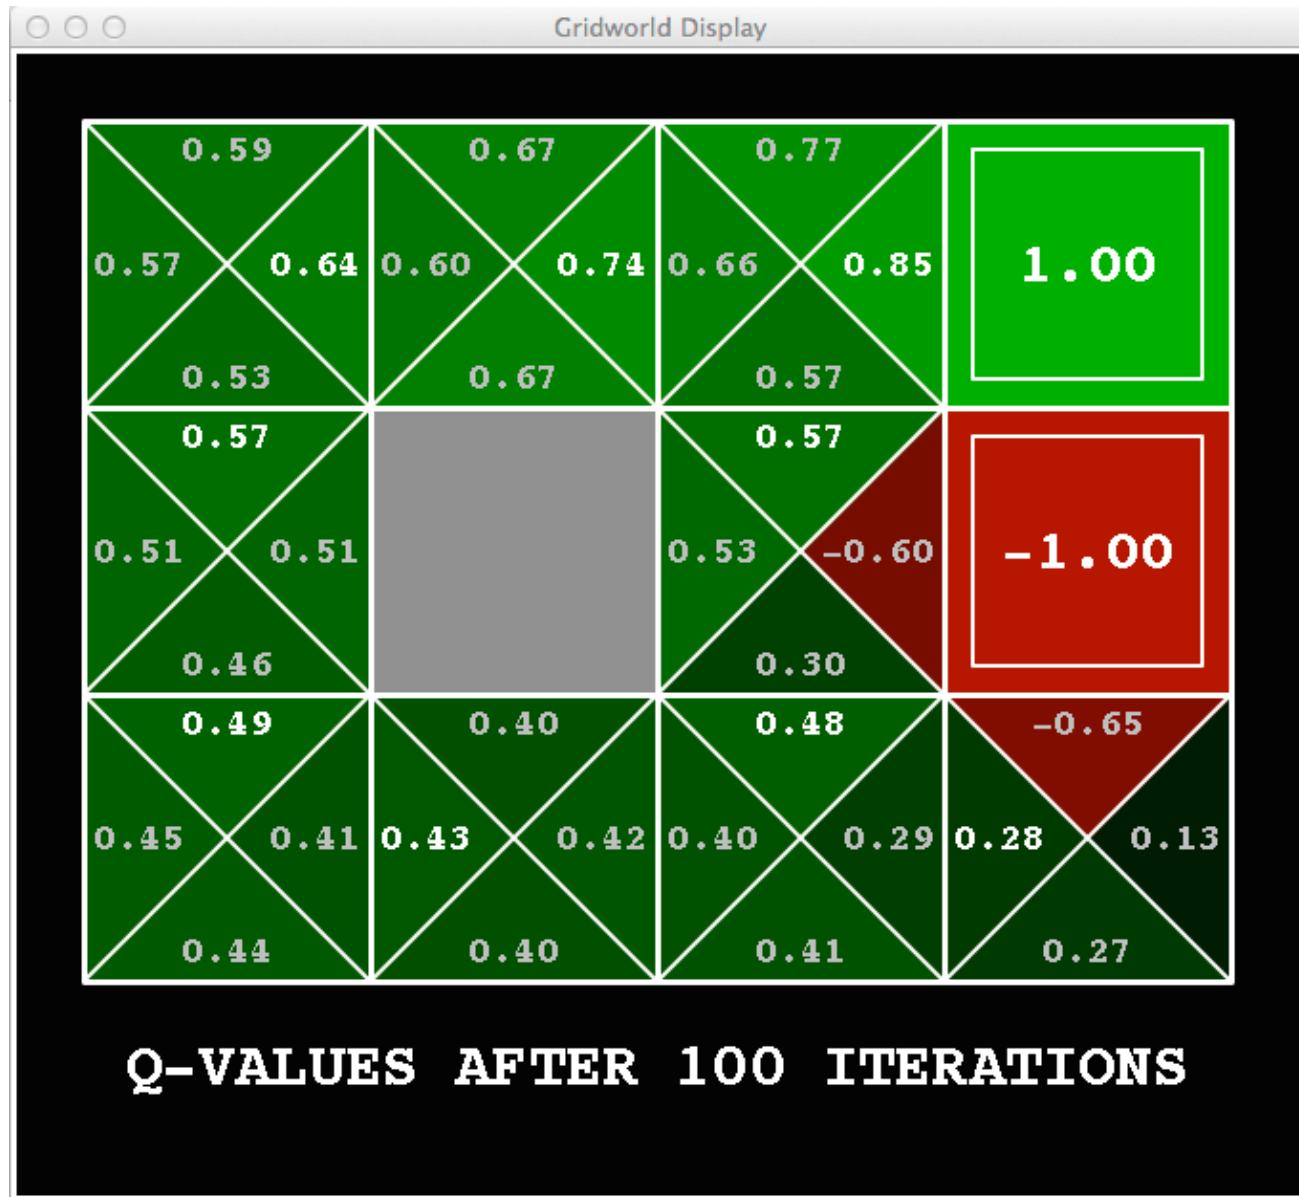
# Snapshot of Demo – Gridworld Values



Value of a state is a more “long term” quantity whereas reward is an immediate quantity.

Noise = 0.2  
Discount = 0.9  
Living reward = 0

# Snapshot of Demo – Gridworld Q Values



Q-Value is also a long term quantity

Noise = 0.2  
Discount = 0.9  
Living reward = 0



# A Mathematical Remark

- Given a potentially infinite state sequence:

$$\sigma = \{s_0, s_1, s_2, \dots\}$$

- The utility of the path is

$$U(\sigma) = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots = r_0 + \gamma U(\sigma'), \sigma' = \{s_1, s_2, \dots\}$$

- Then the expected utility, or the value, of the state  $s_0$  given a policy  $\pi$  is:

$$V(s_0) = \sum_{s'} P(s'|s_0, \pi(s_0)) (R(s_0, \pi(s_0), s') + \gamma V(s'))$$

- or if  $R(s, a, s') = r_0$

$$V(s_0) = r_1 + \gamma \sum_{s'} P(s'|s_0, \pi(s_0)) V(s')$$

# Bellman Equations

- $V^*(s)$ : Expected utility starting in  $s$  and acting optimally:

$$V^*(s) = \sum_{s'} (P(s'|s, \pi^*(s)) (R(s, \pi^*(s), s') + \gamma V^*(s')))$$

- $\pi^*(s)$ : For a single state, pick the action that maximizes the expected utility

$$\pi^*(s) = \arg_a \max(\sum_{s'} (P(s'|s, a) (R(s, a, s') + \gamma V^*(s'))))$$

- $Q^*(s, a)$ : Expected utility of starting in  $s$  with action  $a$  and acting optimally

$$Q^*(s, a) = \sum_{s'} (P(s'|s, a) (R(s, a, s') + \gamma V^*(s')))$$

- These are related:

$$V^*(s) = \max_a (Q^*(s, a))$$

$$V^*(s) = \max_a (\sum_{s'} (P(s'|s, a) (R(s, a, s') + \gamma V^*(s'))))$$

# Some Observations

$$V^*(s) = \max_a \left( \sum_{s'} \left( P(s'|s, a) \left( \overset{\text{Immediate reward}}{\boxed{R(s, a, s')}} + \gamma \overset{\text{Discounted neighbor value}}{\boxed{V^*(s')}} \right) \right) \right)$$

- If we calculate the values of states, we immediately get the policy
- These quantities are defined recursively
- The value/utility of a state is related to its neighboring states
- Why not iteratively calculate the values of states, using current estimates of the neighboring states?

# Bellman Updates

- General Case:

$$V_{t+1}(s) = \max_a \left( \sum_{s'} \left( P(s'|s, a) (R(s, a, s') + \gamma V_t(s')) \right) \right)$$

- If the rewards are only for states:

$$V_{t+1}(s) = R(s) + \gamma \max_a \left( \sum_{s'} (P(s'|s, a) V_t(s')) \right)$$

- The idea is to find the unknown function  $V(s)$  given the MDP
- Note that the policy is not explicitly calculated

# Value Iteration

- Start with  $V_0(s) = 0$
- Use the Bellman update to recursively calculate  $V(s)$   
$$V_{t+1}(s) = \max_a \left( \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma V_t(s')) \right)$$
  - At each step, go over all the states
- Repeat until convergence
- Complexity of one iteration:  $O(|S|^2 |A|)$  - Why?
- Theorem:  $V(s)$  will converge to optimal values,  $V^*(s)$ 
  - The book has a proof – **read it!**
- Note: The resulting policy may converge long before the values! – Why?

# Value Iteration

**function** VALUE-ITERATION( $mdp, \epsilon$ ) **returns** a utility function

**inputs:**  $mdp$ , an MDP with states  $S$ , transition model  $T$ , reward function  $R$ , discount  $\gamma$   
 $\epsilon$ , the maximum error allowed in the utility of any state

**local variables:**  $U, U'$ , vectors of utilities for states in  $S$ , initially zero  
 $\delta$ , the maximum change in the utility of any state in an iteration

**repeat**

$U \leftarrow U'; \delta \leftarrow 0$

**for each** state  $s$  **in**  $S$  **do**

$$U'[s] \leftarrow R[s] + \gamma \max_a \sum_{s'} T(s, a, s') U[s']$$

**if**  $|U'[s] - U[s]| > \delta$  **then**  $\delta \leftarrow |U'[s] - U[s]|$

**until**  $\delta < \epsilon(1 - \gamma)/\gamma$

**return**  $U$


Bellman update equation  
Could also use the general version here

Comes from the convergence analysis!

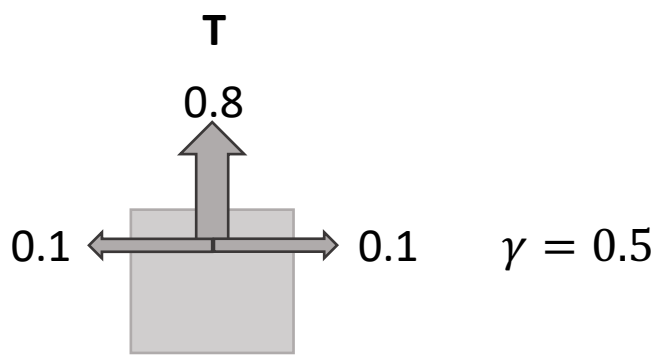
# Exercise: Value Iteration

- MDP:
  - States: Represented by the grid
  - Actions: Up, Down, Left, Right
  - Rewards: As seen on the grid
  - Transition Model: 80-10-10 (see the figure)
  - Discount:  $\gamma = 0.5$
  - (2,2) is a terminal state
- Question: Starting from  $V(s_{11}) = V(s_{12}) = V(s_{21}) = 0.1$  and  $V(s_{22}) = 1.0$ , do one step of value iteration

-0.04	+1.0
-0.04	-0.04



R	
-0.04	+1.0
-0.04	-0.04



$V_0$	
0.1	+1.0
0.1	0.1

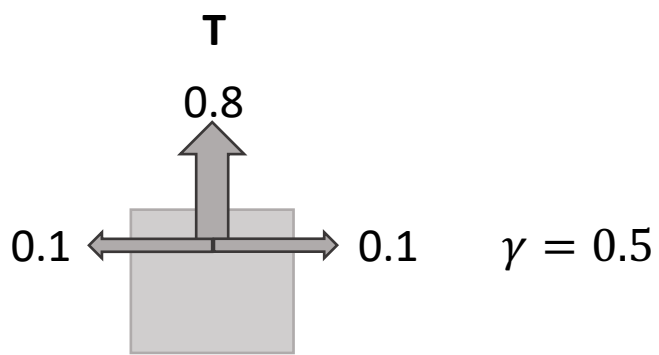
$$V_{t+1}(s) = R(s) + \gamma \max_a \left( \sum_{s'} (P(s'|s, a) V_t(s')) \right)$$

$$V_1(s_{21}) = -0.04 + 0.5 \max_a (P(s_{11}|s_{21}, UP)V_0(s_{11}) + P(s_{21}|s_{21}, UP)V_0(s_{21}) + P(s_{22}|s_{21}, UP)V_0(s_{22}), \\ P(s_{11}|s_{21}, DN)V_0(s_{11}) + P(s_{21}|s_{21}, DN)V_0(s_{21}) + P(s_{22}|s_{21}, DN)V_0(s_{22}), \\ P(s_{11}|s_{21}, LT)V_0(s_{11}) + P(s_{21}|s_{21}, LT)V_0(s_{21}) + P(s_{22}|s_{21}, LT)V_0(s_{22}), \\ P(s_{11}|s_{21}, RT)V_0(s_{11}) + P(s_{21}|s_{21}, RT)V_0(s_{21}) + P(s_{22}|s_{21}, RT)V_0(s_{22}))$$

$$V_1(s_{21}) = -0.04 + 0.5 \max_a (a = UP: 0.1 \cdot 0.1 + 0.1 \cdot 0.1 + 0.8 \cdot 1.0 = 0.82, \\ a = DN: 0.1 \cdot 0.1 + 0.9 \cdot 0.1 + 0.0 \cdot 1.0 = 0.1, \\ a = LT: 0.8 \cdot 0.1 + 0.1 \cdot 0.1 + 0.1 \cdot 1.0 = 0.19, \\ a = RT: 0.0 \cdot 0.1 + 0.9 \cdot 0.1 + 0.1 \cdot 1.0 = 0.19) \\ = -0.04 + 0.5 \cdot 0.82 = 0.37$$



R	
-0.04	+1.0
-0.04	-0.04



$V_0$	
0.1	+1.0
0.1	0.1


$$V_{t+1}(s) = R(s) + \gamma \max_a \left( \sum_{s'} (P(s'|s, a) V_t(s')) \right)$$

$$V_1(s_{11}) = -0.04 + 0.5 \max_a (P(s_{12}|s_{11}, UP)V_0(s_{12}) + P(s_{21}|s_{11}, UP)V_0(s_{21}) + P(s_{11}|s_{11}, UP)V_0(s_{11}), \\ P(s_{12}|s_{11}, DN)V_0(s_{12}) + P(s_{21}|s_{11}, DN)V_0(s_{21}) + P(s_{11}|s_{11}, DN)V_0(s_{11}), \\ P(s_{12}|s_{11}, LT)V_0(s_{12}) + P(s_{21}|s_{11}, LT)V_0(s_{21}) + P(s_{11}|s_{11}, LT)V_0(s_{11}), \\ P(s_{12}|s_{11}, RT)V_0(s_{12}) + P(s_{21}|s_{11}, RT)V_0(s_{21}) + P(s_{11}|s_{11}, RT)V_0(s_{11}))$$

$$V_1(s_{11}) = -0.04 + 0.5 \max_a (a = UP: 0.8 \cdot 0.1 + 0.1 \cdot 0.1 + 0.1 \cdot 0.1 = 0.1, \\ a = DN: 0.0 \cdot 0.1 + 0.1 \cdot 0.1 + 0.9 \cdot 0.1 = 0.1, \\ a = LT: 0.1 \cdot 0.1 + 0.0 \cdot 0.1 + 0.9 \cdot 0.1 = 0.1, \\ a = RT: 0.8 \cdot 0.1 + 0.1 \cdot 0.1 + 0.1 \cdot 0.1 = 0.1)$$

$$= 0.01$$

R	
-0.04	+1.0
-0.04	-0.04



<b>V<sub>0</sub></b>	
0.1	+1.0
0.1	0.1

$$V_{t+1}(s) = R(s) + \gamma \max_a \left( \sum_{s'} (P(s'|s, a) V_t(s')) \right)$$

$$V_1(s_{12}) = -0.04 + 0.5 \max_a (a = UP: 0.0 \cdot 0.1 + 0.9 \cdot 0.1 + 0.1 \cdot 1.0 = 0.19,$$

$$a = DN: 0.8 \cdot 0.1 + 0.1 \cdot 0.1 + 0.1 \cdot 1.0 = 0.1,$$

$$a = LT: 0.1 \cdot 0.1 + 0.9 \cdot 0.1 + 0.0 \cdot 1.0 = 0.1,$$

$$a = RT: 0.1 \cdot 0.1 + 0.1 \cdot 0.1 + 0.8 \cdot 1.0 = 0.82)$$

$$= -0.04 + 0.5 \cdot 0.82 = 0.37$$

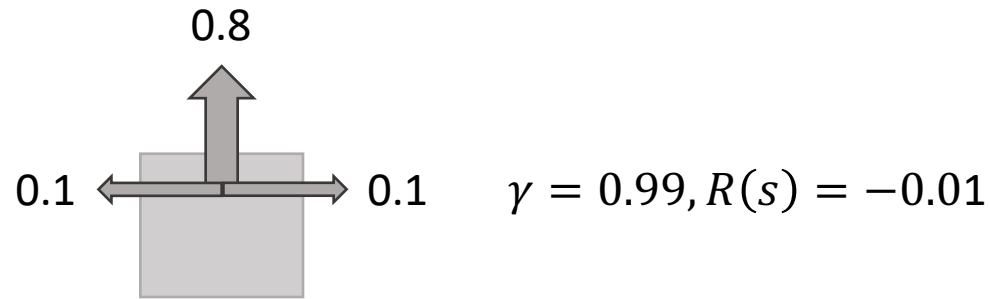
<b>V<sub>1</sub></b>	
0.37	+1.0
0.01	0.37

<b>V<sub>5</sub></b>	
0.376	+1.0
0.12	0.376

# Value Function to Policy (Policy Extraction)

$$\pi^*(s) = \arg_a \max(\sum_{s'} (P(s'|s, a)(R(s, a, s') + \gamma V^*(s'))))$$

0.90	0.93	0.95	+1.0
0.88		0.79	-1.0
0.85	0.83	0.81	0.13



$$\arg_a \max( a = UP: 0.8 \cdot 0.79 + 0.1 \cdot 0.83 + 0.1 \cdot 0.13$$

$$a = DN: 0.8 \cdot 0.81 + 0.1 \cdot 0.83 + 0.1 \cdot 0.13$$

$$a = LT: 0.8 \cdot 0.83 + 0.1 \cdot 0.81 + 0.1 \cdot 0.79$$


$$a = RT: 0.8 \cdot 0.13 + 0.1 \cdot 0.81 + 0.1 \cdot 0.79$$

With some manipulation  
( $\gamma$  and  $R(s)$  same for all)

# Value Function to Policy (Policy Extraction)

$$\pi^*(s) = \arg_a \max(\sum_{s'} (P(s'|s, a)(R(s, a, s') + \gamma V^*(s'))))$$

0.90	0.93	0.95	+1.0
0.88		0.79	-1.0
0.85	0.83	0.81	0.13



$$\arg_a \max( a = UP: 0.8 \cdot 0.95 + 0.1 \cdot 0.79 - 0.1 \cdot 1.0 = 0.739$$

$$a = DN: 0.8 \cdot 0.81 + 0.1 \cdot 0.83 - 0.1 \cdot 1.0$$

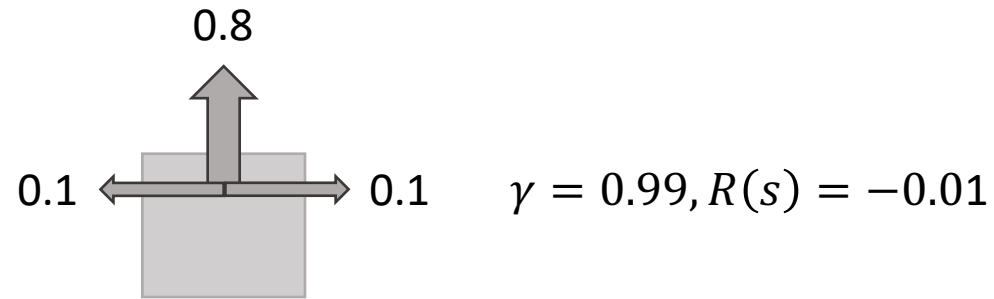
$$a = LT: 0.8 \cdot 0.79 + 0.1 \cdot 0.81 + 0.1 \cdot 0.95 = 0.808$$

$$a = RT: -0.8 \cdot 1.0 + 0.1 \cdot 0.81 + 0.1 \cdot 0.95$$

# Exercise at Home

$$\pi^*(s) = \arg_a \max(\sum_{s'} (P(s'|s, a)(R(s, a, s') + \gamma V^*(s'))))$$

0.90	0.93	0.95	+1.0
0.88		0.79	-1.0
0.85	0.83	0.81	0.13




→	→	→	+1.0
↑		←	-1.0
←	←	←	?

# Problems with Value Iteration

$$V_{t+1}(s) = \max_a \left( \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma V_t(s')) \right)$$

1. It's slow,  $O(|S|^2|A|)$  per step
2. The “max” at each state rarely changes after a certain point
3. As a result, the policy converges before the value function!



# Policy Based Methods


- In value iteration we have used (drum roll) values to get the optimal policy
- Policy iteration is an alternate way to get the optimal policy
- Idea:
  - Start with an initial policy
  - Evaluate the policy (Policy Evaluation)
  - Improve the policy (Policy Improvement)
  - Stop when the policy does not change

# Policy Evaluation

- What is the value function for a given policy  $\pi$ ?
- In other words, the expected total discounted rewards starting in  $s$  and following  $\pi$ ,  $V^\pi(s) = ?$
- Note that  $\pi$  does not have to be optimal

$$V^\pi(s) = \sum_{s'} (P(s'|s, \pi(s)) (R(s, \pi(s), s') + \gamma V^\pi(s')))$$


- Like value iteration calculations but simpler since we are not considering all possible actions






# Example: Policy Evaluation

Always Go Right



Always Go Forward




# Example: Policy Evaluation

Always Go Right



Always Go Forward



# Policy Evaluation

$$V^\pi(s) = \sum_{s'} (P(s'|s, \pi(s)) (R(s, \pi(s), s') + \gamma V^\pi(s')))$$

- How to calculate  $V^\pi(s) = ?$
- Idea 1: Calculate it iteratively, like before ( $O(|S|^2)$  per step)

$$V_0^\pi(s) = 0$$

$$V_{t+1}^\pi(s) = \sum_{s'} (P(s'|s, \pi(s)) (R(s, \pi(s), s') + \gamma V_t^\pi(s')))$$

- Idea 2: Without the max, the Equations are just a linear system!
  - Plug it into your favorite linear system solver!
  - Involves a matrix inversion:  $O(|S|^3)$  overall

# Policy Improvement

- Just do policy extraction!

$$\pi_{t+1}(s) = \arg_a \max \left( \sum_{s'} (P(s'|s, a) (R(s, a, s') + \gamma V^{\pi_t}(s'))) \right)$$

# Policy Iteration Summary

- Starting with an initial policy
  - Set  $V^\pi(s) = 0$ , calculate an initial policy
  - Alternatively start with a random policy
- Evaluation: For the current fixed policy calculate the values  $V^\pi(s)$

$$V^\pi(s) = \sum_{s'} (P(s'|s, \pi(s)) (R(s, \pi(s), s') + \gamma V^\pi(s')))$$

- Either through iteration
  - Or through the linear solution
- Improvement: For fixed values, extract a new policy
$$\pi_{t+1}(s) = \arg_a \max(\sum_{s'} (P(s'|s, a) (R(s, a, s') + \gamma V^{\pi_t}(s'))))$$
- Stop when there is no change

# Comparison

- Both value iteration and policy iteration compute the same thing (optimal values)
- In value iteration:
  - Every iteration updates both the values and (implicitly) the policy
  - We don't track the policy, but taking the max over actions implicitly recomputes it
- In policy iteration:
  - We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
  - After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
  - The new policy will be better (or we're done)
- Both are dynamic programs for solving MDPs

# Summary: MDP Algorithms

- So, you want to....
  - Compute optimal values: use value iteration or policy iteration
  - Compute values for a particular policy: use policy evaluation
  - Turn your values into a policy: use policy extraction
- These all look the same!
  - They basically are – they are all variations of Bellman updates
  - They differ only in whether we plug in a fixed policy or max over actions

# Further Reading

- There are other, more efficient methods to calculate values:
  - Asynchronous Value Iteration
  - Modified Policy Iteration
- Demonstration seeded policy iteration
  - Observe a human or another agent solving the same MDP
  - Use their actions as an initial policy
  - Might not cover the entire state space
- Inverse Problem:
  - Given an optimal policy, what are the rewards/values?
- Policy Search
  - If the state and action spaces are large, improve the policy using local search methods
  - Utility calculations are local
- Partially Observable MDPs (POMDPs)
  - We cannot observe the state directly
  - Use a Dynamic BN to represent probability distribution over states