



COMP 341 Intro to AI Reinforcement Learning



Asst. Prof. Barış Akgün
Koç University

MDPs


- Markov decision processes:
 - Set of states S
 - Set of actions A
 - Rewards $R(s,a,s')$ (and discount γ)
 - Transition Model ($P(s' | s,a)$ or $T(s,a,s')$)
 - Start state distribution
 - Sometimes terminal states
- Quantities:
 - Value function
 - Policies
 - Q-values (value for a state action pair)



Calculating Policies

- Goal is to find an optimal policy given the MDP
- Methods:
 - Value Iteration
 - Policy Iteration
- All the calculations are offline!
- Doing this online is called Reinforcement Learning

Reinforcement Learning



- Basic idea:
 - Receive feedback in the form of **rewards**
 - Agent's utility is defined by the reward function
 - Must (learn to) act so as to **maximize expected rewards**
 - All learning is based on observed samples of outcomes!

Learning to Walk



Initial




Final

Reinforcement Learning

- Still assume a Markov decision process (MDP):
 - A set of states $s \in S$
 - A set of actions (per state) A
 - A model $T(s,a,s')$ or $P(s' | s,a)$
 - A reward function $R(s,a,s')$
- Still looking for a policy $\pi(s)$
- New twist: don't know T or R (only receive rewards online) or both
 - i.e. we don't know which states are good or what the actions do
 - Must actually try actions and states out to learn

Offline (MDPs) vs. Online (RL)



Offline Solution



Online Learning

Why Reinforcement Learning?

- It's not that we would prefer to “risk” the agent, but we are often not left a choice
 - Models cannot be approximated reliably (e.g. a spider robot walking on snow)
 - Reward may be unknown (e.g. showing ads to maximize profit)
- It turns out that reinforcement learning is the only feasible approach to train an agent to perform at high levels in complex and uncertain domains
- Also note that we often have either the transition model or the reward function (but not both)
- Even in some cases where we have the models, offline learning borrows some ideas from online methods
- In this case the learning happens online “in simulation”

Motivation

- In the agent way of thinking, rewards can be considered as a sensory input to the agent
- However, the agent must recognize the reward as reward!
- Animals are hardwired to recognize pain and hunger as bad reward and pleasure and food intake as positive reward
- Food can be used to teach tricks to animals!

Motivation


- Animals are not born with “physics” knowledge but somethings are built-in
- They learn how to move their bodies starting from innate knowledge
 - E.g., babies start from reflexes
- They re-learn if their body changes
 - E.g., teenagers’ voice cracks and they are clumsy growing up
 - E.g., when they are injured
- They also learn to manipulate objects through trial-error and rewards

Model-Based Learning

- Model-Based Idea:
 - Learn an approximate model based on experiences (supervised learning)
 - Solve for values as if the learned model were correct
- Step 1: Learn empirical MDP model
 - Count outcomes s' for each s, a
 - Normalize to give an estimate of $P(s'|s, a)$
 - Discover each $R(s, a, s')$ when we experience (s, a, s')
- Step 2: Solve the learned MDP
 - We already know the methods!

Example: Model Based Learning

Input Policy π



Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10

Learned Model

$$\hat{T}(s, a, s')$$

T(B, east, C) = 1.00
T(C, east, D) = 0.75
T(C, east, A) = 0.25
...

$$\hat{R}(s, a, s')$$

R(B, east, C) = -1
R(C, east, D) = -1
R(D, exit, x) = +10
...

Example: Expected Age

Goal: Compute expected age of a group of people

Known $P(A)$

$$E[A] = \sum_a P(a) \cdot a = 0.35 \times 20 + \dots$$

Without $P(A)$, instead collect samples $[a_1, a_2, \dots, a_N]$

Unknown $P(A)$: “Model Based”

$$\hat{P}(a) = \frac{\text{num}(a)}{N}$$

$$E[A] \approx \sum_a \hat{P}(a) \cdot a$$

Why does this work? Because eventually you learn the right model.

Unknown $P(A)$: “Model Free”

$$E[A] \approx \frac{1}{N} \sum_i a_i$$


Why does this work? Because samples appear with the right frequencies.

Model-Free Learning

- Our main goal is to learn a policy
- In both value iteration and policy iteration we learned the policy through calculating the values of the states
- It turns out that we can by-pass learning models and directly learn the values of states in an online fashion
- Model-Based vs Model-Free choice heavily depends on the problem, it is difficult to say in general which one is better
- We will have a discussion after learning more about RL

Passive Reinforcement Learning

- Simplified task: policy evaluation
 - Input: a fixed policy $\pi(s)$
 - You don't know the transitions $T(s,a,s')$
 - You don't know the rewards $R(s,a,s')$
 - **Goal: learn the state values**
- In this case:
 - Learner is “along for the ride”
 - No choice about what actions to take
 - Just execute the policy and learn from experience
 - This is NOT offline planning! You actually take actions in the world.




Direct Utility Estimation

- Goal: Compute values for each state under π
- Idea: Average together observed sample values
 - Act according to π
 - For each visited state, write down what the sum of discounted rewards turned out to be
 - Average the sample values
- This is called direct evaluation

Example: Direct Evaluation

Input Policy π



Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10

Output Values

	-10 A	
+8 B	+4 C	+10 D
	-2 E	

Problems with Direct Evaluation

- What's good about direct evaluation?
 - It's easy to understand
 - It doesn't require any knowledge of T , R
 - It eventually computes the correct average values, using just sample transitions
- What bad about it?
 - It does not use information about state connections
 - Each state must be learned separately
 - So, it takes a long time to learn

Output Values

	-10 A	
+8 B	+4 C	+10 D
	-2 E	

If B and E both go to C under this policy, how can their values be different?


Why not use Policy Evaluation?

- Simplified Bellman updates calculate V for a fixed policy:
 - Each round, replace V with a one-step-look-ahead layer over V

$$V_0^\pi(s) = 0$$

$$V_{k+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$$

- This approach fully exploited the connections between the states
 - Unfortunately, we need T and R to do it!
- Key question: how can we do this update to V without knowing T and R ?
 - In other words, how to we take a weighted average without knowing the weights?



Model-Based Case

- Act according to a fixed policy
- Update the models with the observed episodes
- Do policy evaluation
 - Value iteration OR
 - Solving the linear system
- Called the “Passive Adaptive Dynamic Programming” in the AIMA book
- Note: Anything not labelled Model-Based is Model-Free in these slides

Sample Based Policy Evaluation

- We want to improve our estimate of V by computing these averages:

$$V_{k+1}^{\pi}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^{\pi}(s')]$$

- Idea: Take samples of outcomes s' (by doing the action!) and average


$$sample_1 = R(s, \pi(s), s'_1) + \gamma V_k^{\pi}(s'_1)$$

$$sample_2 = R(s, \pi(s), s'_2) + \gamma V_k^{\pi}(s'_2)$$

...

$$sample_n = R(s, \pi(s), s'_n) + \gamma V_k^{\pi}(s'_n)$$

$$V_{k+1}^{\pi}(s) \leftarrow \frac{1}{n} \sum_i sample_i$$



*Almost! But we can't
rewind time to get sample
after sample from state s .*


Temporal Difference Learning

- Big idea: learn from every experience!
 - Update $V(s)$ each time we experience a transition (s, a, s', r)
 - Likely outcomes s' will contribute updates more often
- Temporal difference learning of values
 - Policy still fixed, still doing evaluation!
 - Move values toward value of whatever successor occurs: running average

Sample of $V(s)$: $sample = R(s, \pi(s), s') + \gamma V^\pi(s')$

Update to $V(s)$: $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$

Same update: $V^\pi(s) \leftarrow V^\pi(s) + \alpha(sample - V^\pi(s))$



- Values do not always converge with a fixed α , it needs to be decayed
- We need to decay it as $O(1/t)$, where t is the iteration number, but not too quickly!

Example: TD Learning

States

	A	
B	C	D
	E	

Assume: $\gamma = 1$, $\alpha = 1/2$

Observed Transitions

B, east, C, -2

	0	
0	0	8
	0	

C, east, D, -2

	0	
-1	0	8
	0	

	0	
-1	3	8
	0	

$$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha [R(s, \pi(s), s') + \gamma V^\pi(s')]$$


Problems with TD-Learning

- TD value learning is a model-free way to do policy evaluation, mimicking Bellman updates with running sample averages
- Once we have the values, can we get to actions? Can we do policy extraction?

$$\pi^*(s) = \arg_a \max(\sum_{s'} (P(s'|s, a) (R(s, a, s') + \gamma V^*(s'))))$$

- We cannot get a new policy in a model-free way!
- Idea: learn Q-values, not values
- Makes action selection model-free too!

$$\pi(s) = \arg \max_a Q(s, a)$$



Detour: Q-Values

- Value iteration: find successive values
 - Start with $V_0(s) = 0$
 - Given V_k , calculate the $k+1$ values for all states:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- But Q-values are more useful for Model-Free RL, so compute them instead
 - Start with $Q_0(s, a) = 0$
 - Given Q_k , calculate the $k+1$ q-values for all q-states:

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q_k(s', a')]$$

Q-Learning

- Q-Learning: sample-based Q-value iteration

$$Q_{k+1}(s, a) = \sum_{s'} P(s'|s, a) \left(R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right)$$

- Learn $Q(s, a)$ values as you go
 - Receive a sample (s, a, s', r)
 - Consider your old estimate: $Q(s, a)$
 - Consider your new sample estimate:
 $sample = R(s, a, s') + \gamma \max_{a'} (Q(s', a'))$
 - Incorporate the new estimate into a running average:
 $Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(sample)$
 $Q(s, a) \leftarrow Q(s, a) + \alpha(sample - Q(s, a))$



Active Reinforcement Learning

- So far the agent has been following a fixed policy
- Full reinforcement learning: optimal policies
 - You don't know the transitions $T(s,a,s')$ / $P(s' | s,a)$
 - You don't know the rewards $R(s,a,s')$
 - You choose the actions now
 - Goal: learn the optimal policy / values
- In this case:
 - Learner makes choices! Which action to take?
 - Fundamental tradeoff: exploration vs. exploitation
 - This is NOT offline planning! You take actions in the world and find out what happens...


Model Based Version

- Learn a model and calculate the values using a fixed policy i.e. do passive RL
- Extract a new policy from the values
- Repeat
- Is this a good idea?
- No, it is not! Learning a model with this greedy strategy is a bad idea
 - Learned model will not be a good approximation of the real model
 - As a result, the extracted policy will be suboptimal in the real environment
- This approach “**exploits**” what we already know
- We also want to “**explore**” as opposed to following the policies to get a better sense of the real model which will in turn let us get more reward!

What about Q-Learning?

- Same arguments go for Q-Learning, we need to “**explore**” enough to get good policies.
- Amazing result: Q-learning converges to optimal policy with enough exploration -- even if you're acting sub-optimally!
- This is called **off-policy learning**
- Caveats:
 - You must explore enough
 - You must eventually make the learning rate small enough
 - ... but not decrease it too quickly
 - Basically, in the limit, it doesn't matter how you select actions (if there is exploration)

Exploration vs. Exploitation



How to Explore?

- Simplest: random actions (ϵ -greedy)
 - Every time step, flip a coin
 - With (small) probability ϵ , act randomly
 - With (large) probability $1-\epsilon$, act on current policy
- Problems with random actions?
 - You do eventually explore the space, but keep thrashing around once learning is done
 - One solution: lower ϵ over time
 - Another solution: exploration functions

How to Explore? – Exploration Functions

- When to explore?
 - Random actions: explore a fixed amount
 - Better idea: explore areas whose badness is not (yet) established, eventually stop exploring
- Exploration function
 - Takes a value estimate u and a visit count n , and returns an optimistic utility, e.g.
$$f(u, n) = u + k/n \text{ or } f(u, n) \begin{cases} R^+, & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$
 - This makes the agent to think that there is high reward in the unexplored regions and forces it to explore those regions
- In essence, reward the unseen states more

Update Equations with Exploration Functions

- Value update (Model-based)

$$V(s) = \max_a \left[f \left(\underbrace{\sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma V(s'))}_{\text{Sum of discounted rewards}}, \underbrace{N(s, a)}_{\text{state-action counter}} \right) \right]$$

- Do not forget to update the models as well!
- Q-value Update (Model free)
 - Regular: $Q(s, a) \leftarrow Q(s, a) + \alpha \left(R(s, a, s') + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$
 - Modified: $Q(s, a) \leftarrow Q(s, a) + \alpha \left(R(s, a, s') + \gamma \max_{a'} (f(Q(s', a'), N(s', a'))) - Q(s, a) \right)$
- These updates propagate the “bonus” back to states that lead to unknown states as well!
- At a given state, s , the agent performs the action that maximizes $V(s)$ or $Q(s, a)$

Some Other Ways to Pick the Action

- ϵ -greedy:
 - With some small chance (ϵ) pick a random action
 - Otherwise pick the action that gives the highest Q-value


- Softmax: Pick the action with probability given by:

$$P_{\pi}(a|s) = \frac{e^{Q(s,a)/\tau}}{\sum_{a'} e^{Q(s,a')/\tau}}$$

- This encourages the agent to pick more promising actions (high Q-value) more often
 - Need to eventually decrease τ
 - $\tau = \infty$: totally random, $\tau = 0$: totally greedy
 - Sample the actions from $P_{\pi}(a|s)$
- Look up “Upper Confidence Bounds” for more sophisticated methods
- There are also other exploration function ideas

Regret

- Even if you learn the optimal policy, you still make mistakes along the way!
- Regret is a measure of your total mistake cost: the difference between your (expected) rewards, including youthful suboptimality, and optimal (expected) rewards
- Minimizing regret goes beyond learning to be optimal – it requires optimally learning to be optimal
- Example: random exploration and exploration functions both end up optimal, but random exploration has higher regret




Minimizing regret is out of scope for this course but search for “Multi-Armed Bandits” if you are interested

Model-Free Learning

- Model-free (temporal difference) learning
 - Experience world through episodes

$$(s, a, r, s', a', r', s'', a'', r'', s'''' \dots)$$

- Update estimates each transition (s, a, r, s')
- Over time, updates will mimic Bellman updates



Q-Learning

- We'd like to do Q-value updates to each Q-state:

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right]$$

- But can't compute this update without knowing T, R
- Instead, compute average as we go
 - Receive a sample transition (s,a,r,s')
 - This sample suggests

$$Q(s, a) \approx r + \gamma \max_{a'} Q(s', a')$$

- But we want to average over results from (s,a) (Why?)
- So keep a running average

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) \left[r + \gamma \max_{a'} Q(s', a') \right]$$

Exploration: Back to ϵ -greedy

- Simplest: random actions (ϵ -greedy)
 - Every time step, flip a coin
 - With (small) probability ϵ , act randomly
 - With (large) probability $1-\epsilon$, act on current policy
- Yields the *soft* policy (given m actions)

$$\pi(a|s) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{m} & \text{if } a = \underset{a}{\operatorname{argmax}}(Q^{\pi_t}(s, a)) \\ \frac{\epsilon}{m} & \text{otherwise} \end{cases}$$

Q-Learning with ϵ -greedy Action Selection

- Estimation/Target policy π is greedy w.r.t. $Q(S, A)$

$$\pi(s_{t+1}) = \underset{a'}{\operatorname{argmax}}(Q(s_{t+1}, a'))$$

- Chose the action via the ϵ -greedy policy with respect to $Q(S, A)$

$$\pi(a|s) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{m} & \text{if } a = \underset{a}{\operatorname{argmax}}(Q^{\pi_t}(s, a)) \\ \frac{\epsilon}{m} & \text{otherwise} \end{cases}$$

- Use this action and observe the transition
- The target sample becomes:

$$sample = R_t + \gamma \underset{a'}{\max}(Q(s_{t+1}, a'))$$

Q-Learning Pseudocode

```
function Q-Learning()  
    //Q(s,a): Q-value estimates  
     $Q_w() \leftarrow \text{init}()$  //randomly initialize all values except terminals  
    forall episodes  
         $s \leftarrow s_0$  //get the start state, can change from episode to episode  
        forall steps or until  $s$  is terminal  
             $a \leftarrow \text{sampleTransition}(s)$  //e.g.  $\epsilon$ -greedy using  $Q$   
             $r, s' \leftarrow \text{doAction}(a)$  //observe a transition  
             $y \leftarrow r + \gamma \max_{a'} (Q(s', a'))$  //Q(s',a')=0 if s' is terminal  
             $Q(s, a) \leftarrow Q(s, a) + \alpha(y - Q(s, a))$  //or  $(1 - \alpha)Q(s, a) + \alpha y$   
             $s \leftarrow s'$ 
```

Not shown: ϵ and α scheduling

Q-Learning Properties

- Amazing result: Q-learning converges to optimal policy -- even if you're acting suboptimally!
- This is called **off-policy learning**
- Caveats:
 - You have to explore enough
 - You have to eventually make the learning rate small enough
 - ... but not decrease it too quickly
 - Basically, in the limit, it doesn't matter how you select actions (given exploration)
- Exploration – Exploitation Tradeoff

Model Based or Model Free

Unknown MDP: Model-Based


Goal	Technique
Compute V^* , Q^* , π^*	VI/PI on approx. MDP
Evaluate a fixed policy π	PE on approx. MDP

Unknown MDP: Model-Free

Goal	Technique
Compute V^* , Q^* , π^*	Q-learning
Evaluate a fixed policy π	Value Learning



- In my opinion the answer is domain/problem dependent
- In practice, both approaches have solved interesting problems
- In some problems, model is already available or can be reliably estimated so a model-free approach does not make sense
- In some problems we do not want to re-calculate the policy every time we get a new percept, so a model-based approach does not make sense

Approximate Q-Learning



Generalizing Across States

- Basic Q-Learning keeps a table of all q-values
- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all in training
 - Too many states to hold the q-tables in memory
- Instead, we want to generalize:
 - Learn in a smaller space from experience
 - Generalize that experience to new, similar situations
 - This is the same as the fundamental idea in machine learning




Example: Pacman

Let's say we discover through experience that this state is bad:



In naïve q-learning, we know nothing about this state:




Or even this one!



Feature-Based Representations

- Solution: describe a state using a vector of features (properties)
 - Features are functions from states to real numbers (often 0/1) that capture important properties of the state
 - Example features:
 - Distance to closest ghost
 - Distance to closest dot
 - Number of ghosts
 - $1 / (\text{dist to dot})^2$
 - Is Pacman in a tunnel? (0/1)
 - etc.
 - Is it the exact state on this slide?
 - Can also describe a q-state (s, a) with features (e.g. action moves closer to food)



Linear Value Functions

- Using a feature representation, we can write a q function (or value function) for any state using a few weights:

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Advantage: our experience is summed up in a few powerful numbers
- Disadvantage: states may share features but actually be very different in value!

Approximate Q-Learning

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Q-learning with linear Q-functions:

$$\text{transition} = (s, a, r, s')$$

$$\text{difference} = \left[r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$$

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$$


Exact Q's

Approximate Q's

- Intuitive interpretation:

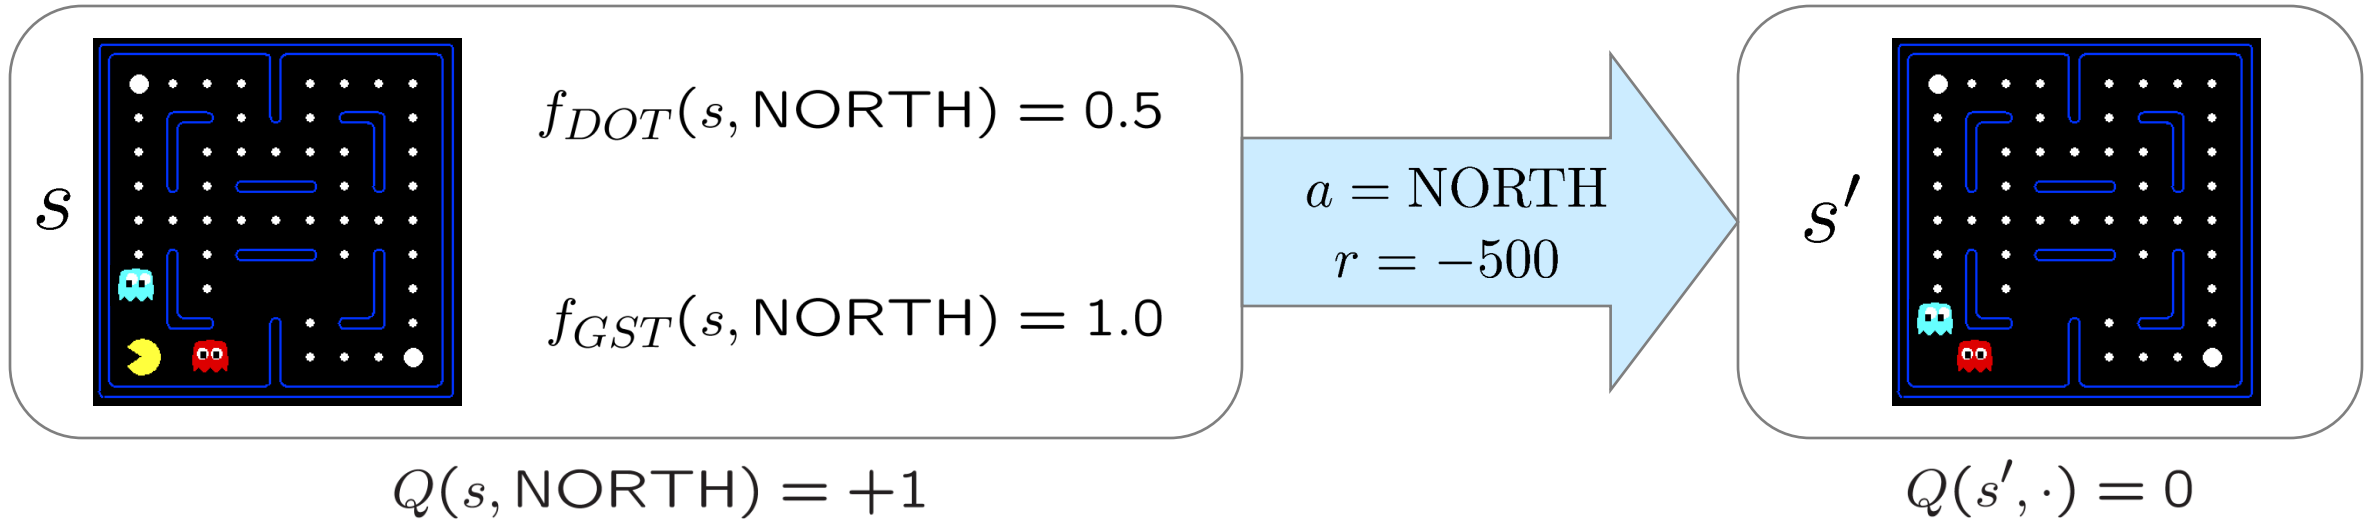
- Adjust weights of active features
- E.g., if something unexpectedly bad happens, blame the features that were ON: disprefer all states with that state's features

- Formal justification: online least squares



Example: Q-Pacman

$$Q(s, a) = 4.0 f_{DOT}(s, a) - 1.0 f_{GST}(s, a)$$



$$Q(s, \text{NORTH}) = +1$$

$$r + \gamma \max_{a'} Q(s', a') = -500 + 0$$

difference = -501

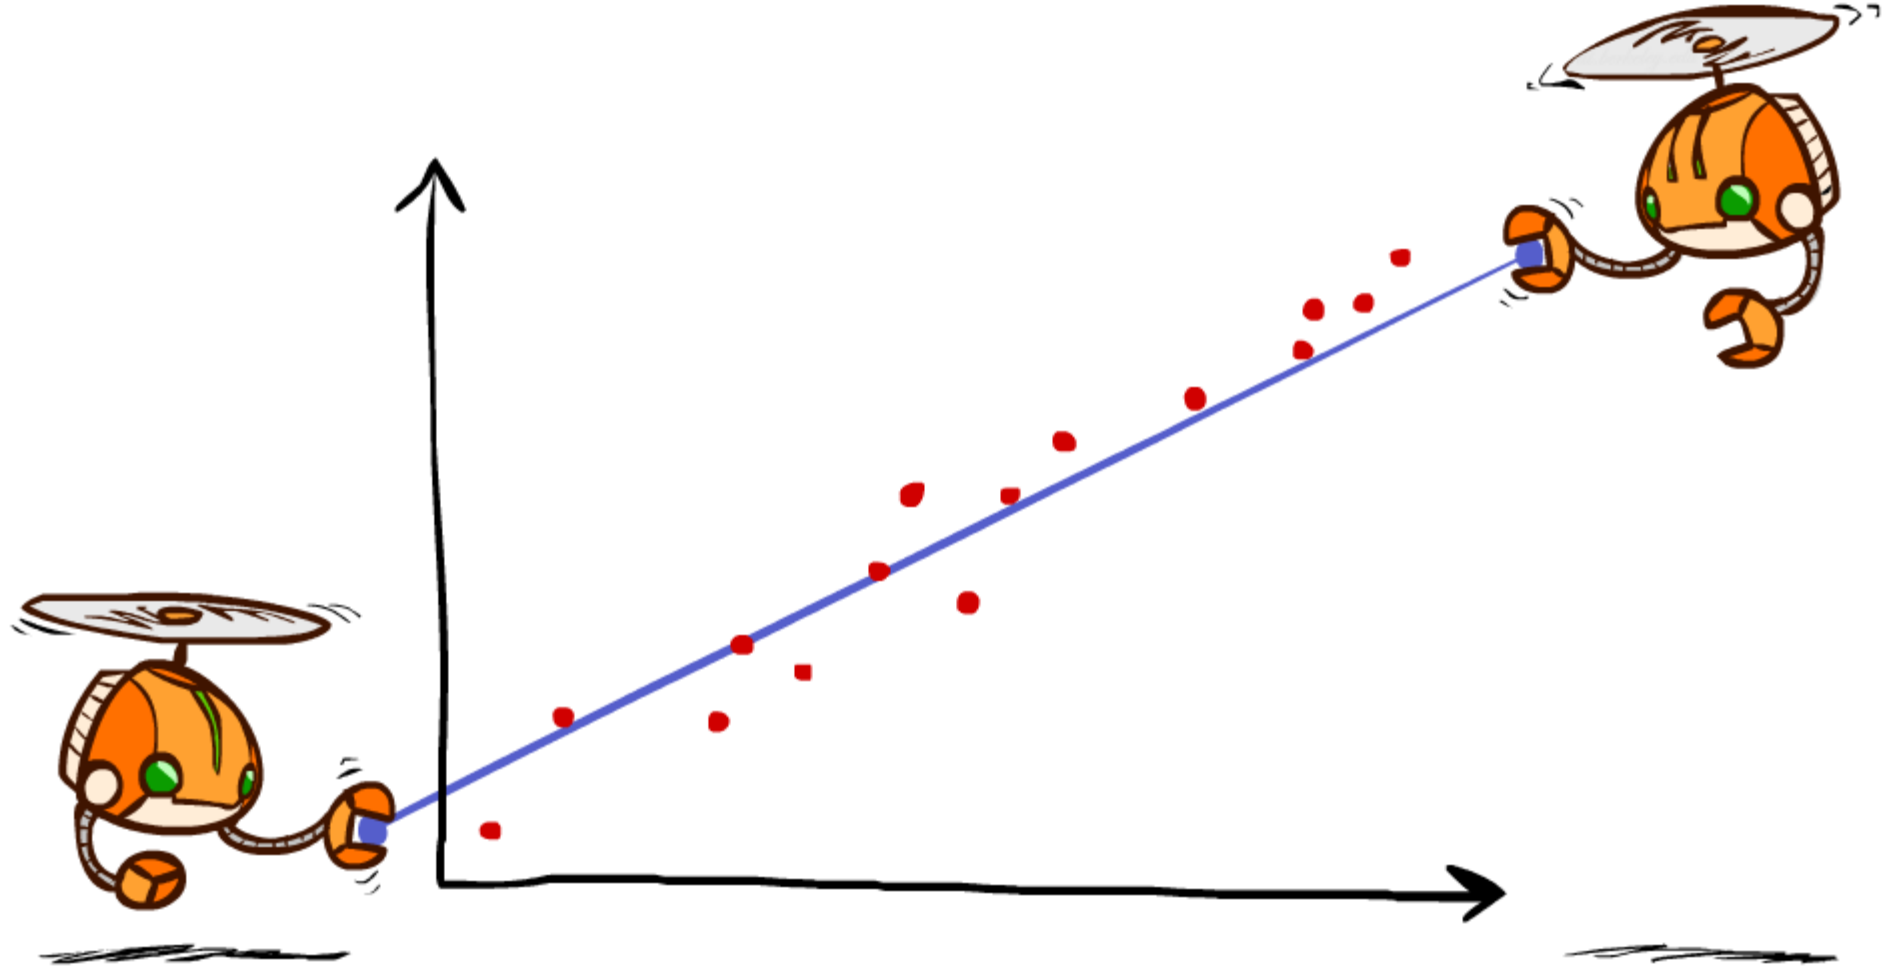


$$w_{DOT} \leftarrow 4.0 + \alpha [-501] 0.5$$


$$w_{GST} \leftarrow -1.0 + \alpha [-501] 1.0$$

$$Q(s, a) = 3.0 f_{DOT}(s, a) - 3.0 f_{GST}(s, a)$$

Q-Learning and Least Squares




Linear Approximation: Regression*



Prediction:

$$\hat{y} = w_0 + w_1 f_1(x)$$




Prediction:

$$\hat{y}_i = w_0 + w_1 f_1(x) + w_2 f_2(x)$$

Optimization: Least Squares*

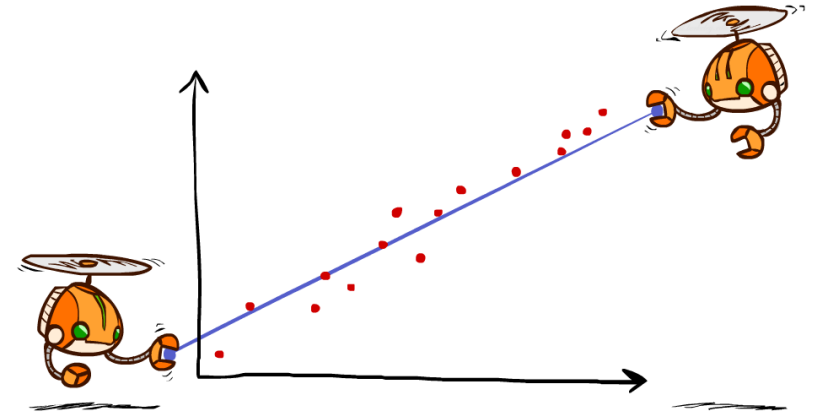
$$\text{total error} = \sum_i (y_i - \hat{y}_i)^2 = \sum_i \left(y_i - \sum_k w_k f_k(x_i) \right)^2$$



Minimizing Error*

Imagine we had only one point x , with features $f(x)$, target value y , and weights w :

$$\begin{aligned}\text{error}(w) &= \frac{1}{2} \left(y - \sum_k w_k f_k(x) \right)^2 \\ \frac{\partial \text{error}(w)}{\partial w_m} &= - \left(y - \sum_k w_k f_k(x) \right) f_m(x) \\ w_m &\leftarrow w_m + \alpha \left(y - \sum_k w_k f_k(x) \right) f_m(x)\end{aligned}$$



Approximate q update explained:

$$w_m \leftarrow w_m + \alpha \left[\underbrace{r + \gamma \max_a Q(s', a')}_{\text{“target”}} - \underbrace{Q(s, a)}_{\text{“prediction”}} \right] f_m(s, a)$$

“target”

“prediction”

Q-Learning with Linear Function Approximation

```
function Q-Learning_LinFunApprox()  
    //Q_w(s,a): approximate Q-value =  $w^T f(s,a) = w_1 f_1(s,a) + \dots + w_n f_n(s,a)$   
    w ← init() //random  
    forall episodes  
        s ← s_0 //get the start state, can change from episode to episode  
        forall steps or until s is terminal  
            a ← sampleTransition(s) //e.g.  $\epsilon$ -greedy using Q_w  
            r, s' ← doAction(a) //observe a transition  
            y ← r +  $\gamma \max_{a'} (Q_w(s', a'))$  //Q_w(s',a')=0 if s' is terminal  
            w ← w +  $\alpha_w (y - Q_w(s,a)) f(s,a)$  //this is a vector operation  
            s ← s'
```

Not shown: ϵ and α_w scheduling

Policy Search

- Problem: often the feature-based policies that work well (win games, maximize utilities) aren't the ones that approximate V / Q best
 - E.g. your value functions from project 2 were probably horrible estimates of future rewards, but they still produced good decisions
 - Q-learning's priority: get Q-values close (modeling)
 - Action selection priority: get ordering of Q-values right (prediction)
- Solution: learn policies that maximize rewards, not the values that predict them
- Policy search: start with an ok solution (e.g. Q-learning) then fine-tune by hill climbing on feature weights

Reinforcement Learning

- Exact RL:
 - Too big to represent and to learn
 - No generalization between similar states (bigger issue!)
- Approximate RL (aka value function approximation):
 - Extract features from the state (compression)
 - Represent V or Q with a function approximator (mostly parametric, usually linear or NN-based)
 - Allows for **generalization** between states if “features are good”
- Policy Search
 - At the end of the day, we mostly care about the policy not the values
 - Learning parameters of a policy representation

Policy Search

