

Programming on Multi-GPUs

Didem Unat


dunat@ku.edu.tr

<http://parcorelab.ku.edu.tr>

Programming Multiple GPUs

- Programming Multiple GPUs on a Host
 - Using a Single Thread
 - Using Multiple Threads
 - Using Multiple Processes
- Programming GPUs on Multiple Hosts
 - Only possible using multiple processes (with MPI)

Programming Multiple GPUs on a Host



- Split/distribute the application data
- Submit a kernel for each device
- Set a device before submitting operations
 - Check number of GPUs with: `cudaGetDeviceCount()`
 - Set a device with: `cudaSetDevice()`

Using a Single Thread

```
...
size_t size = sizeof(double) * 8; // array size
...
double* h_arr; // host array
h_arr = (double*) malloc(size); // host allocation

cudaSetDevice(0);
double* d0_arr;
cudaMalloc((void **)&d0_arr, size/2); // allocation for device 0
cudaMemcpy(d0_arr, h_arr, size/2, cudaMemcpyHostToDevice);


cudaSetDevice(1);
double* d1_arr;
cudaMalloc((void **)&d1_arr, size/2); // allocation for device 1
cudaMemcpy(d1_arr, &h_arr[4], size/2, cudaMemcpyHostToDevice);
...
cudaSetDevice(0);
simpleAdd<<<blocksPerGrid, threadsPerBlock>>>>(d0_arr, 4);

cudaSetDevice(1);
simpleAdd<<<blocksPerGrid, threadsPerBlock>>>>(d1_arr, 4);
...
```

} distribute data to each device

} launch a kernel for each device

What if GPUs need to share data?



- Thread(s) can move data from one GPU to another through host
 - Not ideal, degrades performance with unnecessary host transfers
- If the system employs an **interconnect** between GPUs and supports GPUDirect, data can be shared directly
 - **GPUDirect Peer-to-Peer Direct Access:** directly access another GPU's memory
 - **GPUDirect P2P Direct Transfers:** directly transfer data from another GPU

GPUDirect Peer-to-Peer

- Check if P2P access is available with:
 - `cudaDeviceCanAccessPeer()`
- Enable P2P access with:
 - `cudaDeviceEnablePeerAccess()`
- P2P transfer data with:
 - `cudaMemcpy(..., cudaMemcpyDeviceToDevice)`
 - No need to specify a device number
 - Uses virtual addressing to identify where the data is located
 - `cudaMemcpyPeer()`
 - Need to indicate which device to transfer the data

Types of Communications

- ◆ Device to Device
 - Host-initiated Explicit Transfer(CUDA memcpy PtoP)
 - Case 1.1: P2P cudamemcpy with UVA
 - Case 1.2: P2P cudamemcpy with No-UVA
 - P2P Implicit Transfer
 - Case 2: Unified Virtual Addressing
 - ◆ Direct access to other GPU's memory from the kernel with GPUDirect P2P
 - Case 3: Unified Virtual Memory (Memcpy DtoD)
 - Case 4: Through Host
- ◆ Host To Device and Device to Host (No other device)
 - Case 5: Memcpy (CUDA memcpy DtoH and CUDA memcpy HtoD)
 - Case 6: Unified Virtual Addressing
 - Case 7: Unified Memory (Unified Memory Memcpy HtoD and Unified Memory)

Types of Communication

Device-Device (Peer-to-Peer) Communication		Host-Device Communication
Explicit	Case 1.1: cudaMemcpy with UVA Case 1.2: cudaMemcpyPeer without UVA	Case 1.3: cudaMemcpyPeer & cudaMemcpy (implicit copies through host) Case 1.4: cudaMemcpy with H2D and D2H kinds Case 4: cudaMemcpy with H2D, D2H or cudaMemcpyDefault kinds
	Case 2: Zero-copy Memory Case 3: Unified Memory	Case 5: Zero-copy Memory Case 6: Unified Memory
Peer Access Enabled (a)		Peer Access Disabled (b)

UVA Example-1

```
/* Initial Setup */
// For UVA it is imp to enable peer access
int can_access_peer_0_1, can_access_peer_1_0;
cudaSetDevice(0);
cudaDeviceEnablePeerAccess(&can_access_peer_0_1, 0, 1); //Device 0 can access Device 1
cudaSetDevice(1);
cudaDeviceEnablePeerAccess(&can_access_peer_1_0, 1, 0); //Device 1 can access Device 0

// Allocate buffers
const size_t buf_size = 1024 * 1024 * 16 * sizeof(float);
cudaSetDevice(0);
float *g0; //g0 allocated on Device 0
cudaMalloc(&g0, buf_size);
cudaSetDevice(1);
float *g1; //g1 allocated on Device 1
cudaMalloc(&g1, buf_size);
//h0 allocated and initialized on Host
float *h0;
cudaMallocHost(&h0, buf_size); // Automatically portable with UVA
for (int i=0; i<buf_size / sizeof(float); i++)
    h0[i] = float(i % 4096);
```

UVA Example -1 – cont.

```
cudaSetDevice(0); //Set to Device 0
cudaMemcpy(g0, h0, buf_size, cudaMemcpyDefault);

// Kernel launch configuration
const dim3 threads(512, 1);
const dim3 blocks((buf_size / sizeof(float)) / threads.x, 1);

// Run kernel on GPU 1, reading input from the GPU 0 buffer, writing
// output to the GPU 1 buffer
cudaSetDevice(1); //Set to Device 1
SimpleKernel<<<blocks, threads>>>(g0, g1);
cudaDeviceSynchronize();
```

HtoD memcpy
using UVA (Case 6)

Data transferred
from Host to GPU0

```
//Kernel Example
__global__ void SimpleKernel(float *src, float *dst)
{
    // Just a dummy kernel, doing enough for us to verify that everything
    // worked
    const int idx = blockIdx.x * blockDim.x + threadIdx.x;
    //Here src is g0 and dst is g1
    dst[idx] = src[idx] * 2.0f;
}
```

Implicit PtoP memcpy
(Case 2)

Data is transferred
from GPU0 to GPU1

UVA Example -1 cont.

```
// Run kernel on GPU 0, reading input from the GPU 1 buffer, writing
// output to the GPU 0 buffer
cudaSetDevice(0); //Set to Device 0
SimpleKernel<<<blocks, threads>>>(g1, g0);
cudaDeviceSynchronize();
```

```
__global__ void SimpleKernel(float *src, float *dst)
{
    // Just a dummy kernel, doing enough for us to verify that everything
    // worked
    const int idx = blockIdx.x * blockDim.x + threadIdx.x;
    //Here src is g1 and dst is g0
    dst[idx] = src[idx] * 2.0f;
}
```

```
// Copy data back to the host and verify
cudaMemcpy(h0, g0, buf_size, cudaMemcpyDefault);
for (int i=0; i<buf_size / sizeof(float); i++){
    // Re-generate input data and apply 2x '* 2.0f' computation of both kernel runs
    if (h0[i] != float(i % 4096) * 2.0f * 2.0f) break;
}
```

Implicit PtoP memcpy
(Case 2)

Data transferred from
GPU1 to GPU0

DtoH memcpy
using UVA (Case 6)

Data transferred
from GPU0 to Host

P2P Memcpy with UVA Example

```
// P2P memcpy() benchmark
for (int i=0; i<100; i++){
// With UVA we don't need to specify source and target
// devices, the runtime figures this out by itself from the
// pointers
// Ping-pong copy between GPUs
    if (i % 2 == 0){
        cudaMemcpy(g1, g0, buf_size, cudaMemcpyDefault);
    }
    else
    {
        cudaMemcpy(g0, g1, buf_size, cudaMemcpyDefault);
    }
}
```

Explicit PtoP memcpy
(Case 1.1)

Data transferred from
GPU0 to GPU1

Explicit PtoP memcpy
(Case 1.1)

Data transferred from
GPU1 to GPU0

No UVA Example

```
// Allocate buffers
const size_t buf_size = 1024 * 1024 * 16 * sizeof(float);
// g0_src and g0_dst allocated on GPU 0
cudaSetDevice(0);
float *g0_src, *g0_dst;
cudaMalloc(&g0_src, buf_size);
cudaMalloc(&g0_dst, buf_size);

// g1_src and g1_dst allocated on GPU 1
cudaSetDevice(1);
float *g1_src, *g1_dst;
cudaMalloc(&g1_src, buf_size);
cudaMalloc(&g1_dst, buf_size);

//h0 allocated and initialized on Host
float *h0;
h0 = (float *)malloc(buf_size);
for (int i=0; i<buf_size / sizeof(float); i++)
    h0[i] = float(i % 4096);
```

Case 1.2 & 5: P2P Memcpy (No UVA) Example

```
cudaSetDevice(0); //Set to Device 0
//HtoD copy using cudaMemcpyHostToDevice instead of cudaMemcpyDefault
cudaMemcpy(g0_src, h0, buf_size, cudaMemcpyHostToDevice);
```

HtoD memcpy (Case 5)
Data transferred from
Host to GPU0

```
// Kernel launch configuration
const dim3 threads(512, 1);
const dim3 blocks((buf_size / sizeof(float)) / threads.x, 1);
```

```
cudaSetDevice(1); //Set to Device 1
//Perform P2P memcpy from Device 0 to Device 1 before kernel launch
cudaMemcpyPeer(g1_src, 1, g0_src, 0, buf_size);
SimpleKernel<<<blocks, threads>>>(g1_src, g1_dst);
checkCudaErrors(cudaDeviceSynchronize());
```

Explicit PtoP memcpy
(Case 1.2)
Data transferred from
GPU0 to GPU1

```
cudaSetDevice(0); //Set to Device 0
//Perform P2P memcpy from Device 1 to Device 0 before kernel launch
cudaMemcpyPeer(g0_src, 0, g1_dst, 1, buf_size);
SimpleKernel<<<blocks, threads>>>(g0_src, g0_dst);
cudaDeviceSynchronize();
```

Explicit PtoP memcpy
(Case 1.2)
Data transferred from
GPU1 to GPU0

No UVA Example

```
// Copy data back to the host and verify
//DtoH copy using cudaMemcpyDeviceToHost instead of
cudaMemcpyDefault
cudaMemcpy(h0, g0_dst, buf_size, cudaMemcpyDeviceToHost);

//Verification on Host
int error_count = 0;
for (int i=0; i<buf_size / sizeof(float); i++)
{
    // Re-generate input data and apply 2x '* 2.0f' computation of
both
    // kernel runs
    if (h0[i] != float(i % 4096) * 2.0f * 2.0f)
    {
        if (error_count++ > 10)
        {
            break;
        }
    }
}
```

DtoH memcpy
(Case 5)

Data transferred
from GPU0 to Host

UVM Example

```
const size_t buf_size = 1024 * 1024 * 16 * sizeof(float); //Allocate Buffers
float *g0;
cudaMallocManaged(&g0, buf_size);
float *g1;
cudaMallocManaged(&g1, buf_size);

for (int i=0; i<buf_size / sizeof(float); i++)
{ g0[i] = float(i % 4096); }

// Kernel launch configuration
const dim3 threads(512, 1);
const dim3 blocks((buf_size / sizeof(float)) / threads.x, 1);

// Run kernel on GPU 1
cudaSetDevice(1)); // Set to Device 1
SimpleKernel<<<blocks, threads>>>(g0, g1);
cudaDeviceSynchronize();
```

Buffer allocated and initialized at host in unified memory

```
__global__ void SimpleKernel(float *src, float *dst){
    // Just a dummy kernel, doing enough for us to verify that everything
    // worked
    const int idx = blockIdx.x * blockDim.x + threadIdx.x;
    //Both src and dst are transferred to Device 1 on CPU page fault
    dst[idx] = src[idx] * 2.0f;
}
```

Implicit HtoD Unified Memory Memcpy (Case 7)
g0 and g1 transferred from Host to Device 1

UVM Example – cont.

```
// Run kernel on GPU 0
cudaSetDevice(0); //Set to Device 0
SimpleKernel<<<blocks, threads>>>(g1, g0);
cudaDeviceSynchronize();
```

```
__global__ void SimpleKernel(float *src, float *dst){
// Just a dummy kernel, doing enough for us to verify that everything worked
const int idx = blockIdx.x * blockDim.x + threadIdx.x;
//Both src and dst are brought from Device 1 to Device 0 on GPU page fault
    dst[idx] = src[idx] * 2.0f;
}
```


Implicit PtoP Unified
Memory memcpy (Case 3)

```
int error_count = 0;
for (int i=0; i<buf_size / sizeof(float); i++)
{ // Re-generate input data and apply 2x '* 2.0f' computation of both kernel runs
    if (g0[i] != float(i % 4096) * 2.0f * 2.0f)
    { if (error_count++ > 10) break; }
}
```

Implicit DtoH Unified
Memory Memcpy (Case 7)
g0 transferred from
Device 0 to Host

Lab 10 – Multi-GPU 1D Jacobi

- Simple 1D Jacobi to illustrate the peer to peer communication between two GPU devices
- Input data is divided between two GPUs
- Every iteration, GPUs send single data point to each other



Lab 10 – Multi-GPU 1D Jacobi

- The code is missing some sections in *runJacobi1DCUDA*
 - Need to add necessary **malloc** and **cudaMemcpy** for Device #1
 - For Device 0, these are already set for you as an example
 - Need to add kernel launch for Device 1
 - See the example for Device 0
 - Need to synchronize before you perform communication
 - Next, you need to perform halo updates between two devices
 - See next slide for details

Lab 10 – Multi-GPU 1D Jacobi

- The code is missing the necessary data transfer between the GPUs.
- Fix the code so that
 - First it performs transfers with peer transfer using **cudaMemcpyPeer**
 - Second it performs transfers with memory copy using **cudaMemcpy**
 - Finally it performs the transfers through the host
- Usage for cudaMemcpyPeer
 - http://horacio9573.no-ip.org/cuda/group_CUDA_MEMORY_g046702971bc5a66d9bc6000682a6d844.html
- Usage for cudaMemcpy
 - http://horacio9573.no-ip.org/cuda/group_CUDA_MEMORY_g48efa06b81cc031b2aa6fdc2e9930741.html#g48efa06b81cc031b2aa6fdc2e9930741

Using Multiple Threads

```
// thread 0
cudaSetDevice(0);
double* d_arr;
cudaMalloc((void **)&d_arr, size/2);
cudaMemcpy(d_arr, h_arr, size/2, cudaMemcpyHostToDevice);
...
simpleAdd<<<blocksPerGrid, threadsPerBlock>>>(d_arr, 4);
...
```

thread 0
gets its
own GPU

```
// thread 1
cudaSetDevice(1);
double* d_arr;
cudaMalloc((void **)&d_arr, size/2);
cudaMemcpy(d_arr, &h_arr[4], size/2,
cudaMemcpyHostToDevice);
...
simpleAdd<<<blocksPerGrid, threadsPerBlock>>>(d_arr, 4);
...
```

thread 1
gets its
own GPU

Multi-GPUs with Multi-threads

- We can access multiple GPUs with multiple threads using OpenMP

```
#pragma omp parallel num_threads(2)
{
    int tid = omp_get_thread_num();

    if (tid == 0)
        // do something
    else
        // do something
}
```


Outline

- ✓ Programming Multiple GPUs on a Host
 - ✓ Using a Single Thread
 - ✓ Using Multiple Threads
 - ✓ **Using Multiple Processes**
- Programming GPUs on Multiple Hosts

Using Multiple Processes on a Host

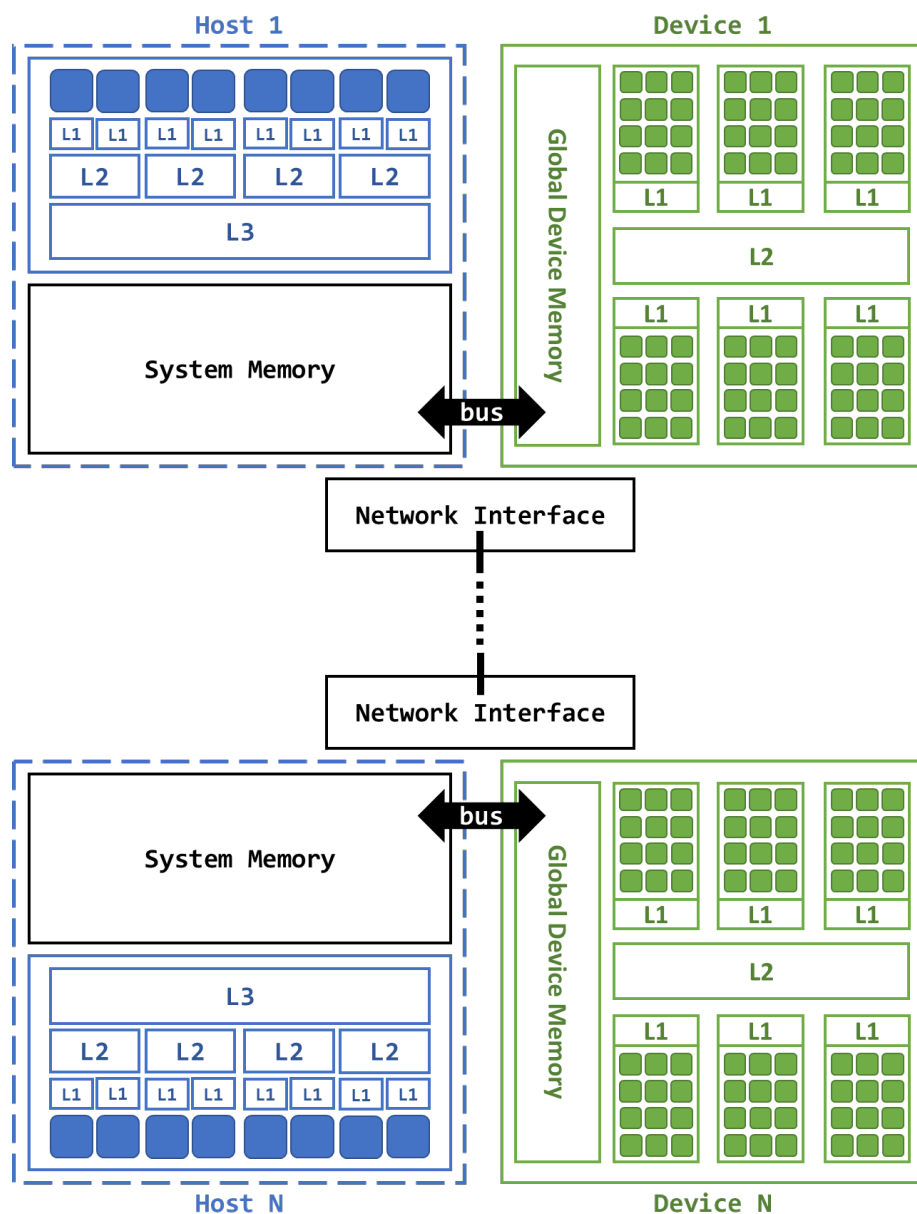
- Similar to programming with multiple threads but
 - Each process either gets its own GPU
 - Or shares a GPU
 - Use message passing (MPI) for communication among processes

What if GPUs need to share data?



- Processes have their own memory address space
 - They cannot directly share pointers as threads
- CUDA Inter-Process Communication (IPC) helps share device memory pointers across processes
 - Get an IPC handle with: `cudaIpcGetMemHandle()`
 - Open IPC handle with: `cudaIpcOpenMemHandle()`

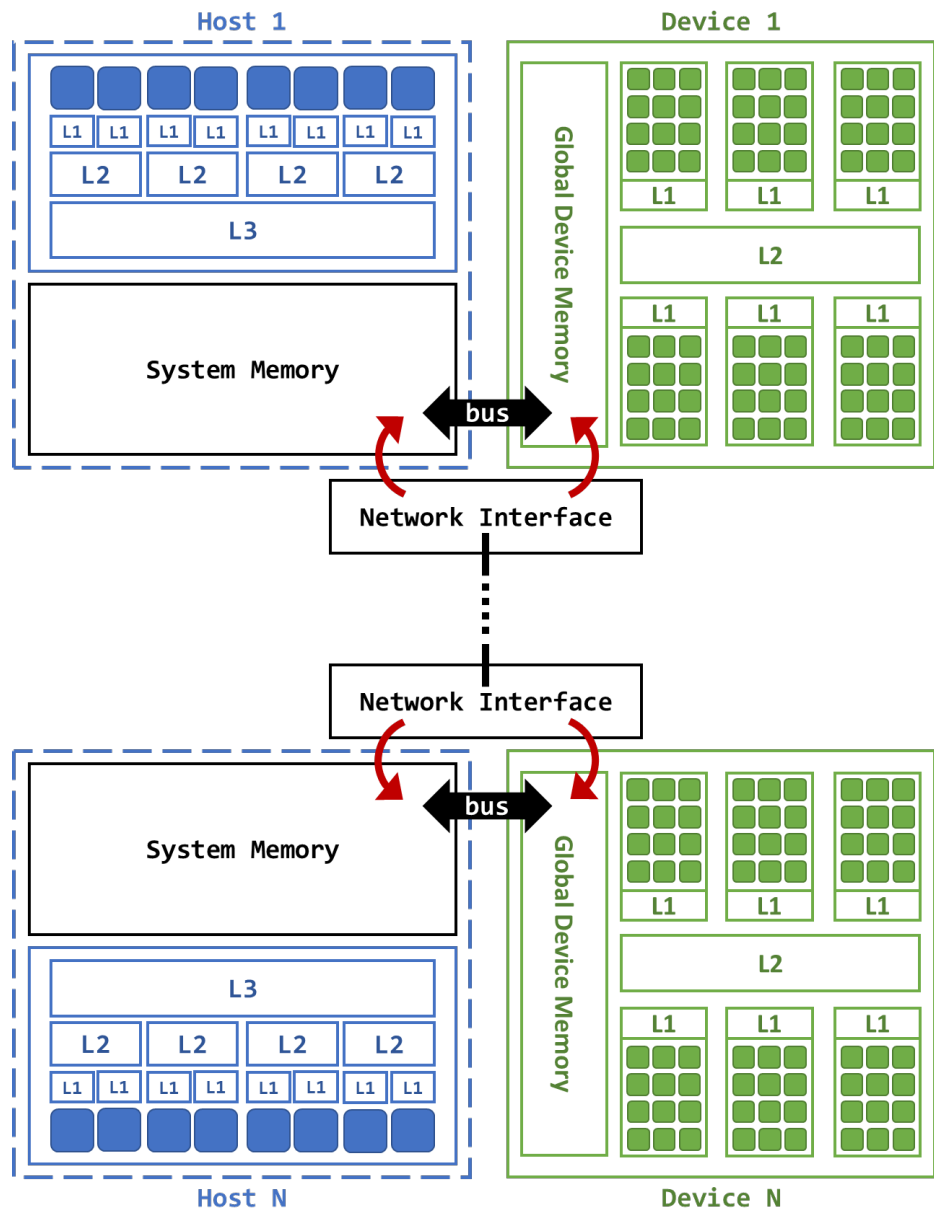
Programming GPUs on Multiple Hosts



- Processes are the only option for programming GPUs on multiple hosts
 - Combine CUDA with MPI
- Communication happens through network interfaces

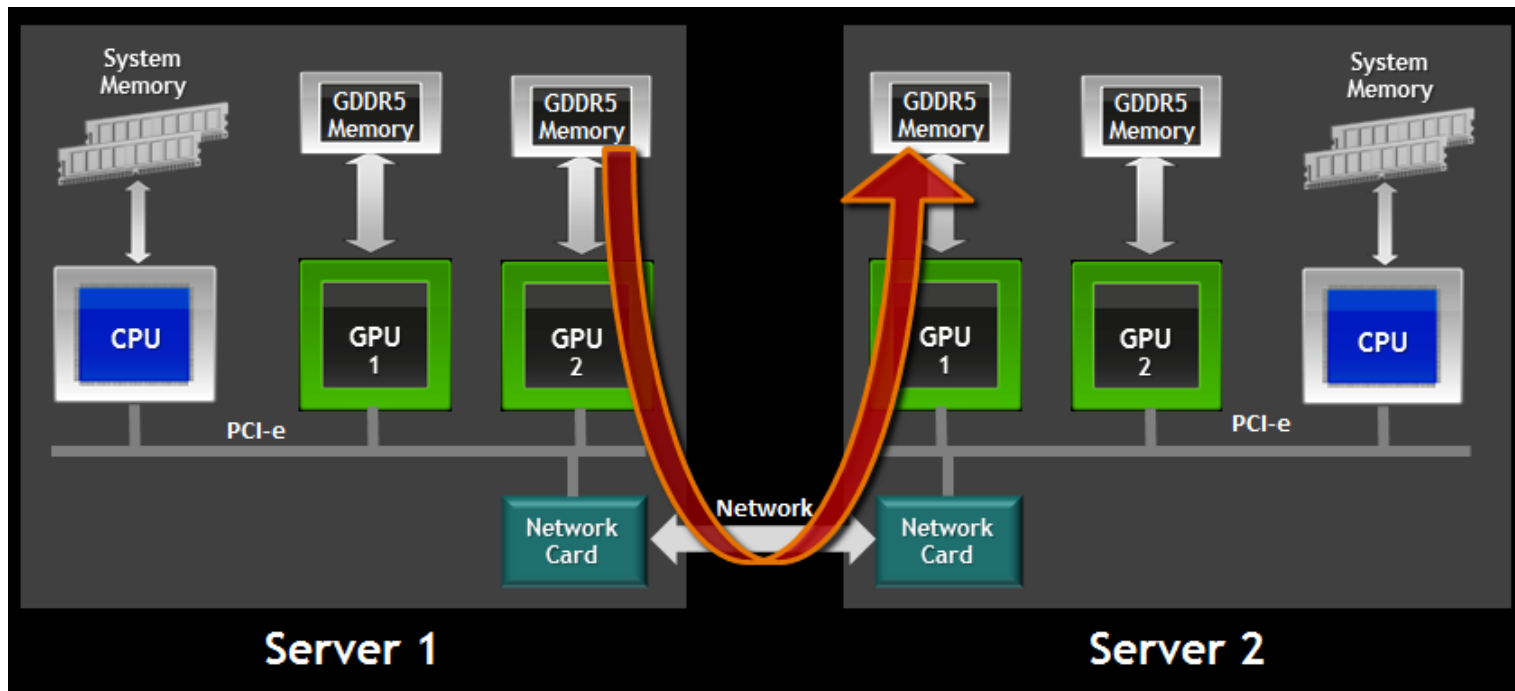
GPUDirect Remote Direct Memory Access (RDMA)

- Specific network interfaces with GPUDirect RDMA support helps avoid unnecessary host transfers
 - It is low-level!
 - There are CUDA-aware MPI implementations
 - Programmer doesn't deal with IPC, etc.



GPUDirect RDMA

- Prevents unnecessary transfers to hosts with
 - Remote Direct Memory Access (RDMA)
 - GPUDirect (pinned to pageable copy on host)



src: NVIDIA

GPU-Aware MPI Library

- Transfer data directly to/from CUDA device memory via MPI calls
- Standard MPI interfaces used for unified data movement
- Takes advantage of Unified Virtual Addressing (\geq CUDA 4.0)
- Overlaps data movement from GPU with RDMA transfers
- Avoids copies through the host

At Sender:

```
MPI_Send(s_devbuf, size, ...);
```

At Receiver:

```
MPI_Recv(r_devbuf, size, ...);
```

Conclusion

- ✓ Programming Multiple GPUs on a Host
 - ✓ Using a Single Thread
 - ✓ Overlap is possible but code size increases
 - ✓ Using Multiple Threads
 - ✓ Use OpenMP
 - ✓ Using Multiple Processes
 - ✓ Can use Multi-process service (MPS) if processes need to share a GPU
 - ✓ **GPUDirect Peer-to-Peer Direct Access:** directly access another GPUs memory
 - ✓ **GPUDirect P2P Direct Transfers:** directly transfer data from another GPU
- Programming GPUs on Multiple Hosts
 - Need to use MPI
 - Transfers to GPU -> host -> host -> GPU
 - GPUDirect RDMA enables GPU <-> GPU transfers