


Partitioned Global Address Space (PGAS)

Didem Unat


COMP 429/529 Parallel Programming

Shared-Memory Programming Model

- Shared-address space programming
 - Threads communicate through shared memory as opposed to messages
 - Threads coordinate through synchronization (also through shared memory).




Recall shared memory system
(can be either UMA, NUMA)



Message Passing Programming Model

- Programs execute as a set of P processes (user specifies P)
- Each processor has its own private address space
 - Processors share data by *explicitly* sending and receiving information (message passing)
 - Coordination is built into message passing primitives (message send and message receive)



Shared Memory vs. Message Passing

Shared Memory

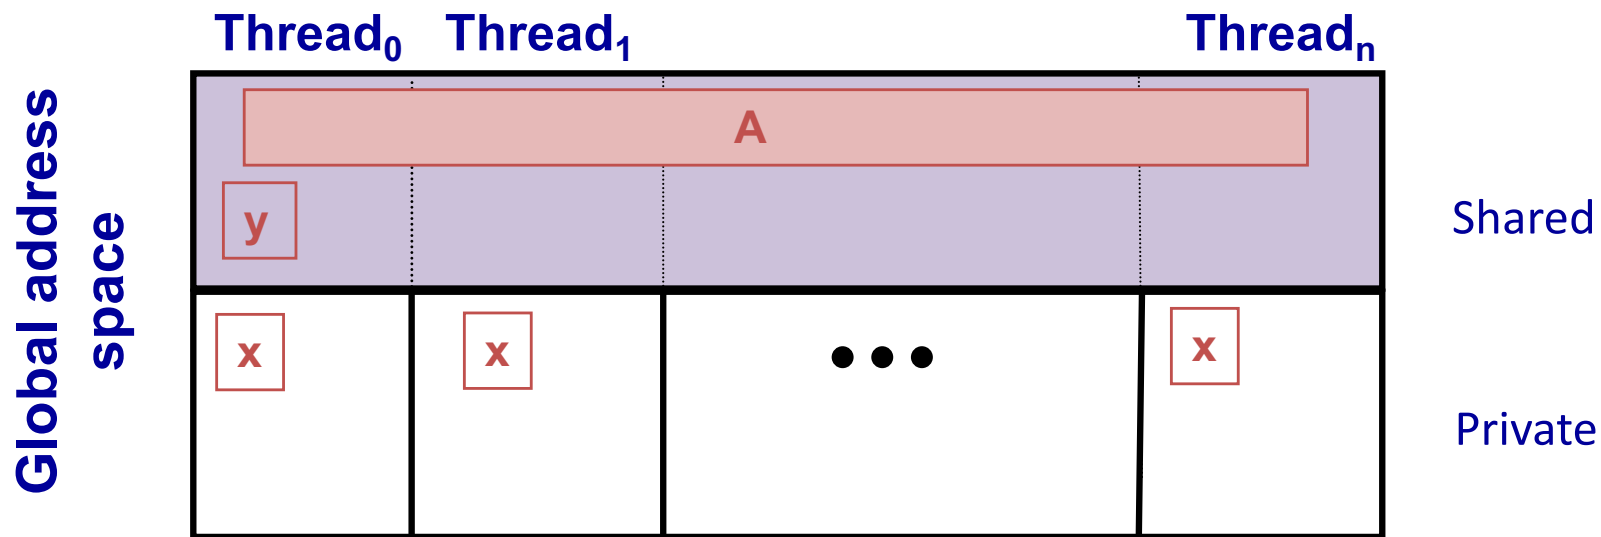
- Advantage: Convenience
 - Can share data structures
 - Just annotate loops
 - Closer to serial code
- Disadvantages
 - No locality control
 - Does not scale
 - Race conditions

Message Passing

- Advantage: Scalability
 - Locality control
 - Communication is all explicit in code (cost transparency)
- Disadvantage
 - Need to rethink entire application / data structures
 - Lots of tedious pack/unpack code

Partitioned Global Address Space (PGAS) Model

- *Global address space*: thread may directly read/write remote data
 - Hides the distinction between shared/distributed address spaces
- *Partitioned*: data is designated as **local or global**
 - Does not hide this: critical for locality and scaling



- Each process has declared a private copy of variable **x**
- Process 0 has declared a shared variable **y**
- A shared array **A** is distributed across the global address space

UPC++

- Unified Parallel C (**UPC++**) is an extension of the C++ programming language designed
 - Provides a uniform programming model for both shared and distributed memory hardware
- A number of threads working independently in a SPMD fashion
 - Number of threads specified at compile-time or run-time; available as program variable **THREADS**
 - **MYTHREAD** specifies thread index (**0..THREADS-1**)
 - **upc_barrier** is a global synchronization: all wait
 - There is a form of parallel loop that we will see later
- There are other PGAS languages such as X10, Chapel and OpenShmem

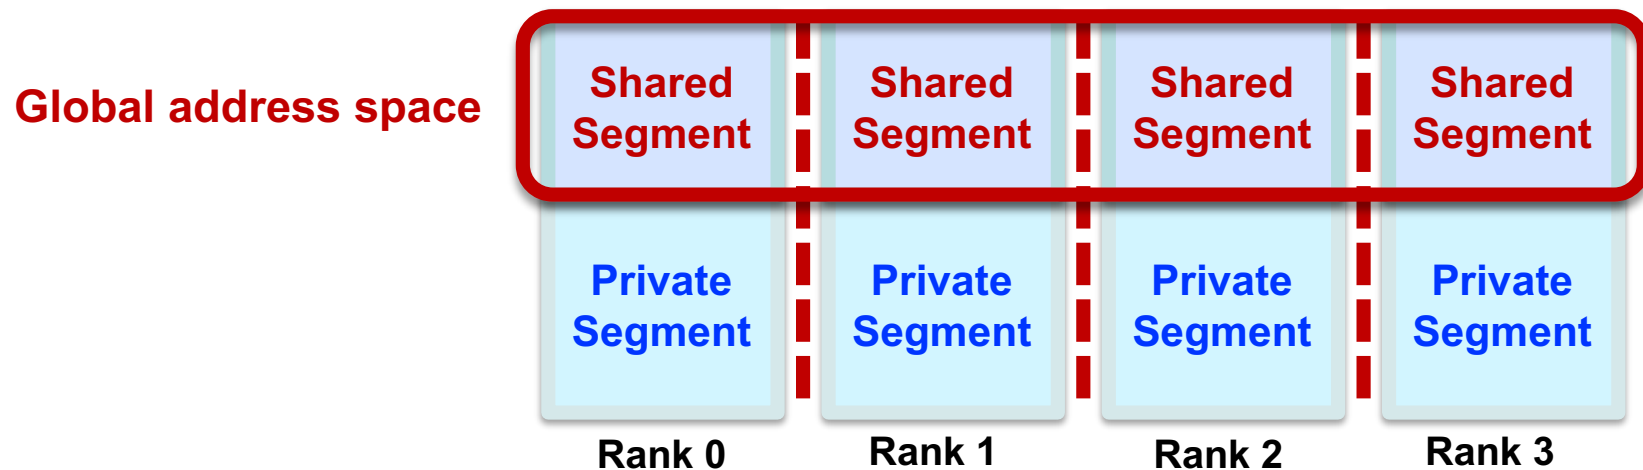
Preliminaries for UPC++

- UPC++ is an SPMD model, like MPI
- A distributed memory parallel computer is an abstract collection of processing elements, an indivisible computing resource with local memory, AKA a *rank*
- The number of ranks is fixed throughout the program




A Partitioned Global Address Space

- Global Address Space (private address spaces only)
 - Ranks read & write *shared segments* of memory, set up at program initialization time, via *1 sided communication*
 - explicit sender & receiver to copy data between private address spaces
- Partitioned (Not applicable to message passing)
 - *Global pointers* to shared segments have affinity to a particular rank
 - Explicitly managed by the programmer to optimize for locality



Example: Monte Carlo Pi Calculation

- Estimate Pi by throwing darts at a unit square
- Calculate percentage that fall in the unit circle
 - Area of square = $r^2 = 1$
 - Area of circle quadrant = $\frac{1}{4} * p r^2 = p/4$
- Randomly throw darts at x,y positions
- If $x^2 + y^2 < 1$, then point is inside circle
- Compute ratio:
 - # points inside / # points total
 - $p = 4 * \text{ratio}$



Independent Estimation of Pi in UPC

```
void main(int argc, char **argv) {
```

```
    int i, hits, trials = 0;  
    double pi;
```

Each thread gets its own copy of these variables

```
    if (argc != 2) trials = 1000000;  
    else trials = atoi(argv[1]);
```

Each thread can use input arguments

```
    srand(MYTHREAD*17);
```

Initialize random in math library

```
    for (i=0; i < trials; i++)  
        hits += hit();  
    pi = 4.0*hits/trials;  
    printf("PI estimated to %f.", pi);
```

```
}
```

Each thread calls “hit” function separately

Helper Code for Pi in UPC

- Required includes:

```
#include <stdio.h>
#include <math.h>
#include <upc.h>
```

- Function to throw dart and calculate where it hits:


```
int hit() {
    int const rand_max = 0xFFFFFFFF;
    double x = ((double) rand()) / RAND_MAX;
    double y = ((double) rand()) / RAND_MAX;
    if ((x*x + y*y) <= 1.0) {
        return(1);
    } else {
        return(0);
    }
}
```

Private vs. Shared Variables in UPC

- Normal C variables and objects are allocated in the private memory space for each thread.

- Shared variables are allocated only once, **with thread 0**

```
shared int ours; // use sparingly: performance  
int mine;
```



Pi in UPC: Shared Memory Style

```
shared int hits;
```

shared variable to record hits

```
void main(int argc, char **argv) {  
    int i, my_trials = 0;  
    int trials = atoi(argv[1]);    divide work up evenly  
    my_trials = (trials + THREADS - 1)/THREADS;  
    srand(MYTHREAD*17);  
    for (i=0; i < my_trials; i++)  
        hits += hit();    accumulate hits  
    upc_barrier;  
    if (MYTHREAD == 0) {  
        printf("PI estimated to %f.", 4.0*hits/trials);  
    }  
}
```

There is a bug in the program. What is the problem with this program?

Pi in UPC: Shared Array Version

- Alternative fix to the race condition
- Have each thread update a separate counter:
 - But do it in a shared array
 - Have one thread compute sum

```
shared int all_hits [THREADS];
```

all_hits is shared by
all processors, just
as hits was

```
main(int argc, char **argv) {
```

```
    ... declarations and initialization code omitted
```

```
    for (i=0; i < my_trials; i++)
```

```
        all_hits[MYTHREAD] += hit();
```

update element with local
affinity

```
    upc_barrier;
```

```
    if (MYTHREAD == 0) {
```

```
        for (i=0; i < THREADS; i++) hits += all_hits[i];
```

```
        printf("PI estimated to %f.", 4.0*hits/trials);
```

```
    }
```

```
}
```

Synchronization - Locks

- Locks in UPC are represented by an opaque type:

`upc_lock_t`

- Locks must be allocated before use:

`upc_lock_t *upc_all_lock_alloc(void) ;`

allocates 1 lock, pointer to all threads

`upc_lock_t *upc_global_lock_alloc(void) ;`

allocates 1 lock, pointer to one thread

- To use a lock:

`void upc_lock(upc_lock_t *l)`

`void upc_unlock(upc_lock_t *l)`

use at start and end of critical region

- Locks can be freed when not in use

`void upc_lock_free(upc_lock_t *ptr) ;`

Pi in UPC: Shared Memory Style

Use locks like pthreads, but use shared accesses judiciously

```
shared int hits;
```

one shared scalar variable

```
main(int argc, char **argv) {
```

other private variables

```
int i, my_hits, my_trials = 0;
```

```
upc_lock_t *hit_lock = upc_all_lock_alloc();
```

create a lock

```
int trials = atoi(argv[1]);
```

```
my_trials = (trials + THREADS - 1)/THREADS;
```

```
srand(MYTHREAD*17);
```

```
for (i=0; i < my_trials; i++)
```

```
    my_hits += hit();
```

accumulate hits
locally

```
upc_lock(hit_lock);
```

```
hits += my_hits;
```

```
upc_unlock(hit_lock);
```

```
upc_barrier;
```

accumulate across
threads using a lock

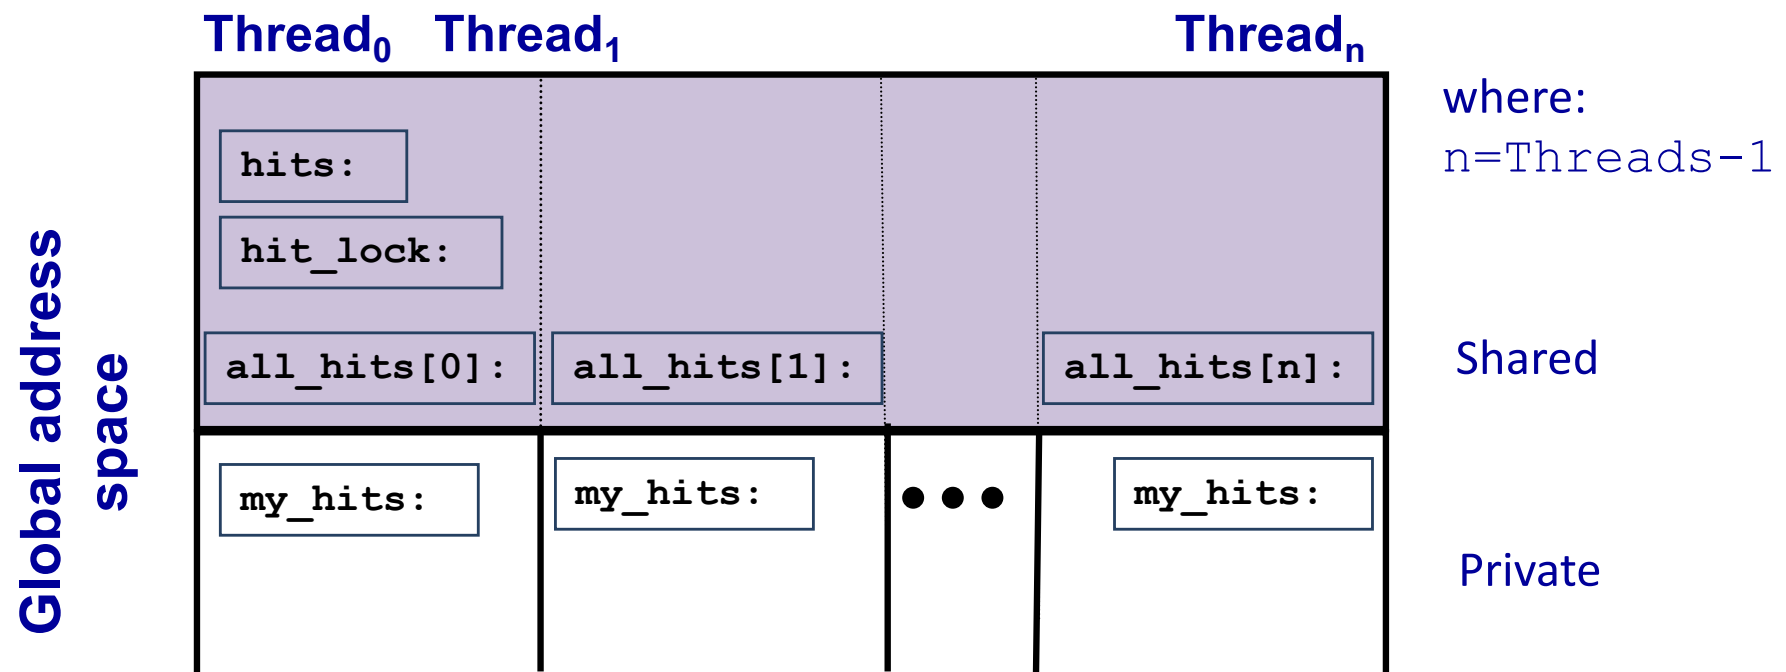
```
if (MYTHREAD == 0)
```

```
    printf("PI: %f", 4.0*hits/trials);
```

```
}
```


Recap: Private vs. Shared Variables in UPC

- We saw several kinds of variables in the pi example
 - Private scalars (`my_hits`)
 - Shared scalars (`hits`)
 - Shared arrays (`all_hits`)
 - Shared locks (`hit_lock`)



One-Sided Communication

- The basic idea of one-sided communication models is to decouple data movement with process synchronization
- Should be able to move data without requiring that the remote process synchronize
- Each process exposes a part of its memory to other processes
- Other processes can directly read from or write to this memory

1-sided Communication in UPC++

- A rank can read or write memory in another address space via a *global pointer*, *target* is unaware
- Also called Remote Memory Access (RMA, MPI, too)
- Unlike message passing, no explicit sender & receiver


```
global_ptr<T> gptr1 = . . . ;  
T t1 = rget(gptr1).wait()  
rput(gptr2,x,2).wait();
```

Wait returns with result to the
requestor when rget completes

Global address space


Start the get

Private memory



One-Sided vs Two-Sided

one-sided put message



- A one-sided put/get message can be handled directly by a network interface with RDMA support
 - Avoid interrupting the CPU or storing data from CPU (preposts)
 - Put() and Get()
- A two-sided messages needs to be matched with a receive to identify memory address to put data
 - Send() and Receive()

Remote Memory Access (RMA)


- If communication *pattern* is not known *a priori*, but the data locations are known, the send-receive model requires an extra step to determine how many sends- receives to issue
- RMA, however, can handle it easily because only the origin or target process needs to issue the put or get call
- This makes dynamic communication easier to code in RMA

MPI support for One-Sided Comm

- **MPI_Win_create** exposes local memory to RMA operation by other processes in a communicator
 - Collective operation
 - Creates window object
- **MPI_Win_free** deallocates window object
- **MPI_Put** moves data from local memory to remote memory
- **MPI_Get** retrieves data from remote memory into local memory
- **MPI_Accumulate** updates remote memory using local values
- Data movement operations are non-blocking
- Subsequent synchronization on window object needed to ensure operation is complete

1-D Point Jacobi in UPC++

- Each rank allocates solution U in global memory




```
global_ptr<double> U = new_array<double>(n);
```

```
global_ptr<double> uL = ... , uR = ... ;
```

1-D Point Jacobi in UPC++

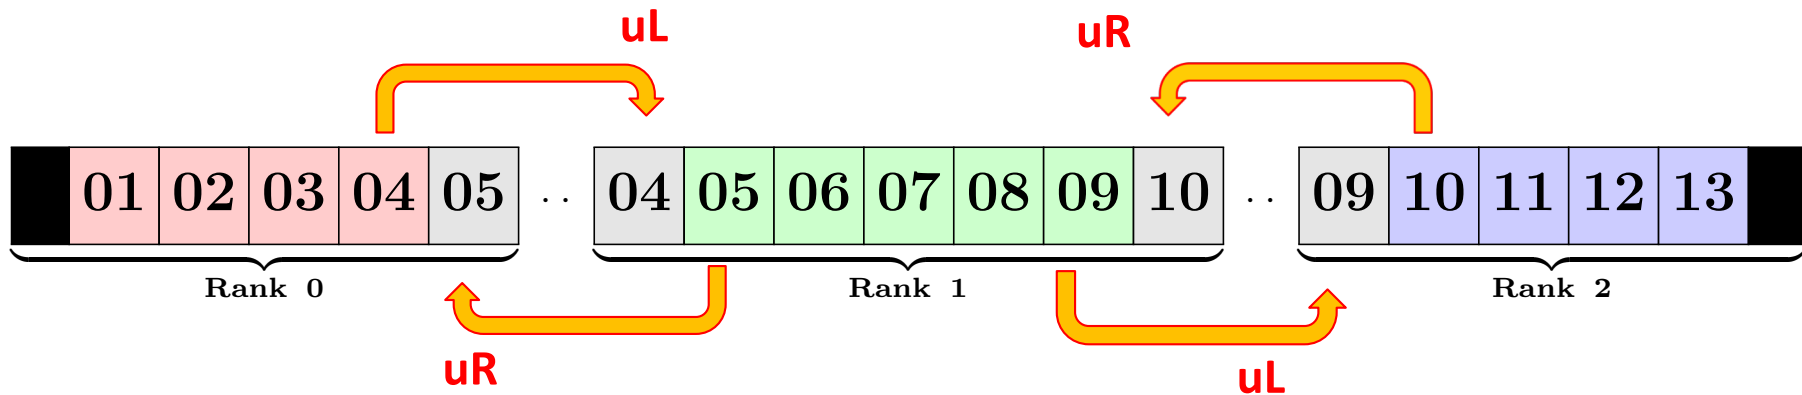
- Each rank allocates solution U in global memory
- u is a raw pointer to local data in the global segment



```
global_ptr<double> U = new_array<double>(n);  
double *u = U.local() ; // Unew allocated similarly  
global_ptr<double> uL = ... , uR = ... ;  
for (i = 1 to MaxIter ){  
    fillGhost(u,n, uL, uR);  
    Sweep(Unew, U);  
    Swap pointers...  
}
```


FillGhost in UPC++

- uL and uR have been set up previously to point to left and right neighboring solution arrays



```
fillGhost(u,n,uL,uR):
```

```
if ( !(uL.is_null()))  
    u[0] = rget(uL).wait(); //get 1-sided comm
```


```
if ( !(uR.is_null()))  
    u[n-1] = rget(uR).wait(); //get 1-sided comm
```

The UPC++ Memory Model

- UPC++ runs under a different memory model than message passing
- Like shared memory, there can be *race conditions*
- A race condition arises because we are reading and writing global storage
- The timing of the accesses can affect the outcome
- We say that we have a *non-deterministic computation* when the outcome can vary from run to run
- Message passing avoids the race condition, since data movement is coupled with synchronization
- We'll illustrate a race condition (and the solution) in 1-D Point Jacobi

Where is the Race Condition?

- Since ranks run at different rates, its possible that a neighboring rank can capture its ghost cells and move on to the next iteration ($i+1$) before its neighbors do
- A *straggler* that's still in iteration i could obtain data from a neighbor that is computing a future iteration $i+1$
- This behavior is unpredictable, & we may not observe it



```
for (i = 1 to MaxIter) {  
  fillGhost(u,n, uL, uR);  
  Sweep();  
  Swap pointers...  
}
```

```
fillGhost(u,n,uL,uR):  
  if ( !(uL.is_null()) )  
    u[0] = rget(uL).wait();  
  if ( !(uR.is_null()) )  
    u[n-1] = rget(uR).wait();
```

The Solution: a Barrier

- No rank can proceed past the barrier until all have arrived, ensuring that no rank runs ahead of others
- Barrier runs in time $\log_2(P)$
allreduce = reduce + bcast

```
for (i = 1 to MaxIter ){  
    fillGhost(u ,n, uL, uR);  
    Sweep();  
    Swap pointers...  
    barrier();  
}
```




www.shutterstock.com - 37858410



Comparing message passing with PGAS

- Recall that message passing is 2-sided
 - There is an explicit sender and receiver
 - Data movement and synchronization are coupled
 - Messages have an associated context (e.g. tag) that must be matched to handle incoming messages correctly
- Ordering guarantees are not semantically matched to the hardware
- PGAS change various communication attributes
 - Execute fewer instructions to perform a transfer (reduce α)

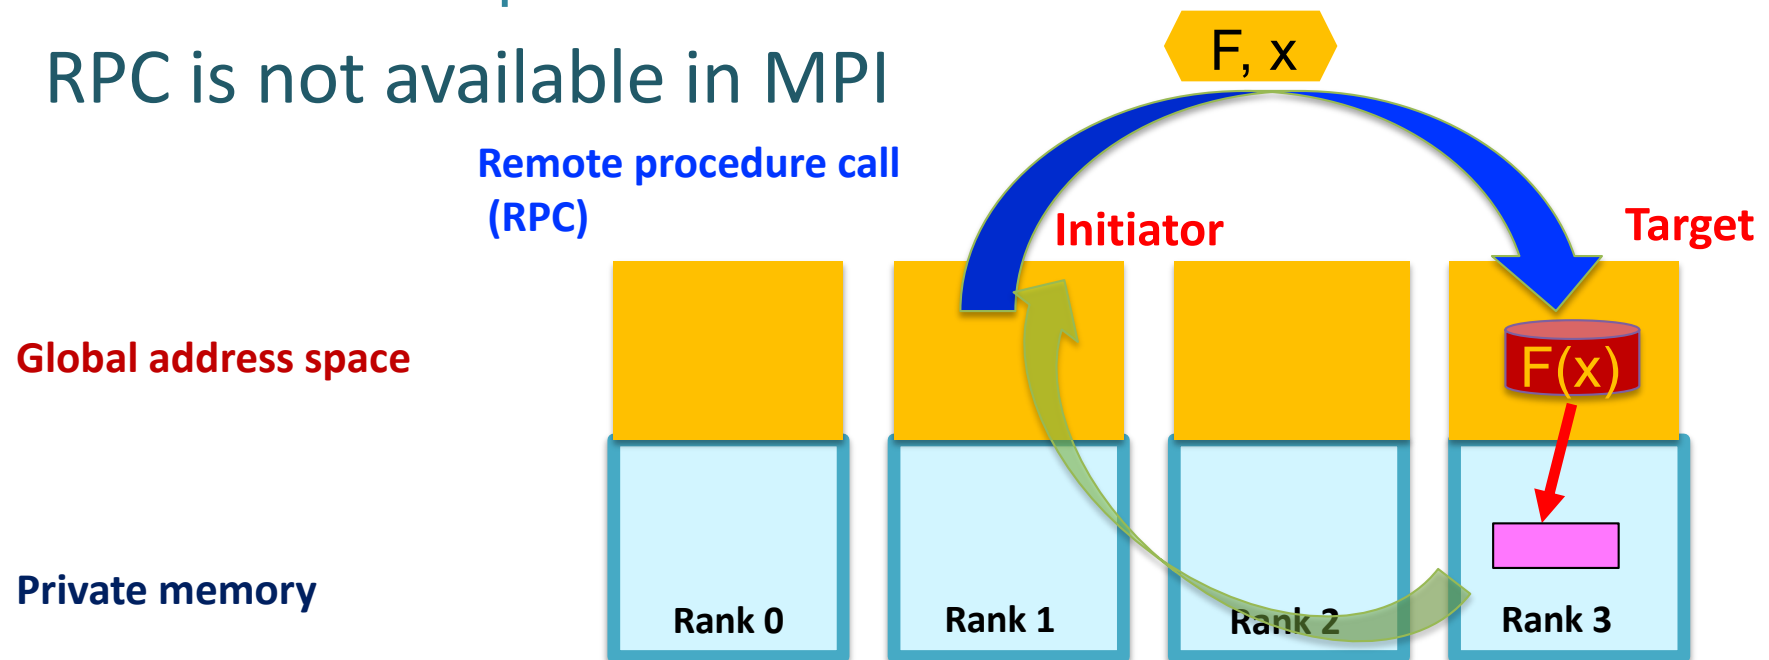


UPC++ reduces the Overheads

- RMA lets each rank directly access one another's memory via a global reference
 - We don't need to match sends to receives
 - We don't need to guarantee message ordering
 - MPI also supports RMA, too
- Looks like shared memory, so we need to handle with race conditions
 - Unlike message passing, synchronization and data movement are separate
- Technology trends provide support that benefits RMA
 - Modern network hardware provides the capability to directly access memory on another node:
Remote Direct Memory Access (RDMA)

Remote Procedure Call


- Remote Procedure call (RPC) is different from RMA
 - Executes a function on another rank (the *target*), sending any arguments
 - Returns an optional result to the initiator
- RPC is not available in MPI



Issuing a Remote Procedure Call

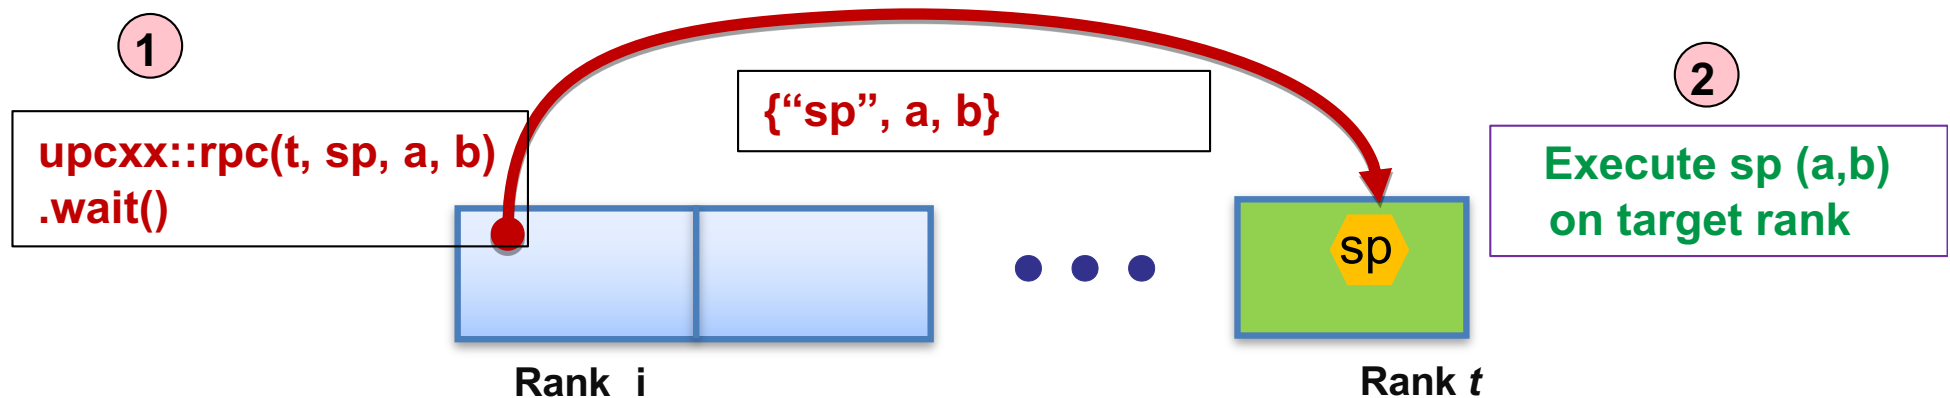
- Rank i executes `sp()` on rank t via an RPC
- `int sp(int a, int b) { return a + b; }`
- `int sum = rpc(t, sp, a, b).wait();`

1



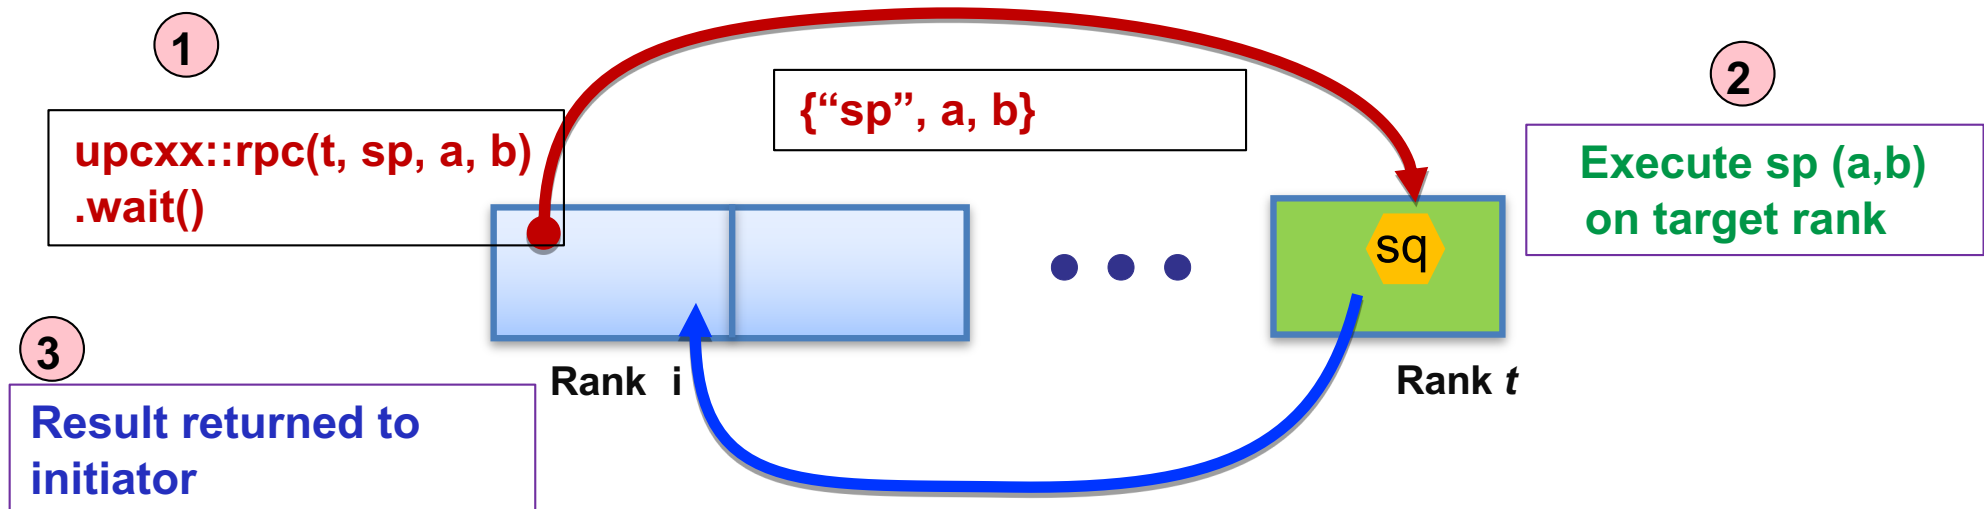
Issuing a Remote Procedure Call

- Rank i executes `sp()` on rank t via an RPC ①
`int sp(int a, int b) { return a + b; }`
- `int sum = rpc(t , sp, a , b).wait();`
- The target rank t will execute the handler function `sp()` at some future time determined at the target (“makes progress”) ②
- The details of progress are hidden from the programmer



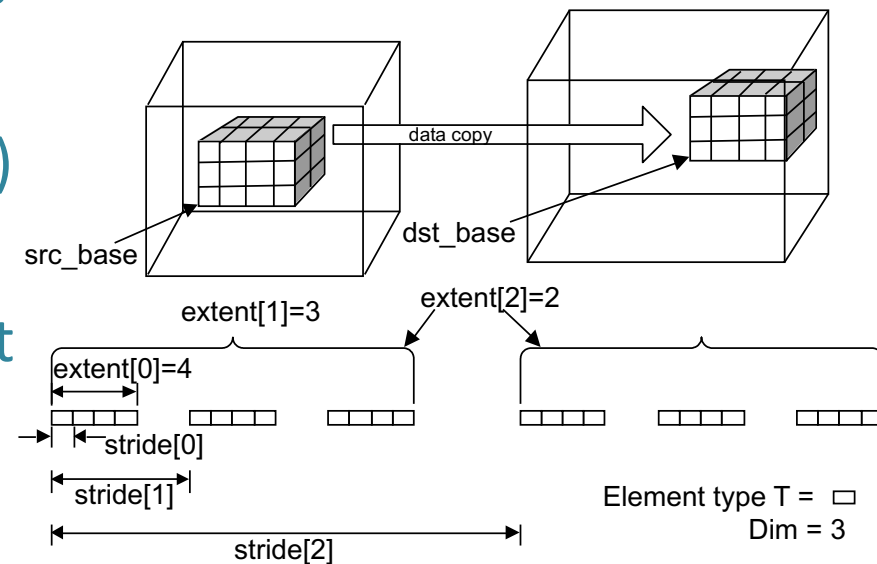
Issuing a Remote Procedure Call

- Rank i executes $sp()$ on rank t via an RPC ①
 - `int sp(int a, int b) { return a + b; }`
- `int sum = rpc(t , sp , a , b), wait();`
- The target rank t will execute the handler function $sp()$ at some future time determined at the target (“makes progress”) ②
- When the handler completes, result returned to rank i ③



Other Features of UPC++

- Remote Atomic operations
- Completions
 - Know when the source memory can be modified, when the operation has completed at the target
 - Attach an RPC to RMA completion
- Non-contiguous transfers
- Teams
(like MPI communicators)
- Memory Kinds
Unified treatment of host
& device memory



Acknowledgments

- These slides are inspired and partly adapted from
 - Kathy Yelick (Univ. of California, Berkeley, LBL)
 - Jim Demmel (Univ. of California, Berkeley)
 - William Gropp (UIUC)
 - Scott B. Baden (c) 2020 / PGAS Programming